

分布式计算


Java网络编程

Weixiong Rao 饶卫雄
Tongji University 同济大学软件学院
2023 秋季
wxrao@tongji.edu.cn

网络编程

- 网络基础知识
- Socket, Client/Server

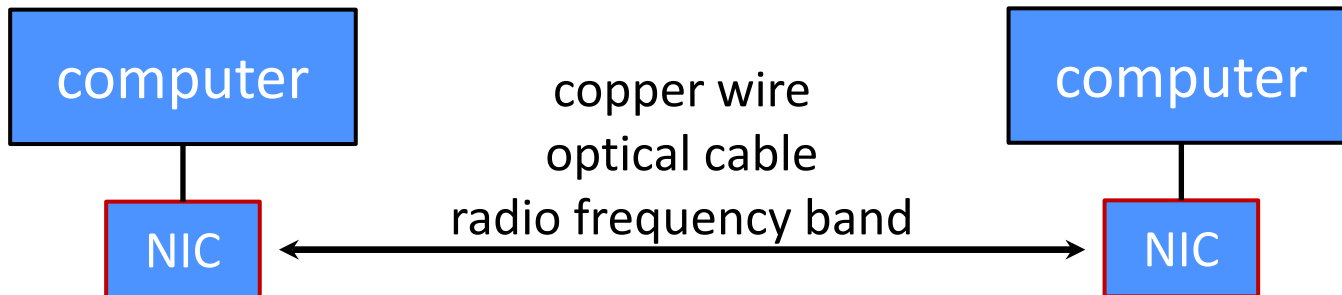

Networks From 10,000 ft



The Physical Layer


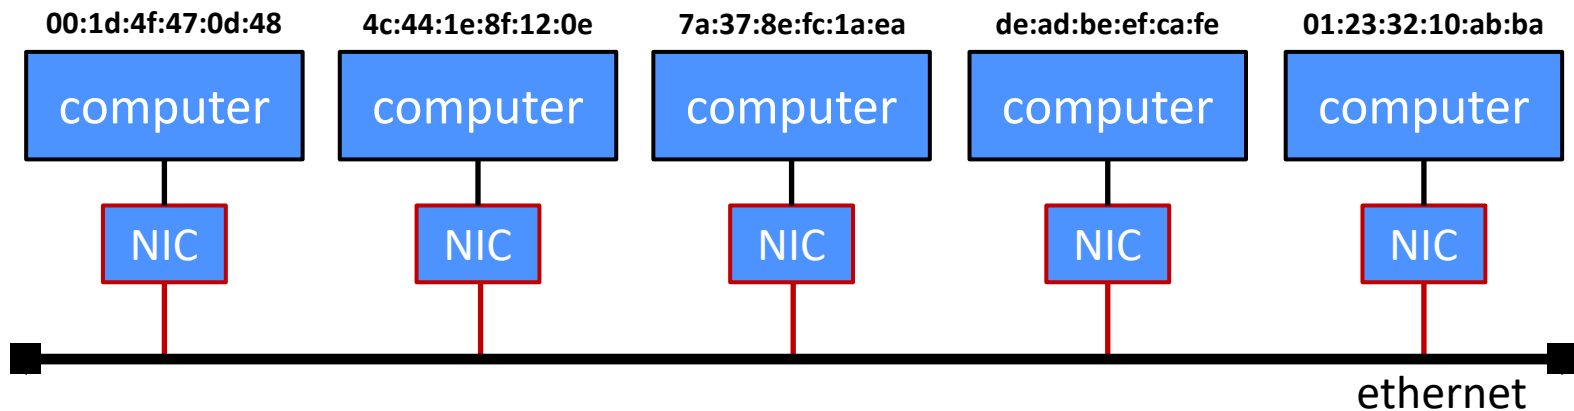
- Individual bits are modulated onto a wire or transmitted over radio
 - ◆ Physical layer specifies how bits are encoded at a signal level
 - ◆ Many choices, e.g., encode “1” as +1v, “0” as -0v; or “0”=+1v, “1”=-1v, ...

0 1 0
1



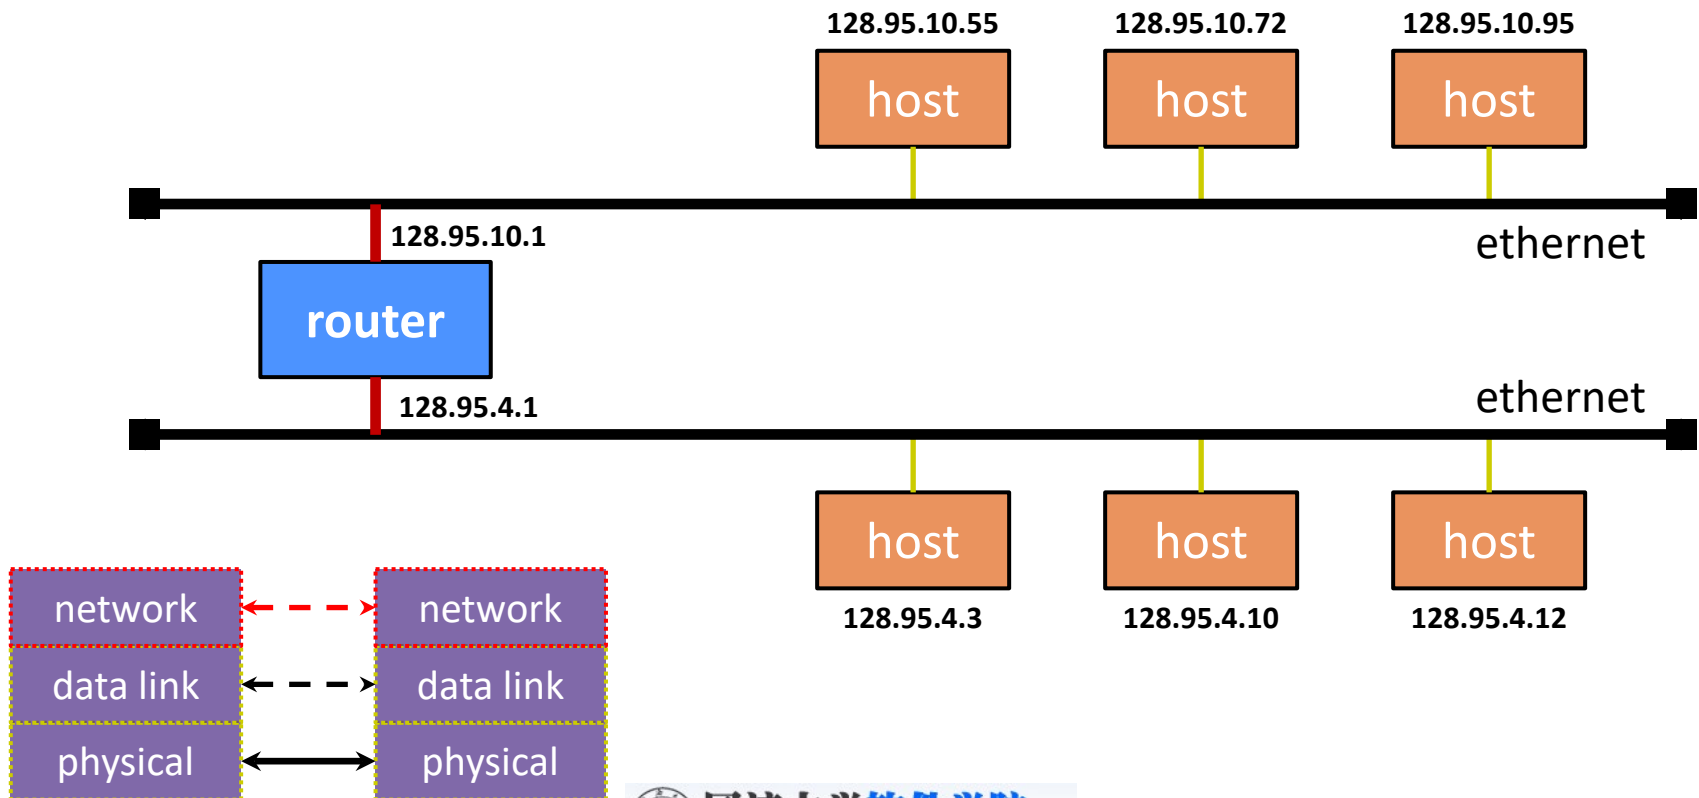
The Data Link Layer

- Multiple computers on a LAN contend for the network medium
 - ◆ Media access control (MAC) specifies how computers cooperate
 - ◆ Link layer also specifies how bits are “packetized” and network interface controllers (NICs) are addressed



The Network Layer (IP)



- Internet Protocol (IP) routes packets across multiple networks
 - ◆ Every computer has a unique IP address
 - ◆ Individual networks are connected by routers that span networks



The Network Layer (IP)

- There are protocols to:



- Let a host map an IP to MAC address on the same network
- Let a router learn about other routers to get IP packets one step closer to their destination



The Network Layer (IP)


■ Packet encapsulation:

- ◆ An IP packet is encapsulated as the payload of an Ethernet frame
- ◆ As IP packets traverse networks, routers pull out the IP packet from an Ethernet frame and plunk it into a new one on the next network



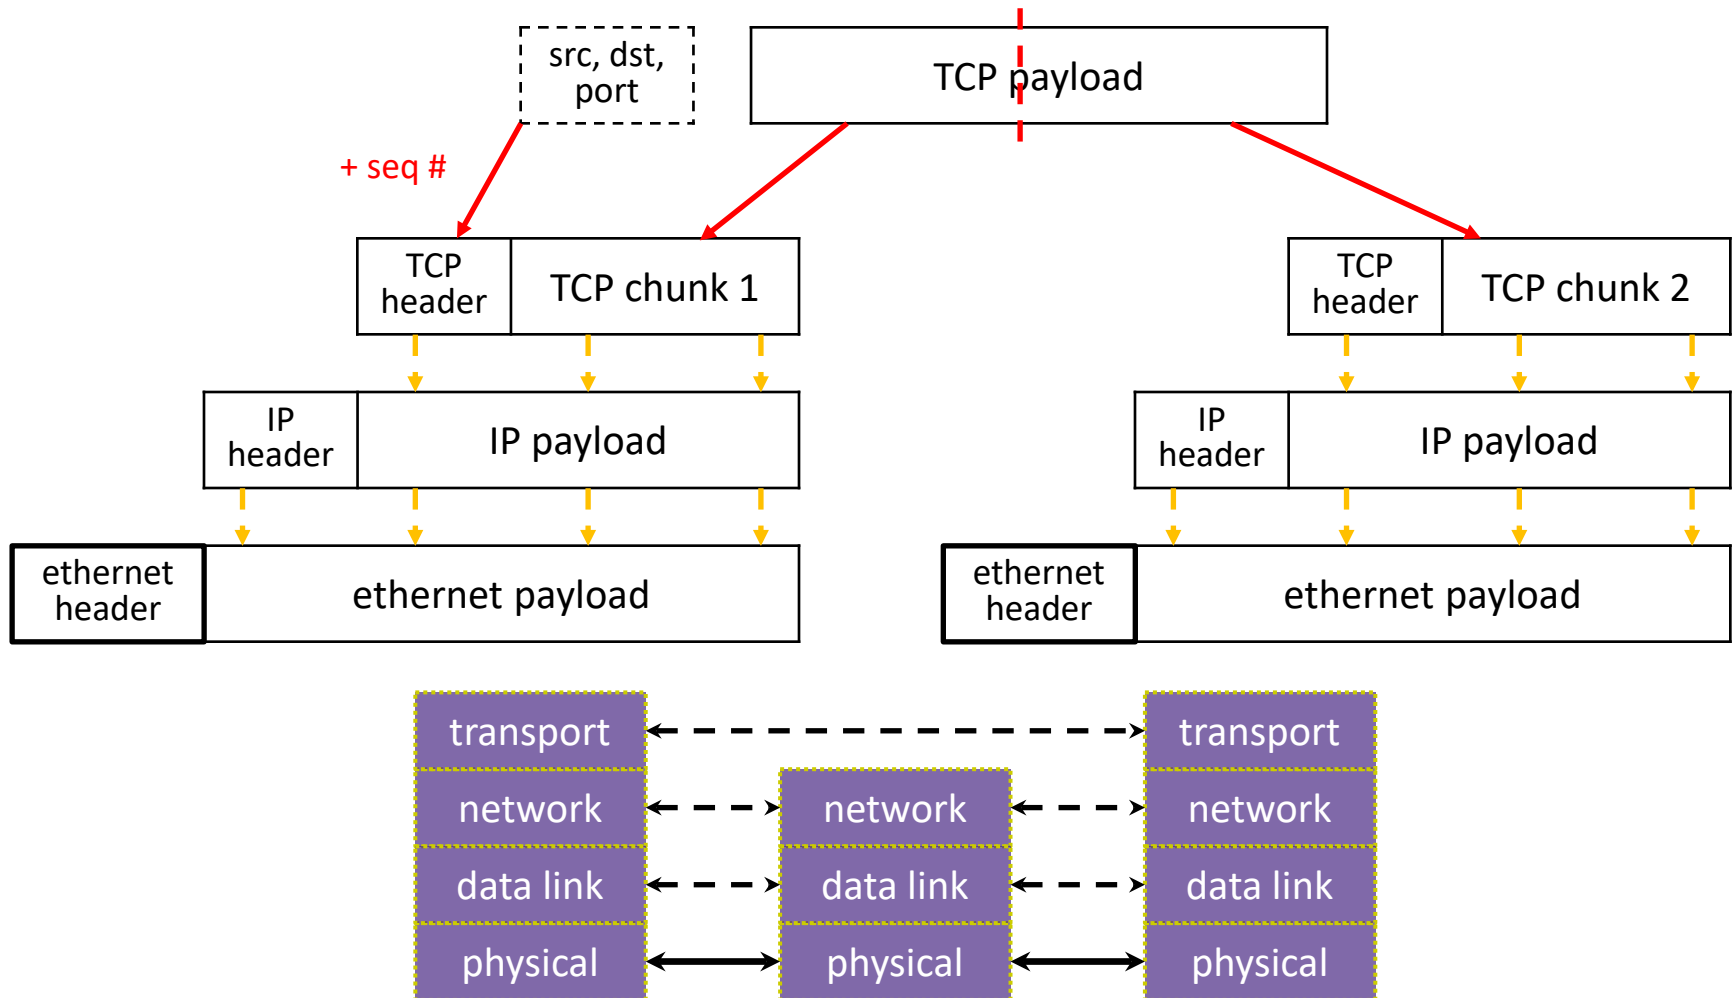
The Transport Layer (TCP)

- Transmission Control Protocol (TCP):
 - ◆ Provides applications with reliable, ordered, congestion-controlled byte streams
 - Sends stream data as multiple IP packets (differentiated by sequence numbers) and retransmits them as necessary
 - When receiving, puts packets back in order and detects missing packets
 - ◆ A single host (IP address) can have up to $2^{16} = 65,535$ “ports”
 - Kind of like an apartment number at a postal address (your applications are the residents who get mail sent to an apt. #)




The Transport Layer (TCP)

- Packet encapsulation – one more nested layer!



The Transport Layer (TCP)


- Applications use OS services to establish TCP streams:
 - ◆ The “Berkeley sockets” API
 - ▣ A set of OS system calls
 - ◆ Clients **connect** () to a server IP address + application port number
 - ◆ Servers **listen** () for and **accept** () client connections
 - ◆ Clients and servers **read** () and **write** () data to each other



The Transport Layer (UDP)

■ User Datagram Protocol (UDP):

- ◆ Provides applications with *unreliable* packet delivery
- ◆ UDP is a really thin, simple layer on top of IP
 - Datagrams still are fragmented into multiple IP packets




The (Mostly Missing) Layers 5 & 6

■ Layer 5: Session Layer

- ◆ Supposedly handles establishing and terminating application sessions
- ◆ Remote Procedure Call (RPC) kind of fits in here

■ Layer 6: Presentation Layer


- ◆ Supposedly maps application-specific data units into a more network-neutral representation
- ◆ Encryption (SSL) kind of fits in here



The Application Layer

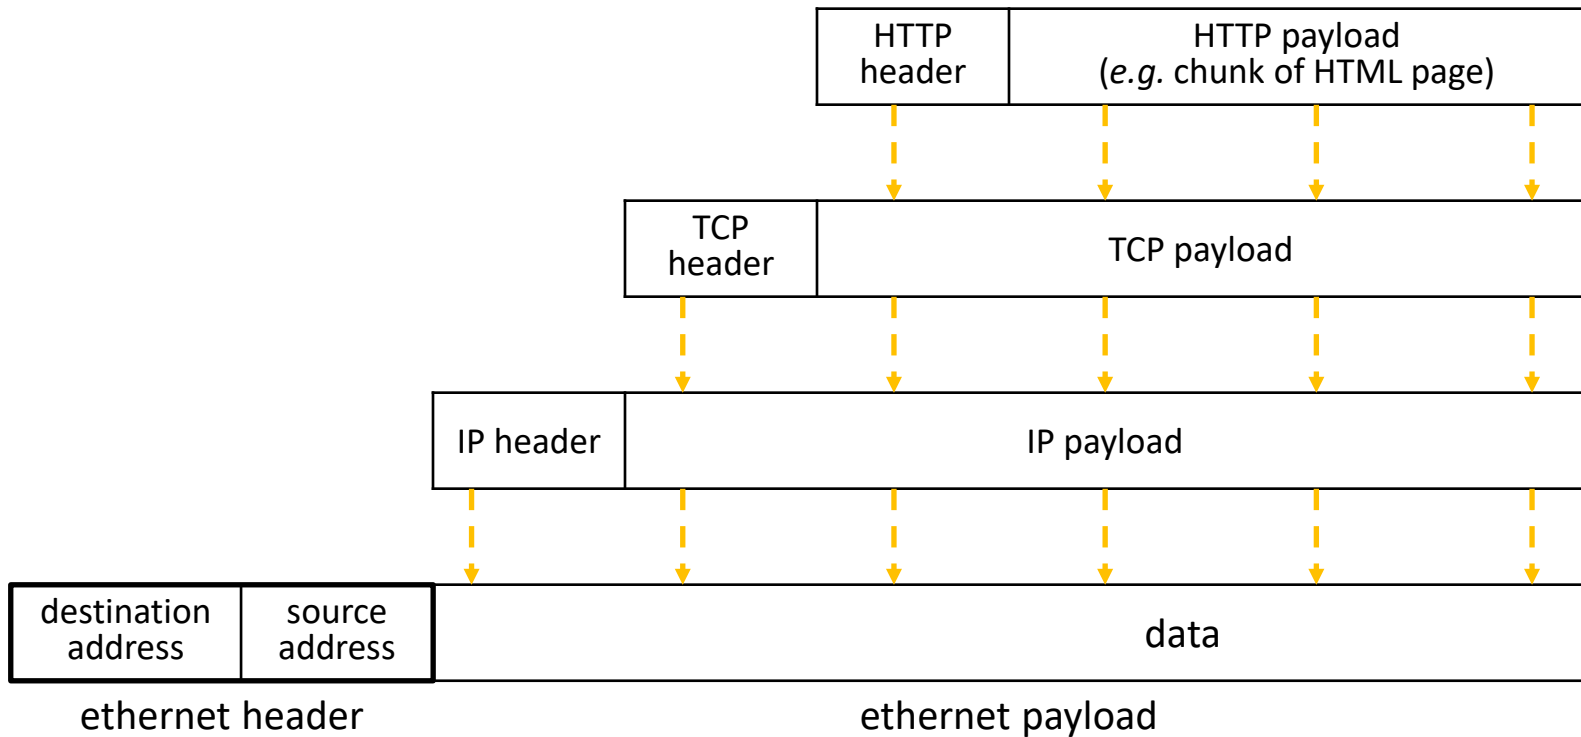
■ Application protocols

- ◆ The format and meaning of messages between application entities
- ◆ Example: HTTP is an application-level protocol that dictates how web browsers and web servers communicate
 - HTTP is implemented *on top of* TCP streams




The Application Layer

■ Packet encapsulation:



The Application Layer

- Packet encapsulation:



The Application Layer

- Popular application-level protocols:
 - ◆ **DNS:** translates a domain name (e.g. www.google.com) into one or more IP addresses (e.g. 74.125.197.106)
 - Domain Name System
 - An hierarchy of DNS servers cooperate to do this
 - ◆ **HTTP:** web protocols
 - Hypertext Transfer Protocol
 - ◆ **SMTP, IMAP, POP:** mail delivery and access protocols
 - Secure Mail Transfer Protocol, Internet Message Access Protocol, Post Office Protocol
 - ◆ **SSH:** secure remote login protocol
 - Secure Shell
 - ◆ **bittorrent:** peer-to-peer, swarming file sharing protocol

网络编程

- 网络基础知识
- Socket, Client/Server

<https://docs.oracle.com/javase/tutorial/networking/index.html>

Files and File Descriptors

- Remember `open()`, `read()`, `write()`, and `close()`?
 - ◆ POSIX system calls for interacting with files
 - ◆ `open()` returns a file descriptor
 - An integer that represents an open file
 - This file descriptor is then passed to `read()`, `write()`, and `close()`
 - ◆ Inside the OS, the file descriptor is used to index into a table that keeps track of any OS-level state associated with the file, such as the file position.

POSIX的故事

<https://unix.org/version3/online.html>

- 可移植操作系统接口Portable Operating System Interface of UNIX → POSIX
- POSIX是IEEE为要在各种UNIX操作系统上运行的软件而定义的一系列API标准的总称，POSIX.1 已经被国际标准化组织（International Standards Organization, ISO）所接受，被命名为 ISO/IEC 9945-1:1990 标准。
- POSIX 标准的制定最后投票敲定阶段大概是 1991~1993 年间，而此时正是 Linux 刚刚起步的时候，这个 UNIX 标准为 Linux 提供了极为重要的信息，使得 Linux 能够在标准的指导下进行开发，并能够与绝大多数 UNIX 操作系统兼容。
- 在最初的 Linux 内核源码（0.01版、0.11版）中就已经为 Linux 系统与 POSIX 标准的兼容做好了准备工作。
- 在 Linux 0.01 版内核 /include/unistd.h 文件中就已经定义了几个有关 POSIX 标准要求的符号常数，而且 Linus 在注释中已写道：“OK，这也许是个玩笑，但我正在着手研究它呢”。
- 正是由于Linux支持POSIX标准，无数可以在unix上运行的程序都陆续的移植到Linux上，而此时unix因为版权问题，官司打的不可开交，使得Linux后来者居上

POSIX历史

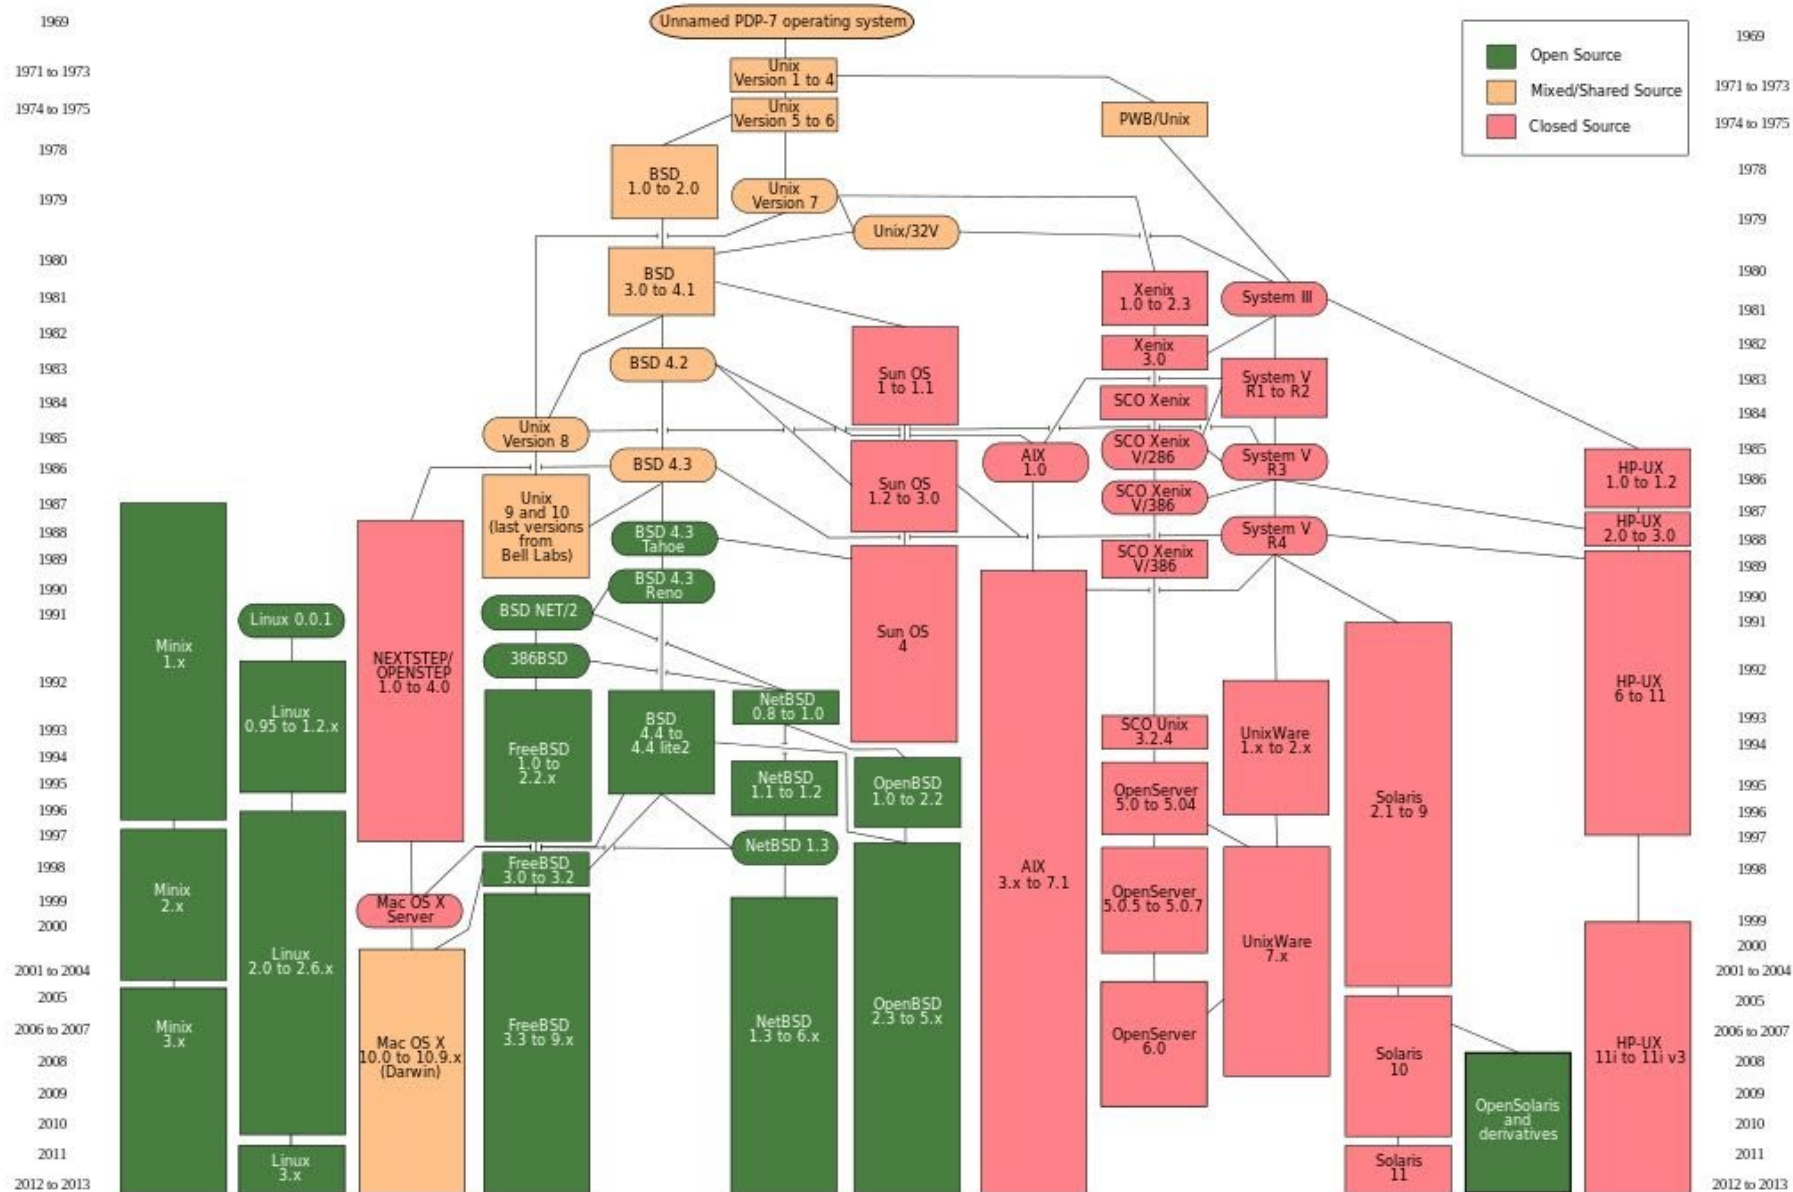
- 1974年，贝尔实验室正式对外发布Unix。因为涉及到反垄断等各种原因，加上早期的Unix不够完善，于是贝尔实验室以慷慨的条件向学校提供源代码，所以Unix在大专院校里获得了很多支持并得以持续发展。
- 于是出现了独立开发的与Unix基本兼容但又不完全兼容的OS→Unix-like OS。
 - ◆ 美国加州大学伯克利分校的Unix4.xBSD(Berkeley Software Distribution)。
 - ◆ 贝尔实验室发布的自己的版本，称为System V Unix。
 - ◆ 其他厂商的版本，比如Sun Microsystems的Solaris系统,则是从这些原始的BSD和System V版本中衍生而来。
 - ◆ 20世纪80年代中期，Unix厂商试图通过加入新的、往往不兼容的特性来使它们的程序与众不同
- 局面非常混乱，麻烦也就随之而来了。
 - ◆ 为了提高兼容性和应用程序的可移植性，阻止这种趋势，IEEE(电气和电子工程师协会)开始努力标准化Unix的开发，后来由 Richard Stallman命名为“Posix”。

谁遵循这个标准呢？

UNIX



POSIX历史

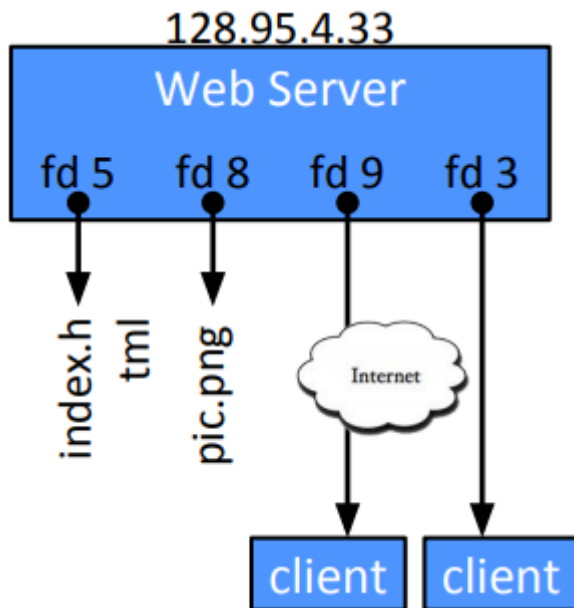


Networks and Sockets

- UNIX likes to make all I/O look like file I/O
 - ◆ You use `read()` and `write()` to communicate with remote computers over the network!
 - ◆ A file descriptor use for network communications is called a **socket**
 - ◆ Just like with files:
 - Your program can have multiple network channels open at once
 - You need to pass a file descriptor to `read()` and `write()` to let the OS know which network channel to use

File Descriptor Table

OS's File Descriptor Table for the Process



File Descriptor	Type	Connection
0	pipe	stdin (console)
1	pipe	stdout (console)
2	pipe	stderr (console)
3	TCP socket	local: 128.95.4.33:80 remote: 44.1.19.32:7113
5	file	index.html
8	file	pic.png
9	TCP socket	local: 128.95.4.33:80 remote: 102.12.3.4:5544

Types of Sockets

- Stream sockets

- ◆ For connection-oriented, point-to-point, reliable byte streams
 - ▣ Using TCP, SCTP, or other stream transports

- Datagram sockets

- ◆ For connection-less, one-to-many, unreliable packets
 - ▣ Using UDP or other packet transports

- Raw sockets

- ◆ For layer-3 communication (raw IP packet manipulation)

Stream Sockets

- Typically used for client-server communications
 - ◆ **Client**: An application that establishes a connection to a server
 - ◆ **Server**: An application that receives connections from clients
 - ◆ Can also be used for other forms of communication like peer-to-peer

1) Establish connection:



2) Communicate:




3) Close connection:




Datagram Sockets

- Often used as a building block
 - ◆ No flow control, ordering, or reliability, so used less frequently
 - ◆ e.g. streaming media applications or DNS lookups

1) Create sockets:



2) Communicate:



The Sockets API

- Berkeley sockets originated in 4.2BSD Unix (1983)
 - ◆ It is the standard API for network programming
 - Available on most OSs
 - ◆ Written in C
- POSIX Socket API
 - ◆ A slight update of the Berkeley sockets API
 - A few functions were deprecated or replaced
 - Better support for multi-threading was added

Socket API: Client TCP Connection

- We'll start by looking at the API from the point of view of a client connecting to a server over TCP
- There are five steps:
 - 1) Figure out the IP address and port to which to connect
 - 2) Create a socket
 - 3) Connect the socket to the remote server
 - 4) `read()` and `write()` data using the socket
 - 5) Close the socket

Java Sockets Programming

- The package `java.net` provides support for sockets programming (and more).
- Typically you import everything defined in this package with:

```
import java.net.*;
```

Classes

InetAddress

Socket

ServerSocket

DatagramSocket

DatagramPacket

InetAddress class

- static methods you can use to create new InetAddress objects.
 - ◆ `getByName(String host)`
 - ◆ `getAllByName(String host)`
 - ◆ `getLocalHost()`

```
InetAddress x = InetAddress.getByName(  
                                "www.tongji.edu.cn");
```

❖ Throws `UnknownHostException`

Sample Code: Lookup.java

- Uses InetAddress class to lookup hostnames found on command line.

```
> java Lookup www.tongji.edu.cn  
www.tongji.edu.cn:192.168.66.4
```

```
try {  
  
    InetAddress a = InetAddress.getByName(hostname);  
  
    System.out.println(hostname + ":" +  
                        a.getHostAddress());  
  
} catch (UnknownHostException e) {  
  
    System.out.println("No address found for " +  
                      hostname);  
  
}
```

Socket class

- Corresponds to active TCP sockets only!
 - ◆ client sockets
 - ◆ socket returned by `accept()`;
- Passive sockets are supported by a different class:
 - ◆ `ServerSocket`
- UDP sockets are supported by
 - ◆ `DatagramSocket`

JAVA TCP Sockets


■ `java.net.Socket`

- ◆ Implements **client sockets** (also called just “sockets”).
- ◆ An endpoint for communication between two machines.
- ◆ Constructor and Methods
 - **Socket** (String host, int port): Creates a stream socket and connects it to the specified port number on the named host.
 - InputStream **getInputStream()**
 - OutputStream **getOutputStream()**
 - `close()`

■ `java.net.ServerSocket`

- ◆ Implements **server sockets**.
- ◆ Waits for requests to come in over the network.
- ◆ Performs some operation based on the request.
- ◆ Constructor and Methods
 - **ServerSocket** (int port)
 - Socket **Accept()**: Listens for a connection to be made to this socket and accepts it. This method blocks until a connection is made.

Sockets



Client socket, welcoming socket (passive) and connection socket (active)

Socket Constructors

- Constructor creates a TCP connection to a named TCP server.
 - ◆ There are a number of constructors:

```
Socket(InetAddress server, int port);
```

```
Socket(InetAddress server, int port,  
        InetAddress local, int localport);
```

```
Socket(String hostname, int port);
```

Socket Methods

```
void close();
```

```
InetAddress getAddress();
```

```
InetAddress getLocalAddress();
```

```
InputStream getInputStream();
```

```
OutputStream getOutputStream();
```

- Lots more (setting/getting socket options, partial close, etc.)

Socket I/O

- Socket I/O is based on the Java I/O support
 - ◆ in the package `java.io`
- `InputStream` and `OutputStream` are abstract classes
 - ◆ common operations defined for all kinds of `InputStreams`, `OutputStreams`...

InputStream Basics

// reads some number of bytes and

// puts in buffer array b

```
int read(byte[] b);
```

// reads up to len bytes

```
int read(byte[] b, int off, int len);
```

Both methods can throw IOException.

Both return -1 on EOF.

OutputStream Basics

// writes b.length bytes

```
void write(byte[] b);
```

// writes len bytes starting

// at offset off

```
void write(byte[] b, int off, int len);
```

Both methods can throw IOException.

ServerSocket Class (TCP Passive Socket)

- Constructors:

`ServerSocket` (int port);

`ServerSocket` (int port, int backlog);

`ServerSocket` (int port, int backlog,
 InetAddress bindAddr);

ServerSocket Methods

Socket **accept()**;

void **close()**;

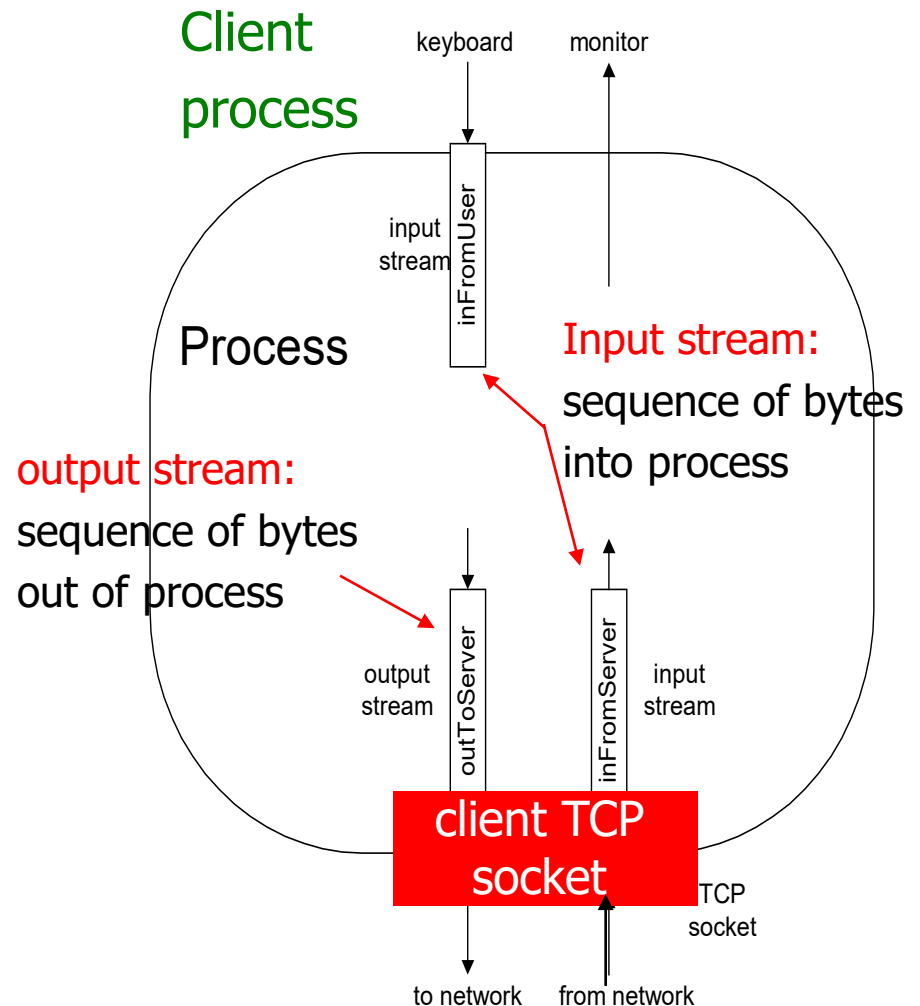
InetAddress **getInetAddress()**;

int **getLocalPort()**;

throw IOException, SecurityException

Socket programming with TCP

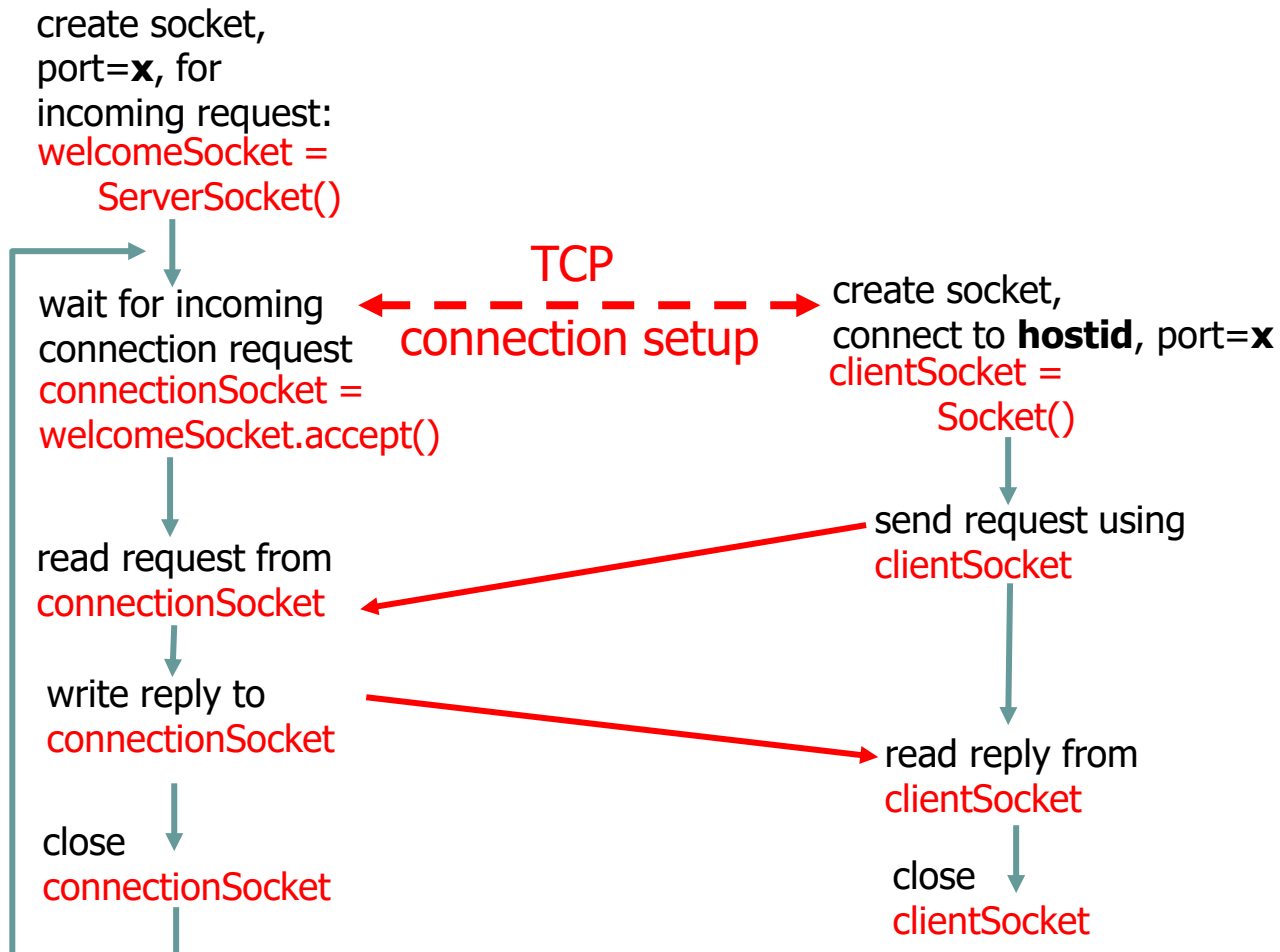
- Example client-server app:
- client reads line from standard input (**inFromUser stream**) , sends to server via socket (**outToServer stream**)
- server reads line from socket
- server converts line to uppercase, sends back to client
- client reads, prints modified line from socket (**inFromServer stream**)



Client/server socket interaction: TCP

Server (running on **hostid**)

Client



TCPClient.java

```
import java.io.*;  
import java.net.*;
```

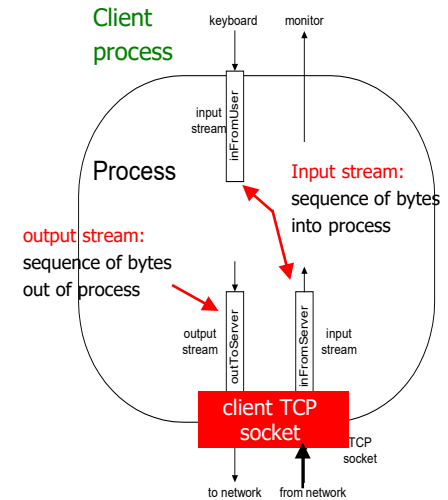
```
class TCPClient {  
    public static void main(String argv[]) throws Exception  
    {
```

```
        String sentence;  
        String modifiedSentence;
```

```
        BufferedReader inFromUser =  
            new BufferedReader(new InputStreamReader(System.in));
```

```
        Socket clientSocket = new Socket("hostname", 6789);
```

```
        DataOutputStream outToServer =  
            new DataOutputStream(clientSocket.getOutputStream());
```



TCPClient.java

```
BufferedReader inFromServer =  
    new BufferedReader(new  
        InputStreamReader(clientSocket.getInputStream()));
```

```
sentence = inFromUser.readLine();
```

```
outToServer.writeBytes(sentence + '\n');
```

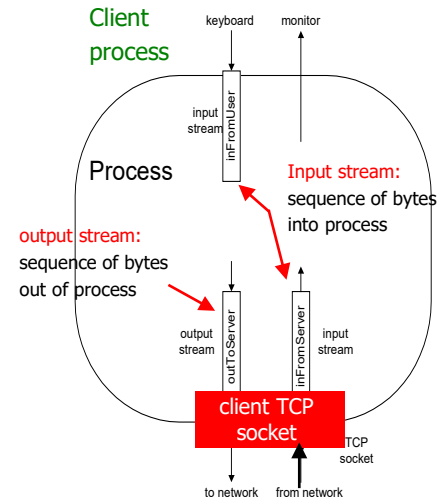
```
modifiedSentence = inFromServer.readLine();
```

```
System.out.println("FROM SERVER: " + modifiedSentence);
```

```
clientSocket.close();
```

```
}
```

```
}
```



TCPServer.java

```
import java.io.*;
import java.net.*;
class TCPServer {
    public static void main(String argv[]) throws Exception
```

```
{
    String clientSentence;
    String capitalizedSentence;
```

```
    ServerSocket welcomeSocket = new ServerSocket(6789);
```

```
    while(true) {
```

```
        Socket connectionSocket = welcomeSocket.accept();
```

```
        BufferedReader inFromClient = new BufferedReader(new
            InputStreamReader(connectionSocket.getInputStream()));
```

```
        DataOutputStream outToClient =
            new DataOutputStream(connectionSocket.getOutputStream());
```

```
        clientSentence = inFromClient.readLine();
```

```
        capitalizedSentence = clientSentence.toUpperCase() + '\n';
```

```
        outToClient.writeBytes(capitalizedSentence);
```

```
    }
```

```
}
```

```
}
```

Server (running on **hostid**)

create socket,
port=**x**, for
incoming request:
welcomeSocket =
ServerSocket()

wait for incoming
connection request
connectionSocket =
welcomeSocket.accept()

read request from
connectionSocket

write reply to
connectionSocket

close
connectionSocket

← TCP
connection setup →

Client

create socket,
connect to **hostid**, port=**x**
clientSocket =
Socket()

send request using
clientSocket

read reply from
clientSocket

close
clientSocket



Sample Echo Server

TCPEchoServer.java

Simple TCP Echo server.

Based on code from:

TCP/IP Sockets in Java

<https://github.com/Jonikiro/TCPEchoApp>

TCPEchoServer.java

```
public class TCPEchoServer {
    public static void main(String[] args) {
        try {
            /* Create ServerSocket, assign it port 12900, give it a
             *queue of 100, and bind it to IP address of localhost. */
            ServerSocket serverSocket = new ServerSocket(12900, 100,
                InetAddress.getByName("localhost"));
            System.out.println("Server started at: " + serverSocket);

            // Keep accepting clients in infinite loop
            while (true) {
                System.out.println("Waiting for a connection...");

                // Accept a connection and assign it to normal Socket
                final Socket activeSocket = serverSocket.accept();

                System.out.println("Received a connection from " +
                    activeSocket);

                // Create a new thread to handle the new connection
                Runnable runnable =
                    () -> handleClientRequest(activeSocket);
                new Thread(runnable).start();
            }
        } catch (IOException ex) {ex.printStackTrace();}
    }
}
```

```
public static void handleClientRequest(Socket socket) {
    BufferedReader socketReader = null;
    BufferedWriter socketWriter = null;

    try {
        // Create a buffered reader/writer for the socket
        socketReader = new BufferedReader(
            new InputStreamReader(socket.getInputStream()));
        socketWriter = new BufferedWriter(
            new OutputStreamWriter(socket.getOutputStream()));

        String inMsg = null;
        while ((inMsg = socketReader.readLine()) != null) {
            System.out.println("Received from client: " + inMsg);

            // Echo the received message to the client
            String outMsg = inMsg;
            socketWriter.write(outMsg + "\n");
            socketWriter.flush();
        }
    } catch (IOException ex) {ex.printStackTrace();}
    finally {
        try {
            socket.close();
        } catch (IOException ex) {ex.printStackTrace();}
    }
}
```

TCPEchoClient.java

```
public class TCPEchoClient {
    public static void main(String[] args) {
        Socket socket = null;
        BufferedReader socketReader = null;
        BufferedWriter socketWriter = null;

        try {
            /* Create a socket that will connect to localhost
             * at port 12900. Note that server must also be
             * running at localhost and 12900. */
            socket = new Socket("localhost", 12900);
            System.out.println("Started client socket at " +
                               socket.getLocalSocketAddress());

            // Create buffered reader/writer using IO streams
            socketReader = new BufferedReader(
                new InputStreamReader(socket.getInputStream()));
            socketWriter = new BufferedWriter(
                new OutputStreamWriter(socket.getOutputStream()));

            // Create buffered reader for user's input
            BufferedReader consoleReader =
                new BufferedReader(new InputStreamReader(System.in));
```

```
String promptMsg = "Please enter a message (Bye to quit):";
String outMsg = null;

System.out.print(promptMsg);
while ((outMsg = consoleReader.readLine()) != null) {
    if (outMsg.equalsIgnoreCase("bye")) {
        break;
    }

    socketWriter.write(outMsg + "\n");
    socketWriter.flush();

    String inMsg = socketReader.readLine();
    System.out.println("Server: " + inMsg);

    System.out.println();
    System.out.print(promptMsg);
} catch (IOException ex) {ex.printStackTrace();}
finally {
    if (socket != null) {
        try {
            socket.close();
        } catch (IOException ex) {ex.printStackTrace();}
    }
}
}
```

UDP Sockets

- DatagramSocket class
- DatagramPacket class needed to specify the payload
 - ◆ incoming or outgoing

Socket Programming with UDP

■ UDP

- ◆ Connectionless and unreliable service.
- ◆ There isn't an initial handshaking phase.
- ◆ Doesn't have a pipe.
- ◆ transmitted data may be received out of order, or lost

■ Socket Programming with UDP

- ◆ No need for a welcoming socket.
- ◆ No streams are attached to the sockets.
- ◆ the sending hosts creates “packets” by attaching the IP destination address and port number to each batch of bytes.
- ◆ The receiving process must unravel to received packet to obtain the packet's information bytes.

JAVA UDP Sockets

■ In Package java.net

◆ java.net.DatagramSocket

- A socket for sending and receiving datagram packets.
- Constructor and Methods
 - **DatagramSocket** (int port): Constructs a datagram socket and binds it to the specified port on the local host machine.
 - void **receive** (DatagramPacket p)
 - void **send** (DatagramPacket p)
 - void **close** ()

DatagramSocket Constructors

- `DatagramSocket ()`;
- `DatagramSocket (int port)`;
- `DatagramSocket (int port, InetAddress a)`;
- All can throw `SocketException` or `SecurityException`

Datagram Methods

```
void connect(InetAddress, int port);
```

```
void close();
```

```
void receive(DatagramPacket p);
```

```
void send(DatagramPacket p);
```

Lots more!

Datagram Packet

- Contain the payload
 - ◆ a byte array
- Can also be used to specify the destination address
 - ◆ when not using connected mode UDP

DatagramPacket Constructors

For **receiving**:

```
DatagramPacket( byte[] buf, int len);
```

For **sending**:

```
DatagramPacket( byte[] buf, int len  
                InetAddress a, int port);
```

DatagramPacket methods

```
byte[] getData();
```

```
void setData(byte[] buf);
```

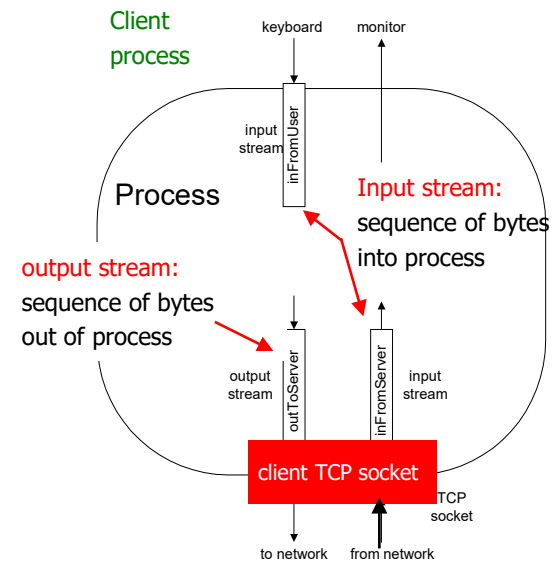
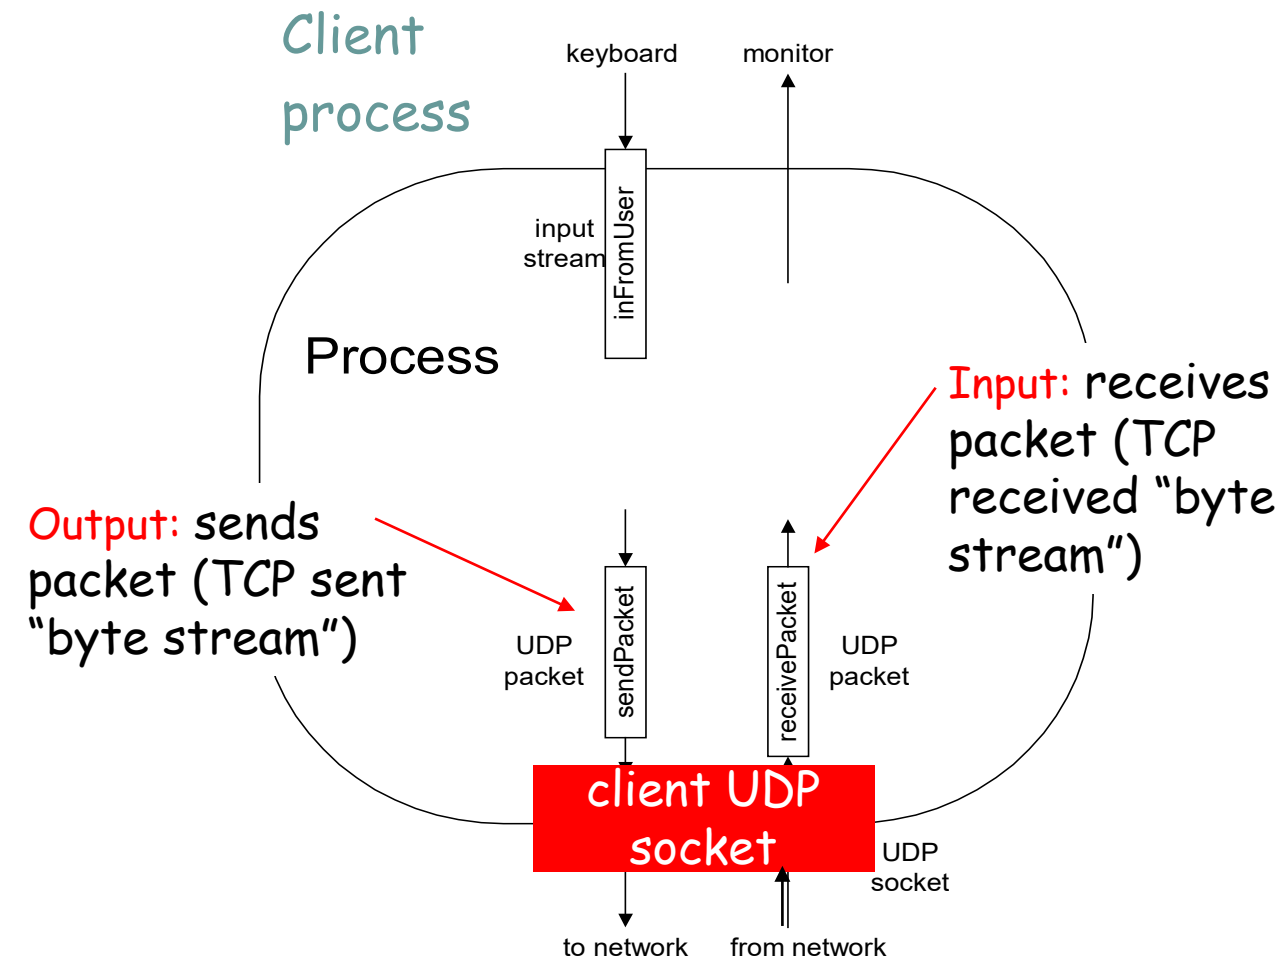
```
void setAddress(InetAddress a);
```

```
void setPort(int port);
```

```
InetAddress getAddress();
```

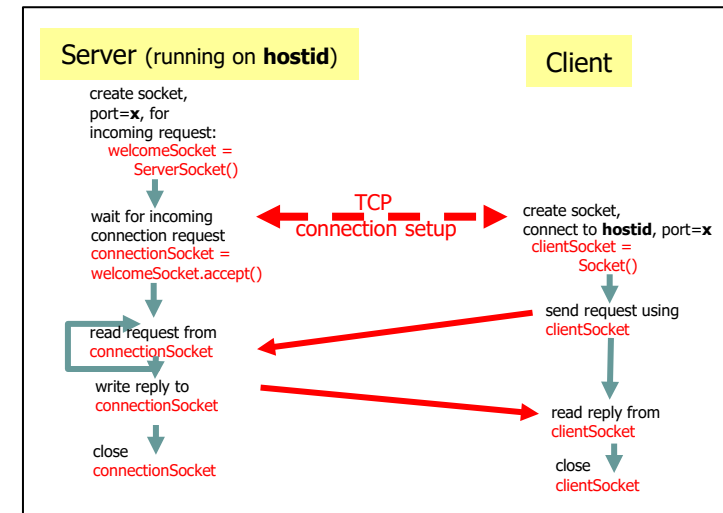
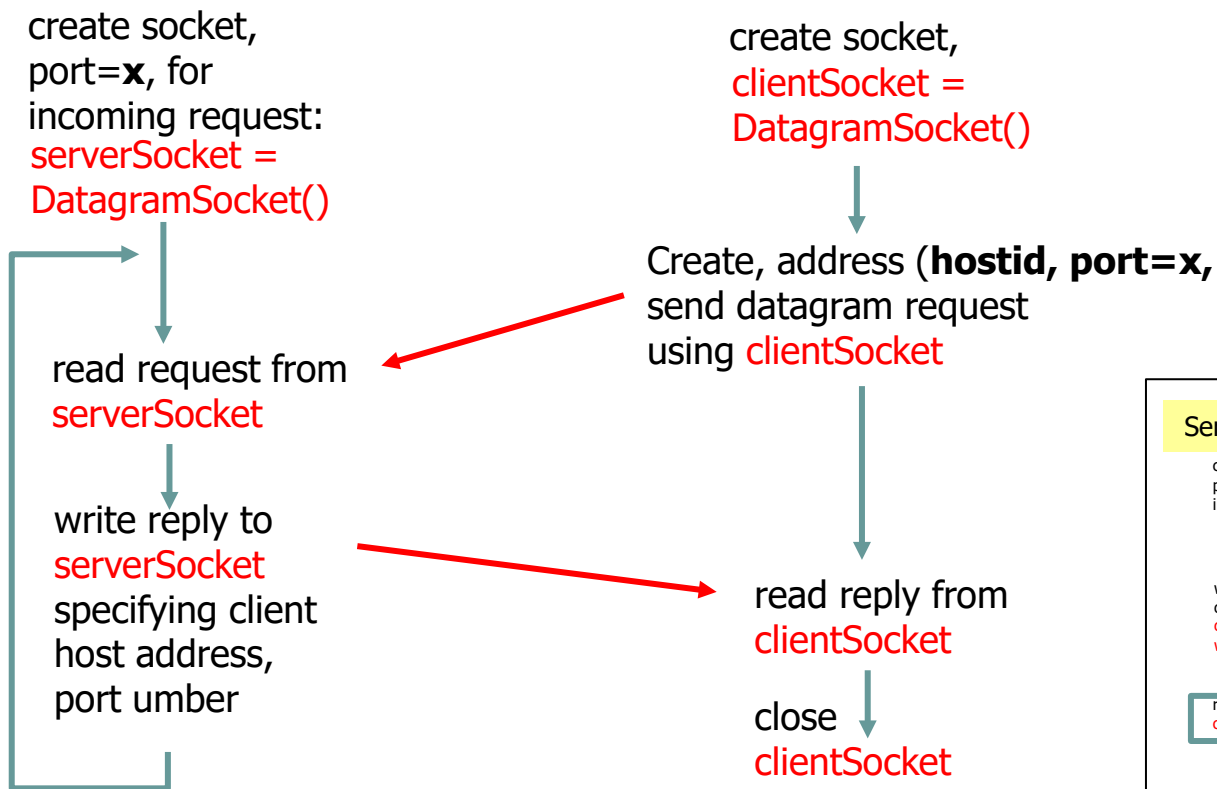
```
int getPort();
```

Example: Java client (UDP)



Client/server socket interaction: UDP

Server (running on **hostid**) Client



UDPClient.java

```
import java.io.*;
import java.net.*;
```

```
class UDPClient {
    public static void main(String args[]) throws Exception
    {
```

```
BufferedReader inFromUser =  
    new BufferedReader(new InputStreamReader(System.in));
```

```
DatagramSocket clientSocket = new DatagramSocket();
```

```
InetAddress IPAddress = InetAddress.getByName("hostname");
```

```
byte[] sendData = new byte[1024];
byte[] receiveData = new byte[1024];
String sentence = inFromUser.readLine();
sendData = sentence.getBytes();
```

```
DatagramPacket sendPacket = new DatagramPacket(sendData, sendData.length, IPAddress, 9876);
```

```
clientSocket.send(sendPacket);
```

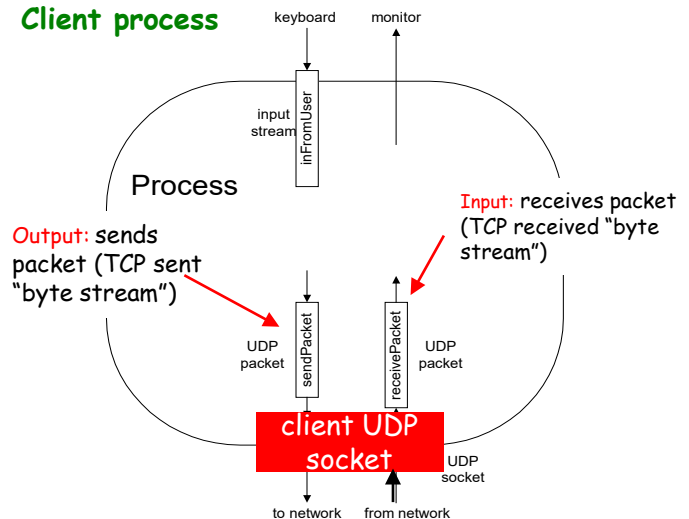
```
DatagramPacket receivePacket = new DatagramPacket(receiveData, receiveData.length);
```

```
clientSocket.receive(receivePacket);
```

```
String modifiedSentence = new String(receivePacket.getData());
```

```
System.out.println("FROM SERVER:" + modifiedSentence);
```

```
clientSocket.close();
```

 $\left. \begin{array}{l} \\ \end{array} \right\}$ 

UDPServer.java

```
import java.io.*;
import java.net.*;
```

```
class UDPServer {
```

```
    public static void main(String args[]) throws Exception
    {
```

```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

```
        byte[] receiveData = new byte[1024];
        byte[] sendData = new byte[1024];
```

```
        while(true)
        {
```

```
            DatagramPacket receivePacket = new DatagramPacket(receiveData, receiveData.length);
```

```
            serverSocket.receive(receivePacket);
```

```
            String sentence = new String(receivePacket.getData());
```

```
            String capitalizedSentence = sentence.toUpperCase();
```

```
            sendData = capitalizedSentence.getBytes();
```

```
            InetAddress IPAddress = receivePacket.getAddress();
```

```
            int port = receivePacket.getPort();
```

```
            DatagramPacket sendPacket = new DatagramPacket(sendData, sendData.length, IPAddress, port);
```

```
            serverSocket.send(sendPacket);
```

```
        }
    }
}
```

Server (running on **hostid**) Client

create socket,
port=**x**, for
incoming request:
serverSocket =
DatagramSocket()

read request from
serverSocket

write reply to
serverSocket
specifying client
host address,
port number

create socket,
clientSocket =
DatagramSocket()

Create, address (**hostid**, **port=x**),
send datagram request
using **clientSocket**

read reply from
clientSocket

close
clientSocket



Sample UDP code

UDPEchoServer.java

Simple UDP Echo server.

Test using nc as the client (netcat):

```
> nc -u hostname port
```

Socket functional calls

socket (): Create a socket

bind(): bind a socket to a local IP address and port #

listen(): passively waiting for connections

connect(): initiating connection to another socket

accept(): accept a new connection

Write(): write data to a socket

Read(): read data from a socket

sendto(): send a datagram to another UDP socket

recvfrom(): read a datagram from a UDP socket

close(): close a socket (tear down the connection)

Java URL Class

- Represents a Uniform Resource Locator
 - ◆ scheme (protocol)
 - ◆ hostname
 - ◆ port
 - ◆ path
 - ◆ query string

Parsing

- You can use a URL object as a *parser*:

```
URL u = new URL("http://www.tongji.edu.cn");
```

```
System.out.println("Proto:" + u.getProtocol());
```

```
System.out.println("File:" + u.getFile());
```

URL construction

- You can also build a URL by setting each part individually:

```
URL u = new URL("http",  
                "www.tongji.edu.cn");
```

```
System.out.println("URL:" + u.toExternalForm());
```

```
System.out.println("URL: " + u);
```

Retrieving URL contents

- URL objects can retrieve the documents they refer to!
 - ◆ actually this depends on the protocol part of the URL.
 - ◆ HTTP is supported
 - ◆ File is supported (“file:///c:/foo.html”)
 - ◆ You can get “Protocol Handlers” for other protocols.
- There are a number of ways to do this:

`Object getContent() ;`

`InputStream openStream() ;`

`URLConnection openConnection() ;`

Getting Header Information

- There are methods that return information extracted from response headers:

```
String getContentType();
```

```
String getContentLength();
```

```
long getLastModified();
```

URLConnection

- Represents the connection (not the URL itself).
- More control than URL
 - ◆ can write to the connection (send POST data).
 - ◆ can set request headers.
- Closely tied to HTTP