

分布式系统

07-分布式系统的一致性与复制 Consistency and Replication in DS

Weixiong Rao 饶卫雄
Tongji University 同济大学软件学院
2023 秋季
wxrao@tongji.edu.cn

Today...

- Last Session
 - Synchronization: Mutual Exclusion and Election Algorithms
- Today's session
 - Consistency and Replication
 - Introduction
 - Data-centric and Client-Centric Consistency Models



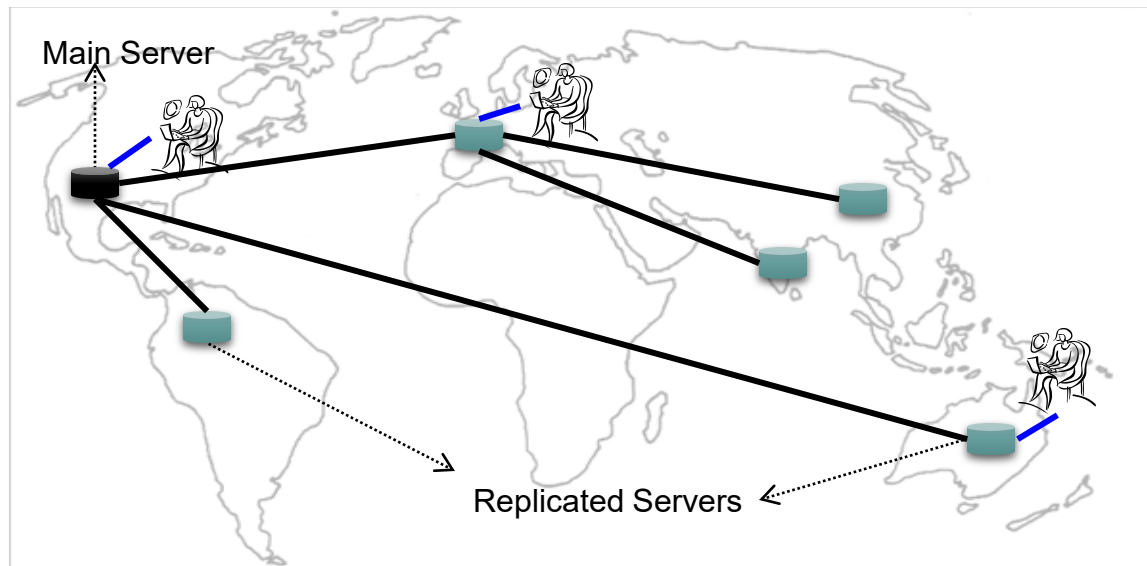
Why Replication?

- Replication is the process of maintaining the data at multiple computers
- Replication is necessary for:
 1. Improving performance
 - A client can access the replicated copy of the data that is near to its location
 2. Increasing the availability of services
 - Replication can mask failures such as server crashes and network disconnection
 3. Enhancing the scalability of the system
 - Requests to the data can be distributed to many servers which contain replicated copies of the data
 4. Securing against malicious attacks
 - Even if some replicas are malicious, secure data can be guaranteed to the client by relying on the replicated copies at the non-compromised servers

1. Replication for Improving Performance

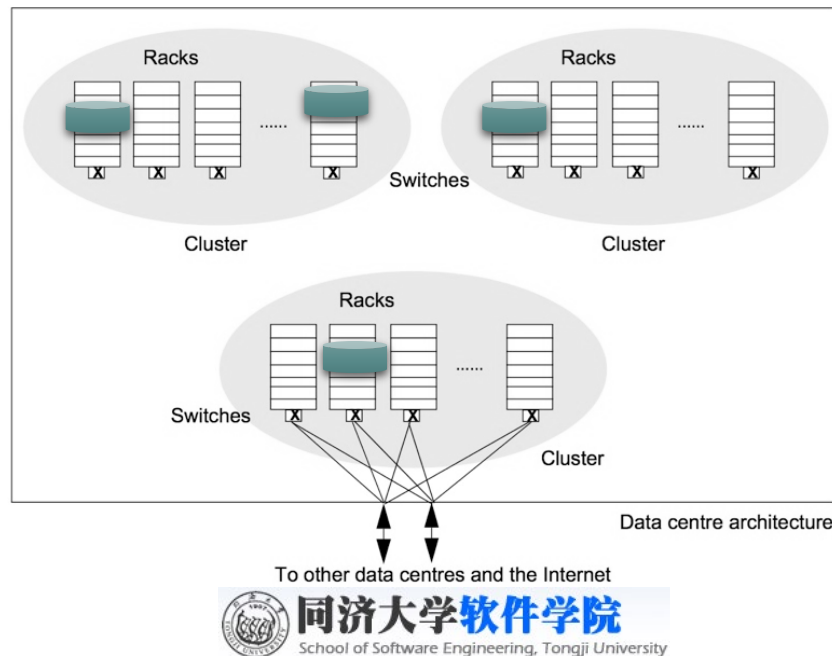
■ Example Applications

- ◆ Caching webpages at the client browser
- ◆ Caching IP addresses at clients and DNS Name Servers
- ◆ Caching in Content Delivery Network (CDNs)
 - Commonly accessed contents, such as software and streaming media, are cached at various network locations



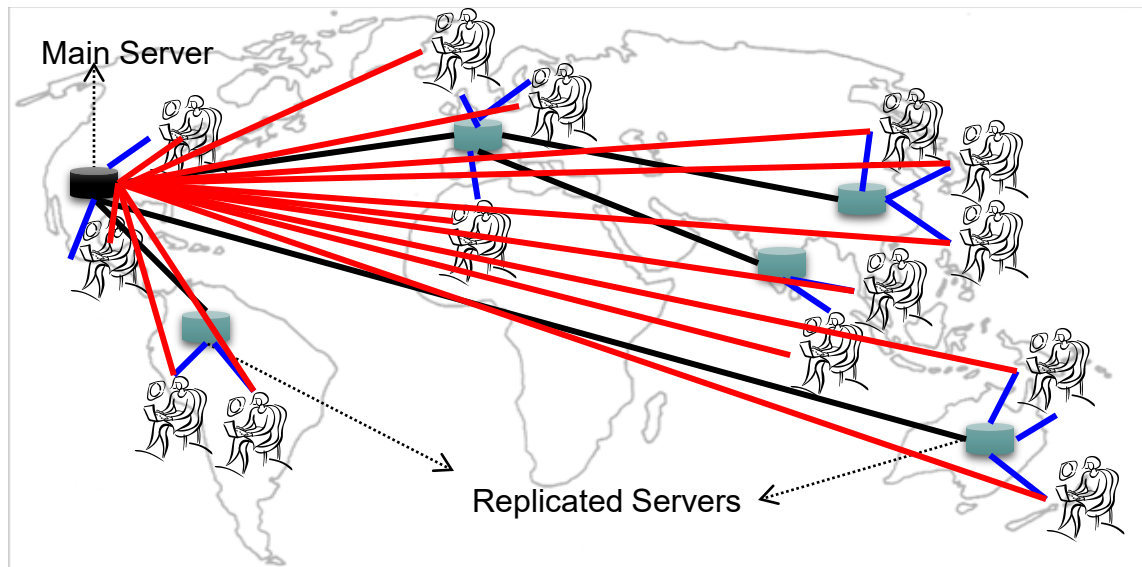
2. Replication for High-Availability

- Availability can be increased by storing the data at replicated locations (instead of storing one copy of the data at a server)
- Example: Google File-System replicates the data at computers across different racks, clusters and data-centers
 - ◆ If one computer or a rack or a cluster crashes, then the data can still be accessed from another source



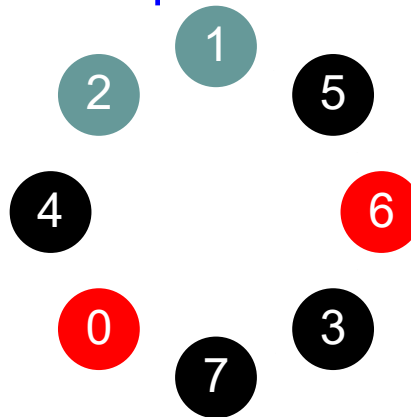
3. Replication for Enhancing Scalability

- Distributing the data across replicated servers helps in avoiding bottle-necks at the main server
 - ◆ It balances the load between the main and the replicated servers
- Example: Content Delivery Networks decrease the load on main servers of the website



4. Replication for Securing Against Malicious Attacks

- If a minority of the servers that hold the data are malicious, the non-malicious servers can outvote the malicious servers, thus providing security
- The technique can also be used to provide fault-tolerance against non-malicious but faulty servers
- Example: In a peer-to-peer system, peers can coordinate to prevent delivering faulty data to the requester



Number of servers with correct data outvote the faulty servers



= Servers that do not have the requested data



= Servers with correct data

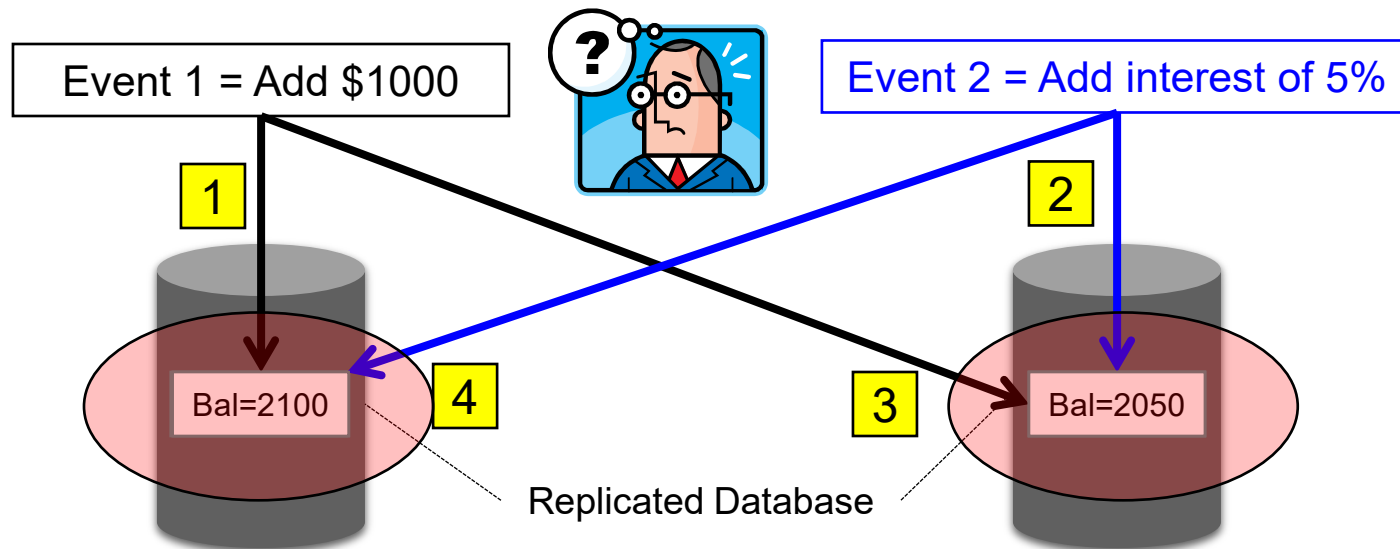


= Servers with faulty data



Why Consistency?

- In a DS with replicated data, one of the main problems is keeping the data consistent
- An example:
 - ◆ In an e-commerce application, the bank database has been replicated across two servers
 - ◆ Maintaining consistency of replicated data is a challenge



Overview of Consistency and Replication

Today's lecture Consistency Models

- ◆ Data-Centric Consistency Models
- ◆ Client-Centric Consistency Models

■ Replica Management

- ◆ When, where and by whom replicas should be placed?
- ◆ Which consistency model to use for keeping replicas consistent?

■ Consistency Protocols

- ◆ We study various implementations of consistency models

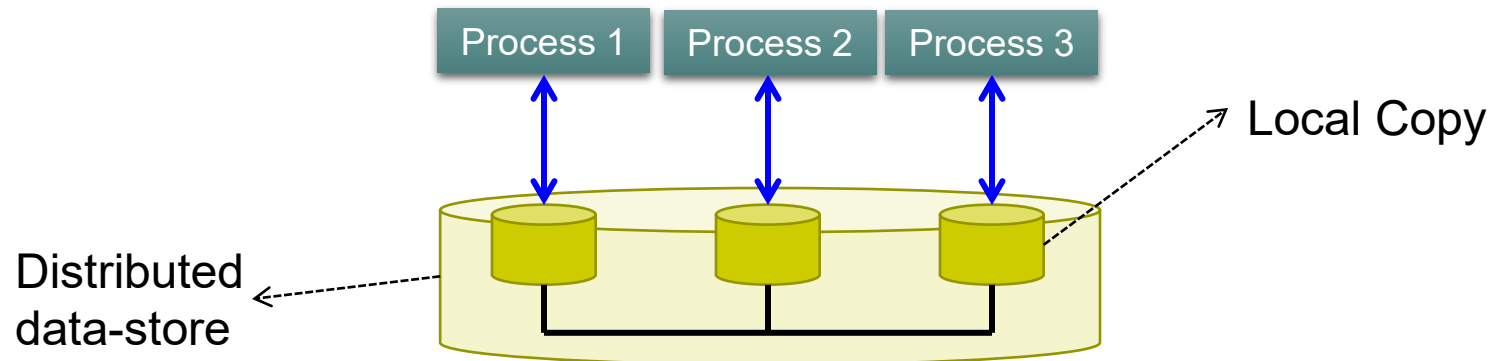
Next lectures

Overview

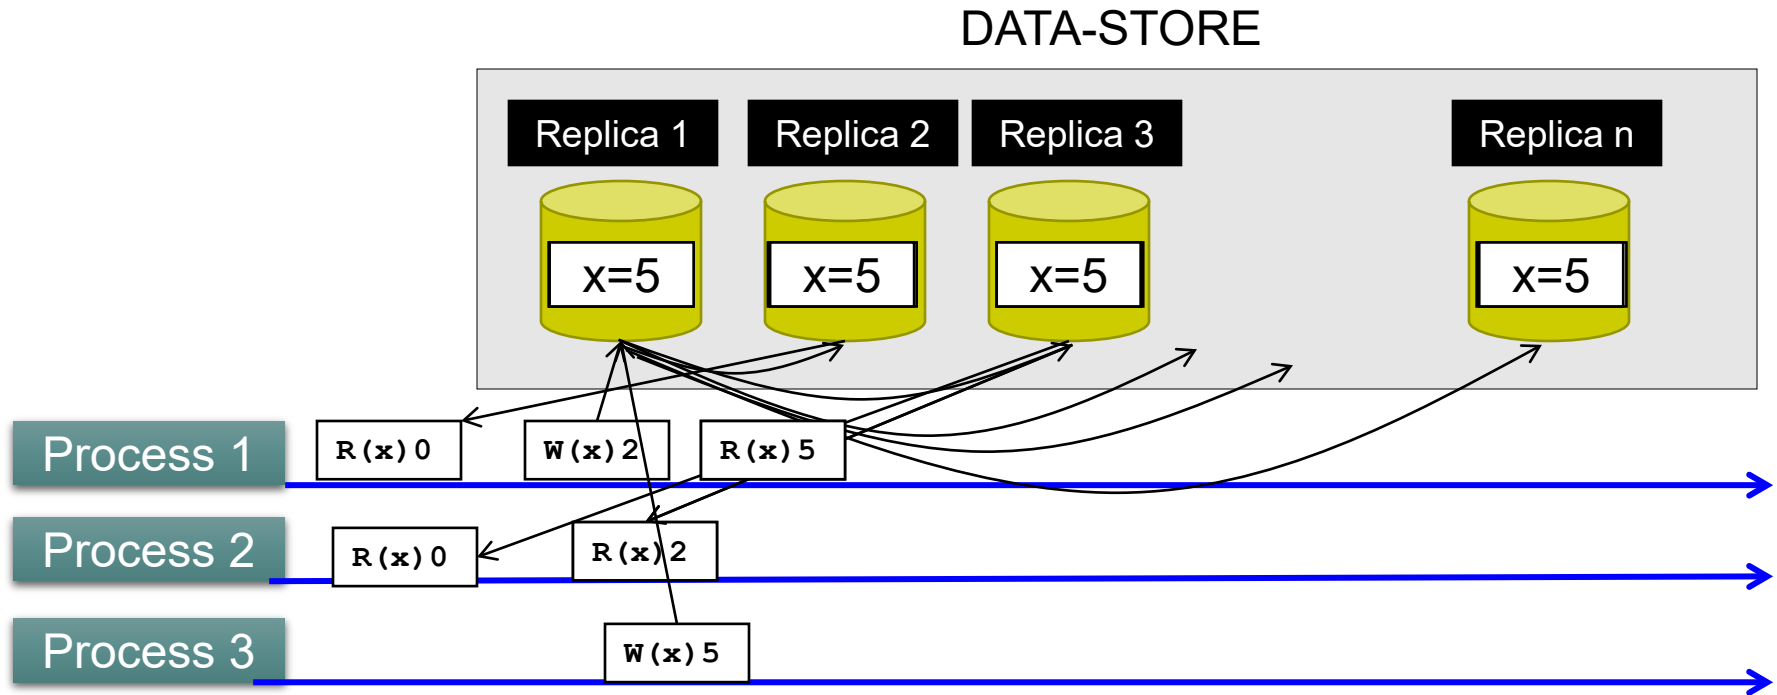
- Consistency Models
 - ◆ Data-Centric Consistency Models
 - ◆ Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

Introduction to Consistency and Replication

- In a distributed system, shared data is typically stored in distributed shared memory, distributed databases or distributed file systems.
 - ◆ The storage can be distributed across multiple computers
 - ◆ Simply, we refer to a series of such data storage units as *data-stores*
- Multiple processes can access shared data by accessing any replica on the data-store
 - ◆ Processes generally perform read and write operations on the replicas



Maintaining Consistency of Replicated Data



Strict Consistency

- Data is always fresh
 - After a write operation, the update is propagated to all the replicas
 - A read operation will result in reading the most recent write
- If there are occasional writes and reads, this leads to large overheads

P1 = Process P1

\longrightarrow = Timeline at P1

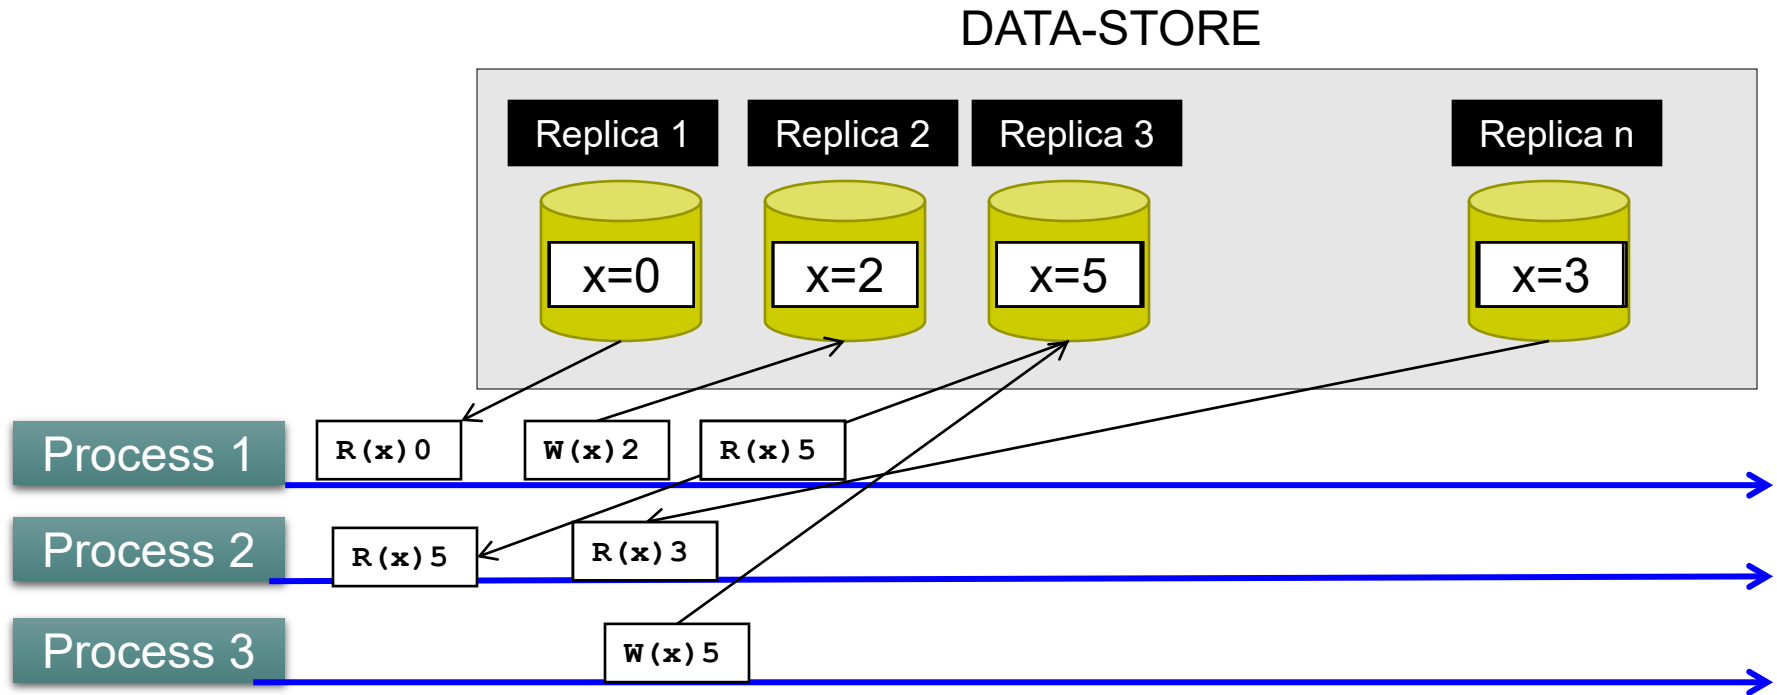
$R(x) b$

= Read variable x;
Result is b

$W(x) b$

= Write variable x;
Result is b

Maintaining Consistency of Replicated Data (cont'd)



Loose Consistency

- Data might be stale
 - A read operation may result in reading a value that was written long back
 - Replicas are generally out-of-sync
- The replicas may sync at coarse grained time, thus reducing the overhead

P1 = Process P1

→ = Timeline at P1

R(x)b

= Read variable x;
Result is b

W(x)b

= Write variable x;
Result is b

Trade-offs in Maintaining Consistency

- Maintaining consistency should balance between the strictness of consistency versus efficiency
 - ◆ Good-enough consistency depends on your application



Consistency Model

- A consistency model is a contract between
 - ◆ the process that wants to use the data, and
 - ◆ the replicated data repository (or data-store)
- A consistency model states the level of consistency provided by the *data-store* to the processes while reading and writing the data

Types of Consistency Models

- Consistency models can be divided into two types:
 - ◆ Data-Centric Consistency Models
 - These models define how the data updates are propagated across the replicas to keep them consistent
 - ◆ Client-Centric Consistency Models
 - These models assume that clients connect to different replicas at different times
 - The models ensure that whenever a client connects to a replica, the replica is brought up to date with the replica that the client accessed previously

Overview

- Consistency Models
 - ◆ Data-Centric Consistency Models
 - ◆ Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

Data-centric Consistency Models

- Data-centric Consistency Models describe how the replicated data is kept consistent, and what the processes can expect
- Under Data-centric Consistency Models, we study two types of models:
 - ◆ Consistency Specification Models:
 - These models enable specifying the consistency levels that are tolerable to the application
 - ◆ Models for Consistent Ordering of Operations:
 - These models specify the order in which the data updates are propagated to different replicas

Overview

- Consistency Models
 - ◆ Data-Centric Consistency Models
 - Consistency Specification Models
 - Models for Consistent Ordering of Operations
 - ◆ Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

Consistency Specification Models

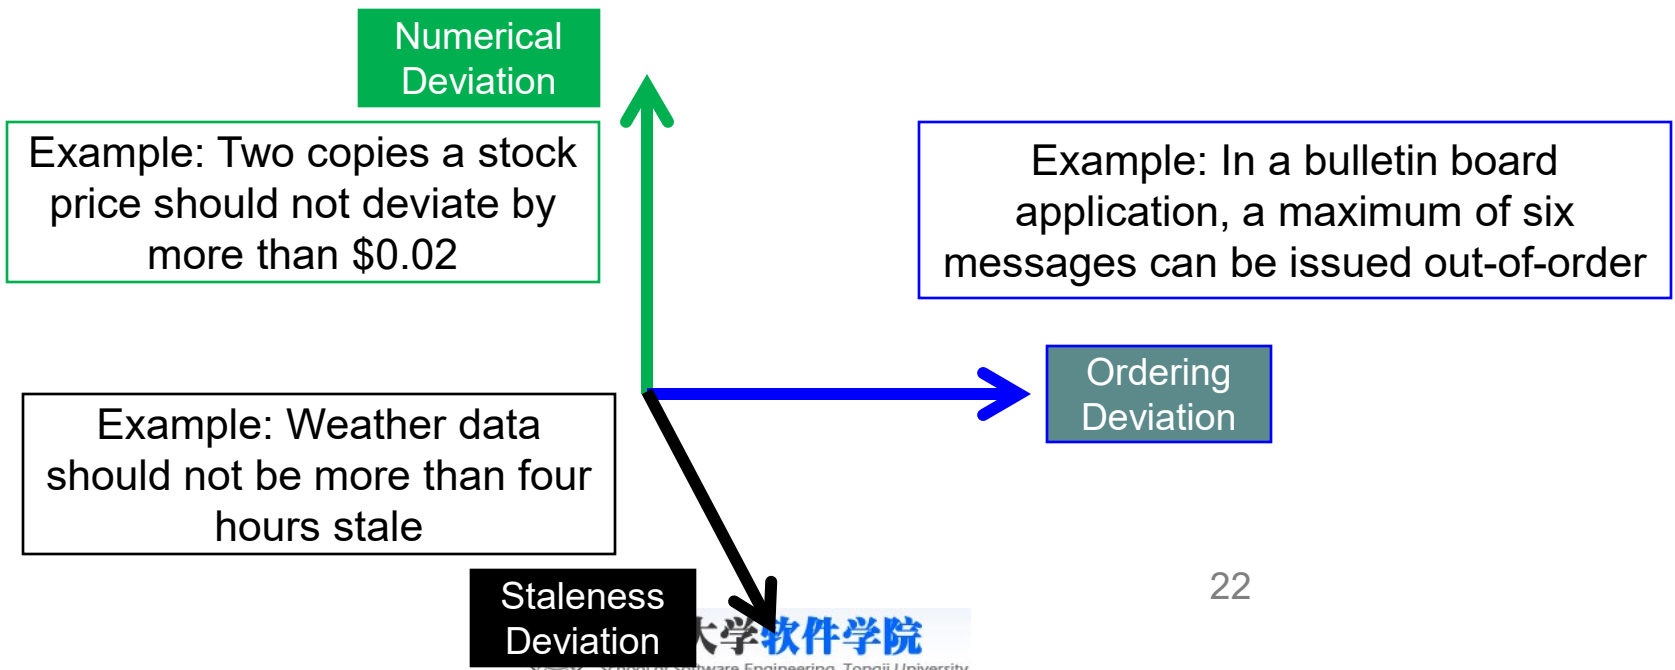
- In replicated data-stores, there should be a mechanism to:
 - ◆ Measure how inconsistent the data might be on different replicas
 - ◆ How replicas and applications can specify the tolerable inconsistency levels
- Consistency Specification Models enable measuring and specifying the level of inconsistency in a replicated data-store
- We study a Consistency Specification Model called *Continuous Consistency Model*

Continuous Consistency Model

- Continuous Consistency Model is used to measure inconsistencies and express what inconsistencies can be expected in the system
- Yu and Vahdat [1] provided a framework for measuring and expressing consistency in replicated data-stores

Continuous Consistency Ranges

- Level of consistency is defined over three independent axes:
 - ◆ **Numerical Deviation**: Deviation in the numerical values between replicas
 - ◆ **Order Deviation**: Deviation with respect to the ordering of update operations
 - ◆ **Staleness Deviation**: Deviation in the staleness between replicas



Consistency Unit (Conit)

- Consistency unit (Conit) specifies the data unit over which consistency is measured
 - ◆ For example, conit can be defined as a record representing a single stock
- Level of consistency is measured by each replica along the three dimensions
 - ◆ Numerical Deviation
 - For a given replica R, how many updates at other replicas are not yet seen at R? What is the effect of the non-propagated updates on local Conit values?
 - ◆ Order Deviation
 - For a given replica R, how many local updates are not propagated to other replicas?
 - ◆ Staleness Deviation
 - For a given replica R, how long has it been since updates were propagated?

Example of Conit and Consistency Measures

Order Deviation at a replica R is the number of operations in R that are not present at the other replicas

Numerical Deviation at replica R is defined as $n(w)$, where
 n = # of operations at other replicas that are not yet seen by R,
 w = weight of the deviation
 = max(update amount of all variables in a Conit)

Replica A					Replica B				
x	y	VC	Ord	Num	x	y	VC	Ord	Num
0	0	(0,0)	0	0(0)	0	0	(0,0)	0	0(0)
0	0	(0,0)	0	1(2)	2	0	(0,5)	1	0(0)
2	0	(1,5)	0	0(0)	2	0	(0,5)	0	0(0)
2	1	(10,5)	1	0(0)	2	0	(0,5)	0	1(1)
2	1	(10,5)	1	1(1)	2	1	(0,16)	1	1(1)
3	1	(14,5)	2	1(1)	2	1	(0,16)	1	2(2)
3	4	(23,5)	3	1(1)	2	1	(0,16)	1	3(4)

Replica A

x; y

Operation		Result
<5,B>	x+=2	x=2
<10,A>	y+=1	y=1
<14,A>	x+=1	x=3
<23,A>	y+=3	y=4

Replica B

x; y

Operation		Result
<5,B>	x+=2	x=2
<16,B>	y+=1	y=1

<5,B> = Operation performed at B when the vector clock was 5

<m,n> = Uncommitted operation

<m,n> = Committed operation

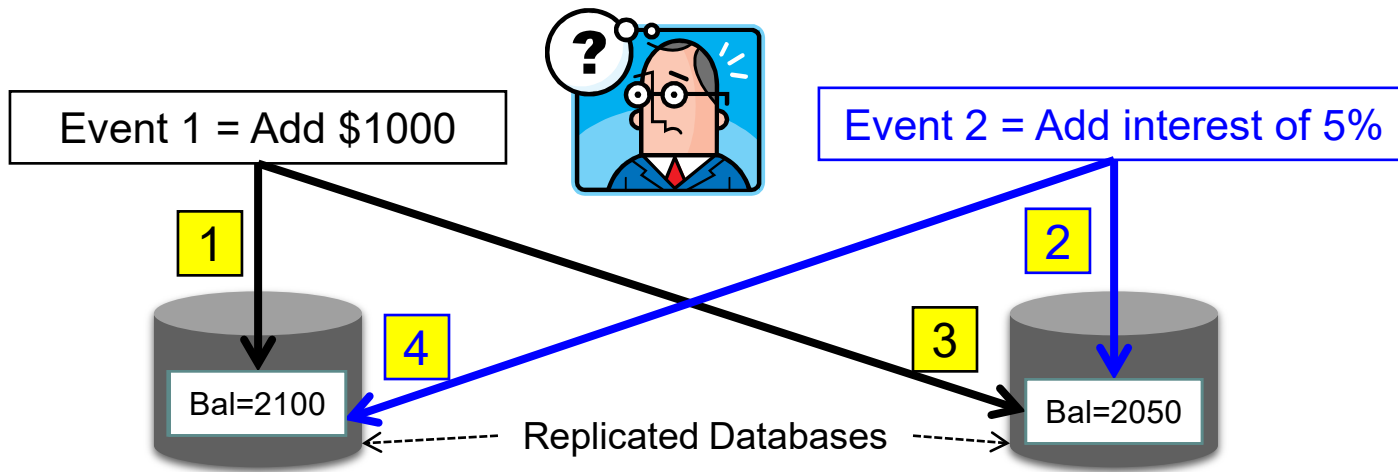
x;y = A Conit

Overview

- Consistency Models
 - ◆ Data-Centric Consistency Models
 - Continuous Specification Models
 - Models for Consistent Ordering of Operations
 - ◆ Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

Why is Consistent Ordering Required in Replication?

- In several applications, the order or the sequence in which the replicas commit to the data store is critical
- Example:



- Continuous Specification Models defined how inconsistency is measured
 - ◆ However, the models did not enforce any order in which the data is committed

Consistent Ordering of Operations (cont'd)

- Whenever a replica is updated, it propagates the updates to other replicas at some point in time
- Updating different replicas is carried out by passing messages between the replica data-stores
- We will study different types of ordering and consistency models arising from these orderings

Types of Ordering

- We will study three types of ordering of messages that meet the needs of different applications:
 1. Total Ordering
 2. Sequential Ordering
 - i. Sequential Consistency Model
 3. Causal Ordering
 - i. Causal Consistency Model

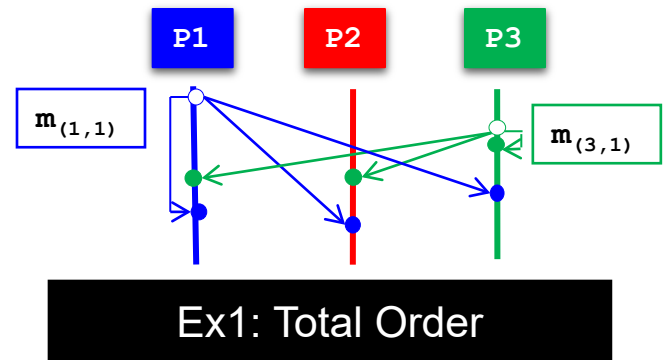
Types of Ordering

1. Total Ordering
2. Sequential Ordering
3. Causal Ordering

Total Ordering

◆ Total Order

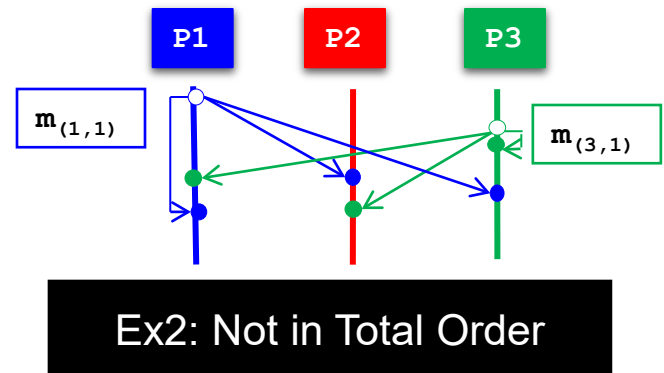
- If process P_i sends a message m_i and P_j sends m_j , and if one correct process delivers m_i before m_j then every correct process delivers m_i before m_j



- Messages can contain replica updates, such as passing the read or write operation that needs to be performed at each replica
 - ◆ In the example Ex1, if P_1 issues the operation $m_{(1,1)}: x=x+1;$ and
 - ◆ If P_3 issues $m_{(3,1)}: \text{print}(x);$
 - ◆ Then, at all replicas P_1, P_2, P_3 the following order of operations are executed

`print(x);`

`x=x+1;`

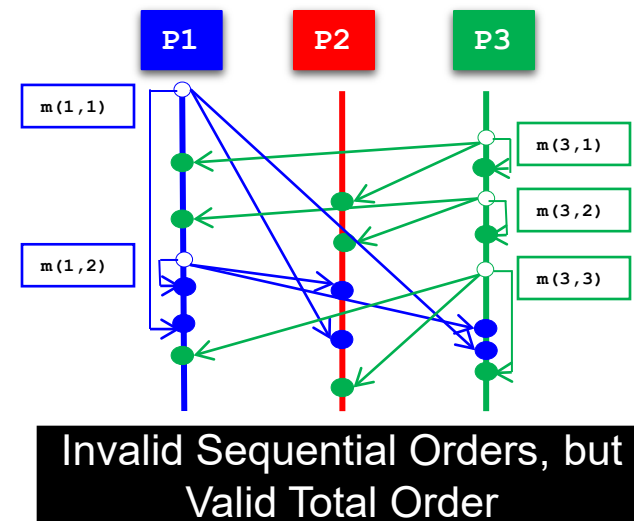
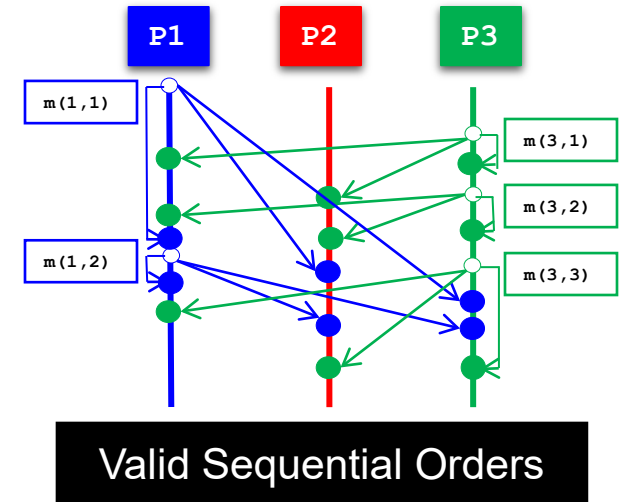


Types of Ordering

1. Total Ordering
2. Sequential Ordering
3. Causal Ordering

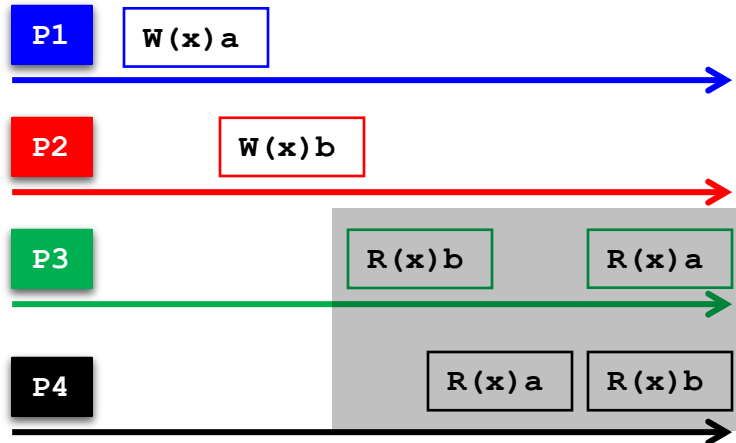
Sequential Ordering

- + If a process P_i sends a sequence of messages $m_{(i,1)}, \dots, m_{(i,n_i)}$, and
- + Process P_j sends a sequence of messages $m_{(j,1)}, \dots, m_{(j,n_j)}$,
- + Then, :
 - + At any process, the set of messages received are in some sequential order
 - + Messages from each individual process appear in this sequence in the order sent by the sender
 - + At every process, $m_{i,1}$ should be delivered before $m_{i,2}$, which is delivered before $m_{i,3}$ and so on...
 - + At every process, $m_{j,1}$ should be delivered before $m_{j,2}$, which is delivered before $m_{j,3}$ and so on...

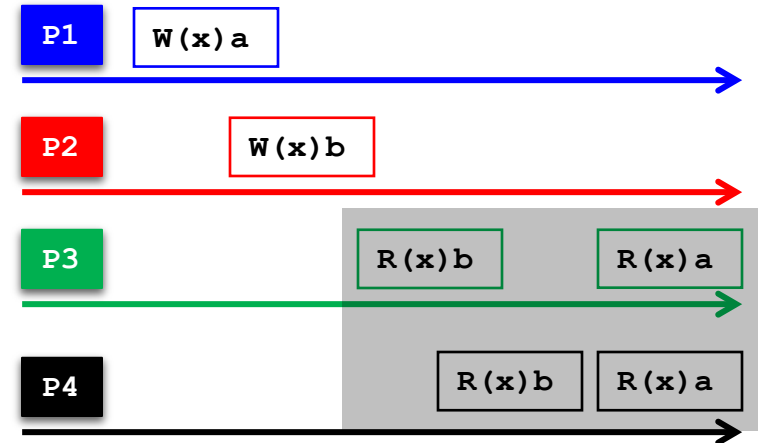


Sequential Consistency Model

- Sequential Consistency Model enforces that all the update operations are executed at the replicas in a sequential order
- Consider a data-store with variable x (Initialized to **NULL**)
 - ◆ In the two data-stores below, identify the sequentially consistent data-store



✗ Results while operating on DATA-STORE-1



✓ Results while operating on DATA-STORE-2

P1

=Process P1



=Timeline at P1

R(x) b

=Read variable x;
Result is b

W(x) b

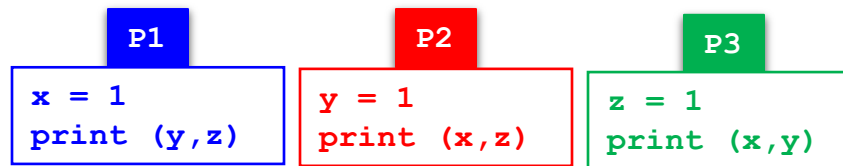
= Write variable x;
Result is b



Sequential Consistency (cont'd)

- Consider three processes P_1 , P_2 and P_3 executing multiple instructions on three shared variables x , y and z

- Assume that x , y and z are set to zero at start



- There are many valid sequences in which operations can be executed at the replica respecting sequential consistency

- Identify the output

	<pre>x = 1 print (y,z) y = 1 print (x,z) z = 1 print (x,y)</pre>	<pre>x = 1 y = 1 print (x,z) print (y,z) z = 1 print (x,y)</pre>	<pre>z = 1 print (x,y) print (x,z) y = 1 x = 1 print (y,z)</pre>	<pre>y = 1 z = 1 print (x,y) print (x,z) x = 1 print (y,z)</pre>
Output	001011	101011	000111	010111



Implications of Adopting Sequential Consistency Model for Applications

- There might be several different sequentially consistent combinations of ordering
 - ◆ Number of combinations for a total of n instructions = $O(n!)$
- The contract between the process and the distributed data-store is that the process must accept all of the sequential orderings as valid results
 - ◆ A process that works for some of the sequential orderings and does not work correctly for others is INCORRECT

Summary

- Replication is necessary for improving performance, scalability and availability, and for providing fault-tolerance
- Replicated data-stores should be designed after carefully evaluating the trade-off between tolerable data inconsistency and efficiency
- Consistency Models describe the contract between the data-store and process about what form of consistency to expect from the system
- Data-centric consistency models:
 - ◆ Continuous Consistency Models provide mechanisms to measure and specify inconsistencies
 - ◆ Consistency Models can be defined based on the type of ordering of operations that the replica guarantees the applications
 - ▣ We studied Sequential Consistency Model

Next Class

- Consistency Models

- ◆ Causal Consistency Model
- ◆ Client-Centric Consistency Models

- Replica Management

- ◆ Replica management studies:
 - ▣ when, where and by whom replicas should be placed
 - ▣ which consistency model to use for keeping replicas consistent

- Consistency Protocols

- ◆ We study various implementations of consistency models

References

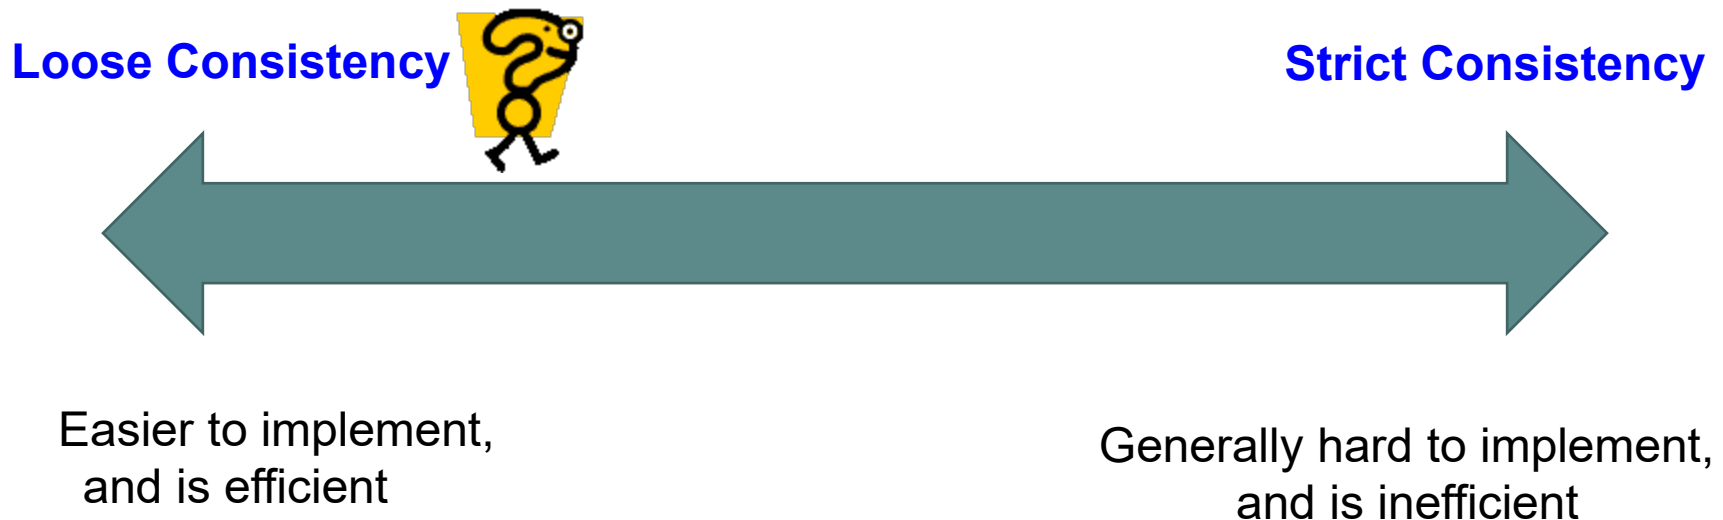
- [\[1\] Haifeng Yu and Amin Vahdat, “Design and evaluation of a conit-based continuous consistency model for replicated services”](#)
- [\[2\] http://tech.amikelive.com/node-285/using-content-delivery-networks-cdn-to-speed-up-content-load-on-the-web/](http://tech.amikelive.com/node-285/using-content-delivery-networks-cdn-to-speed-up-content-load-on-the-web/)
- [\[3\] http://en.wikipedia.org/wiki/Replication_\(computer_science\)](http://en.wikipedia.org/wiki/Replication_(computer_science))
- [\[4\] http://en.wikipedia.org/wiki/Content_delivery_network](http://en.wikipedia.org/wiki/Content_delivery_network)
- [\[5\] http://www.cdk5.net](http://www.cdk5.net)
- [\[6\] http://www.dis.uniroma1.it/~baldoni/ordered%2520communication%25202008.ppt](http://www.dis.uniroma1.it/~baldoni/ordered%2520communication%25202008.ppt)
- [\[7\] http://www.cs.uiuc.edu/class/fa09/cs425/L5tmp.ppt](http://www.cs.uiuc.edu/class/fa09/cs425/L5tmp.ppt)

Today' session

- Last Session
 - Consistency and Replication
 - Introduction and Data-Centric Consistency Models
- **Today's session**
 - Consistency and Replication – Part II
 - Finish Data-centric Consistency Models
 - Client-Centric Consistency Models

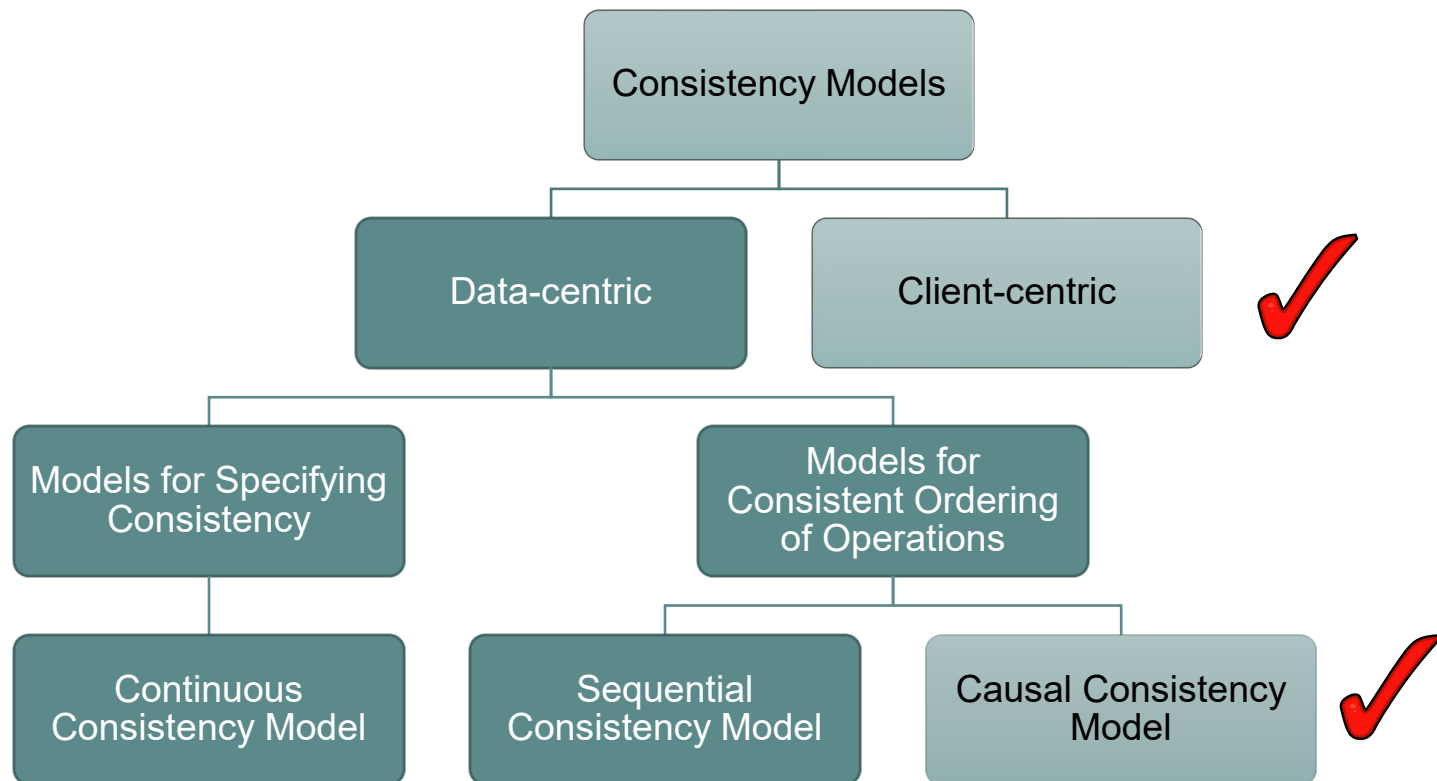
Recap: Trade-offs in Maintaining Consistency

- Maintaining consistency should balance between the strictness of consistency versus efficiency
 - ◆ How much consistency is “good-enough” depends on the application



Consistency Models

- A consistency model states the level of consistency provided by the *data-store* to the processes while reading and writing the data



Types of Ordering

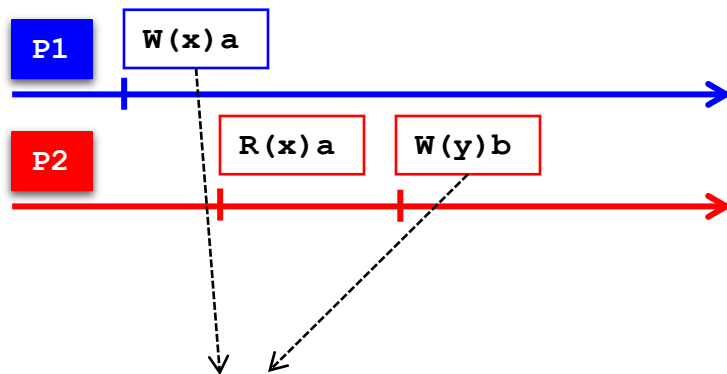
1. Total Ordering
2. Sequential Ordering
3. Causal Ordering

Causality (Recap)

- ◆ Causal relation between two events
 - If a and b are two events such that a happened-before b (i.e., $a \rightarrow b$), and
 - If the (logical) times when events a and b occur at a process P_i are denoted as $C_i(a)$ and $C_i(b)$
 - Then, if we can infer that $a \rightarrow b$ by observing that $C_i(a) < C_i(b)$, then a and b are *causally* related
- Causality can be implemented using Vector Clocks

Causal vs. Concurrent events

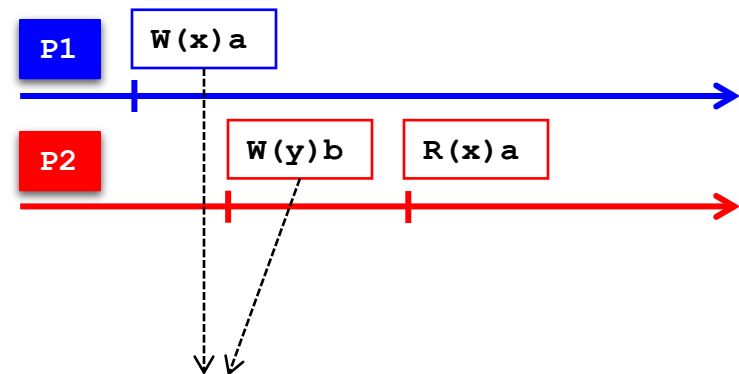
- Consider an interaction between processes P_1 and P_2 operating on replicated data x and y



Events are causally related

Events are not concurrent

- Computation of y at P_2 may have depended on value of x written by P_1



Events are not causally related

Events are concurrent

- Computation of y at P_2 does not depend on value of x written by P_1

P_1 = Process P_1 \longrightarrow = Timeline at P_1

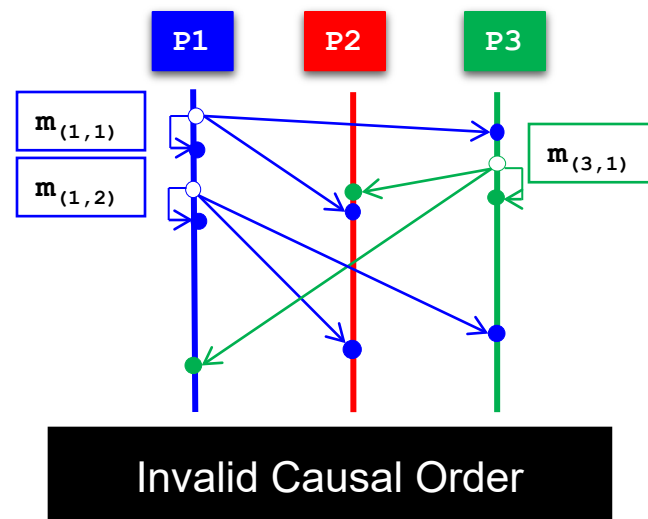
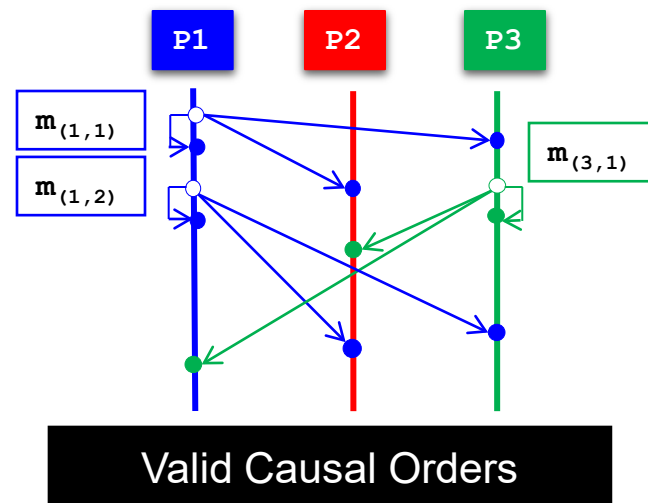
$R(x) b$ = Read variable x ;
Result is b

$W(x) b$ = Write variable x ;
Result is b

Causal Ordering

◆ Causal Order

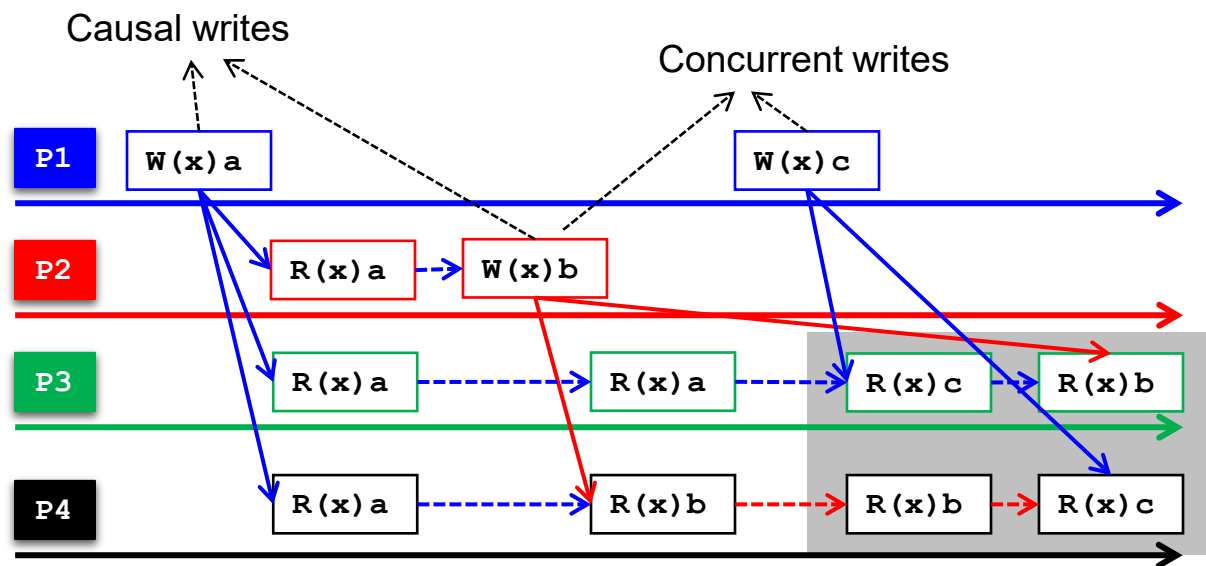
- If process P_i sends a message m_i and P_j sends m_j , and if $m_i \rightarrow m_j$ (operator ' \rightarrow ' is Lamport's **happened-before** relation) then any correct process that delivers m_j will deliver m_i before m_j
- In the example, $m_{(1,1)}$ and $m_{(3,1)}$ are in Causal Order
- **Drawback:**
 - ◆ The **happened-before** relation between m_i and m_j should be induced before communication



Causal Consistency Model

- A data-store is causally consistent if:
 - ◆ Writes that are potentially causally related must be seen by all the processes in the same order
 - ◆ Concurrent writes may be seen in a different order on different machines

Example of a Causally Consistent Data-store



A Causally Consistent
Data-Store

But not a Sequentially
Consistent Data-Store

P1 = Process P1 \rightarrow = Timeline at P1

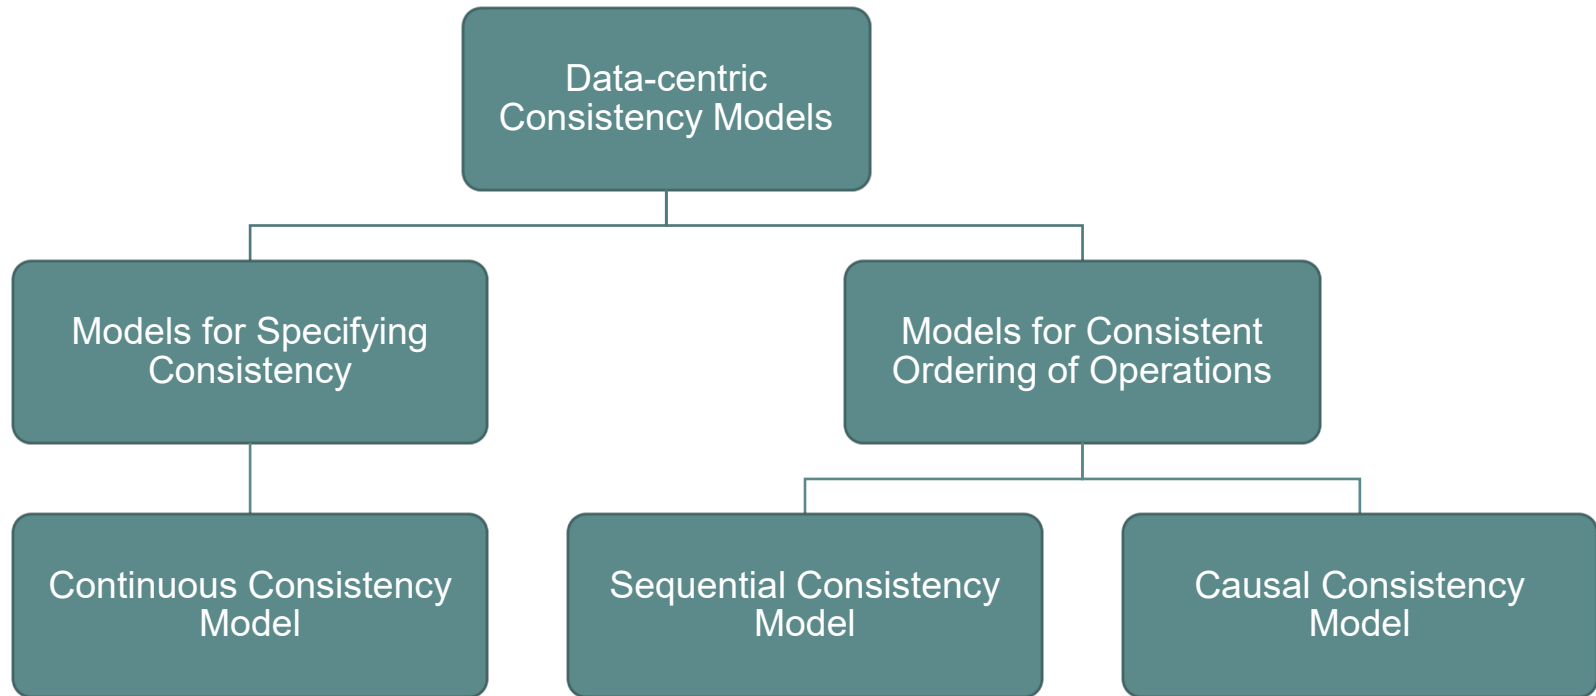
$R(x) b$ = Read variable x ;
Result is b

$W(x) b$ = Write variable x ;
Result is b

Implications of adopting a Causally Consistent Data-store for Applications

- Processes have to keep track of which processes have seen which writes
- This requires maintaining a dependency graph between write and read operations
 - ◆ Vector clocks provides a way to maintain causally consistent data-base

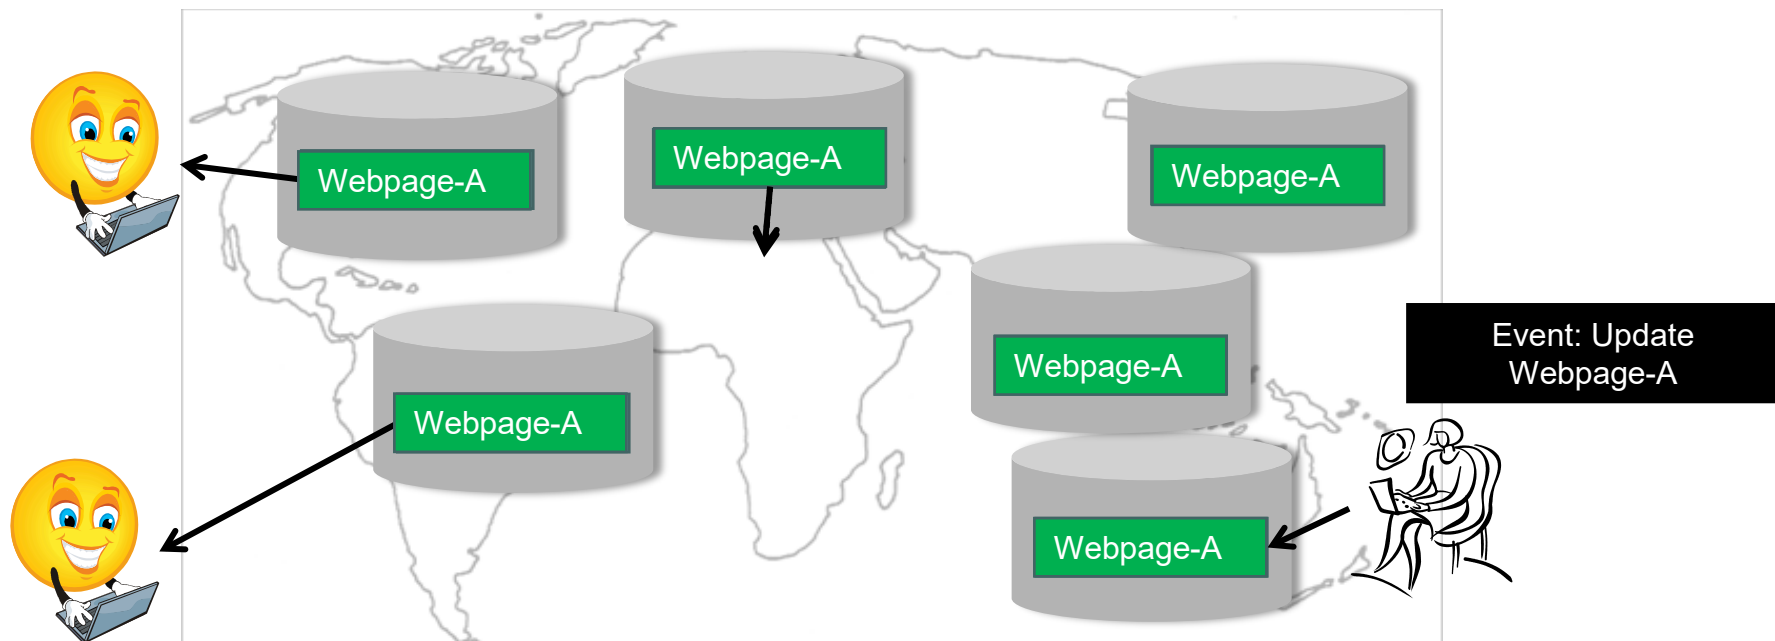
Topics Covered in Data-centric Consistency Models



But, is Data-centric Consistency Model good for all applications?

Applications that Can Use Data-centric Models

- Data-centric models are applicable when many processes are concurrently updating the data-store
- But, do all applications need all replicas to be consistent?



Data-Centric Consistency Model is too strict when

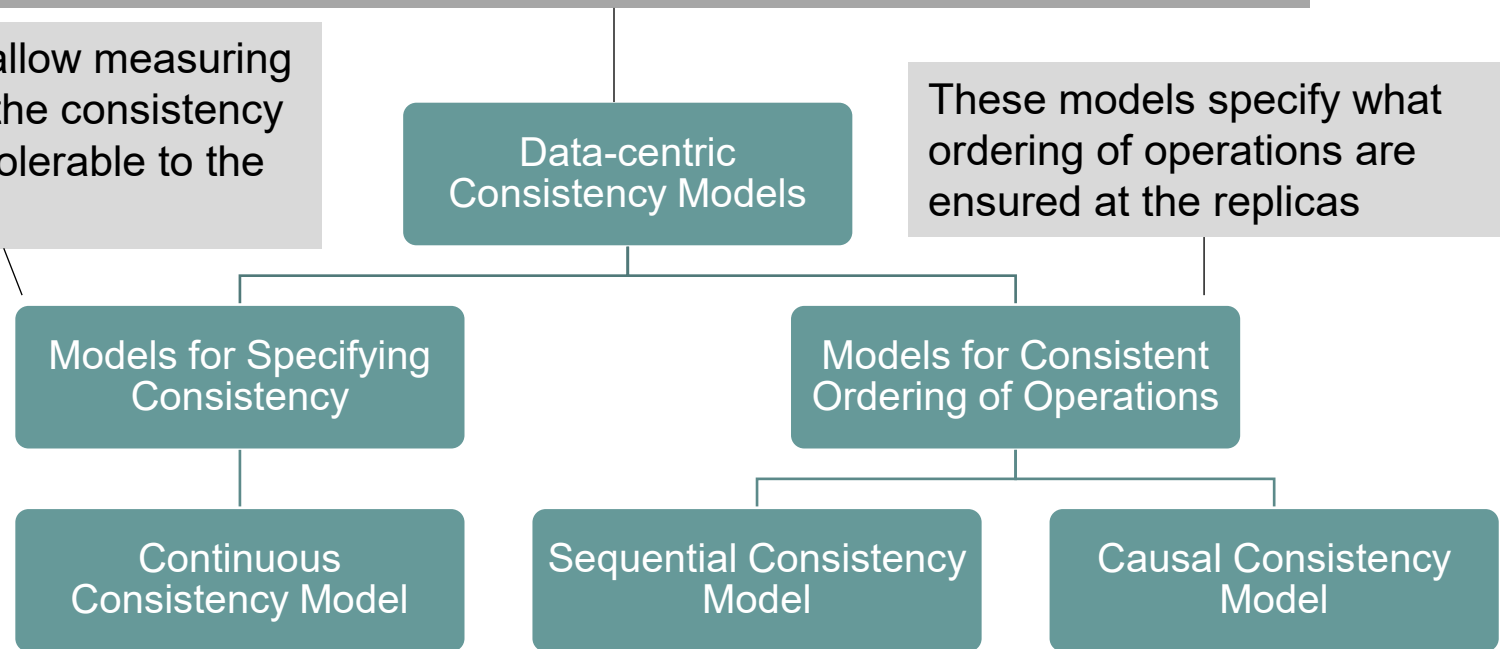
- One client process updates the data
- Other processes read the data, and are OK with reasonably stale data

Summary of Data-Centric Consistency Models

Data-centric consistency models describe how the replicated data is kept consistent across different data-stores, and what a process can expect from the data-store

These models allow measuring and specifying the consistency levels that are tolerable to the application

These models specify what ordering of operations are ensured at the replicas

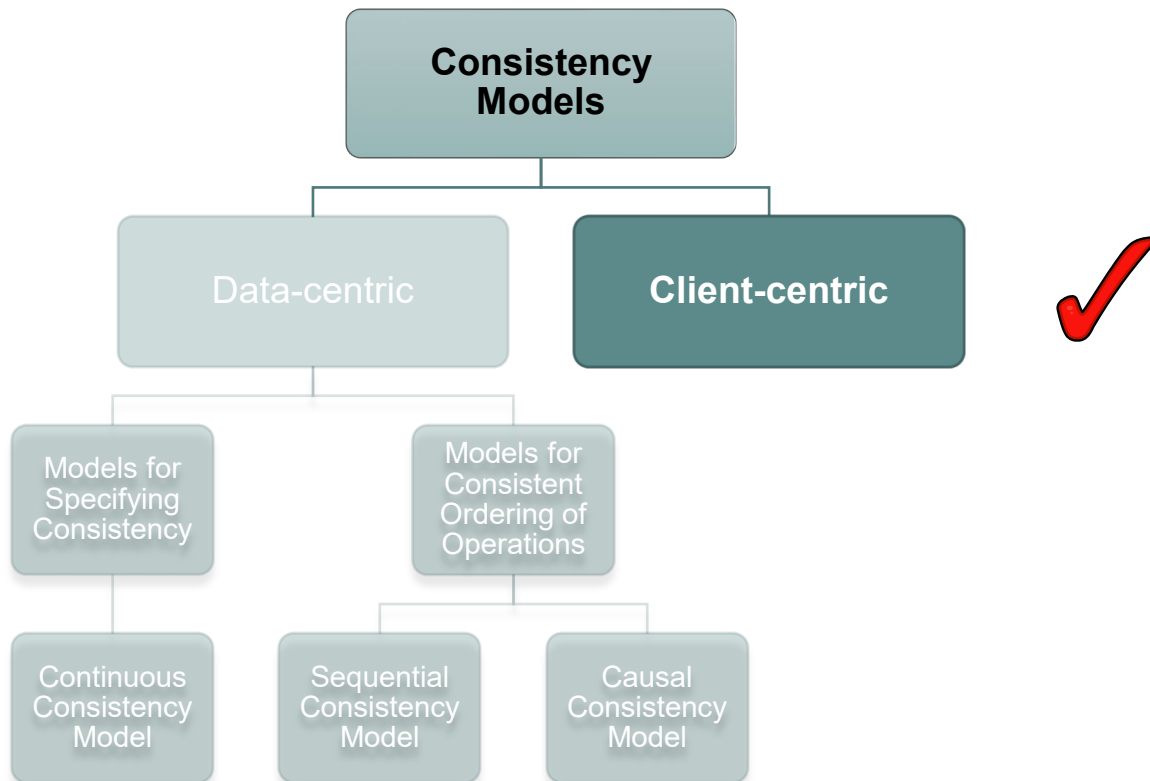


Data-centric models are too strict when:

- Most operations are read operations
- Updates are generally triggered from one client process



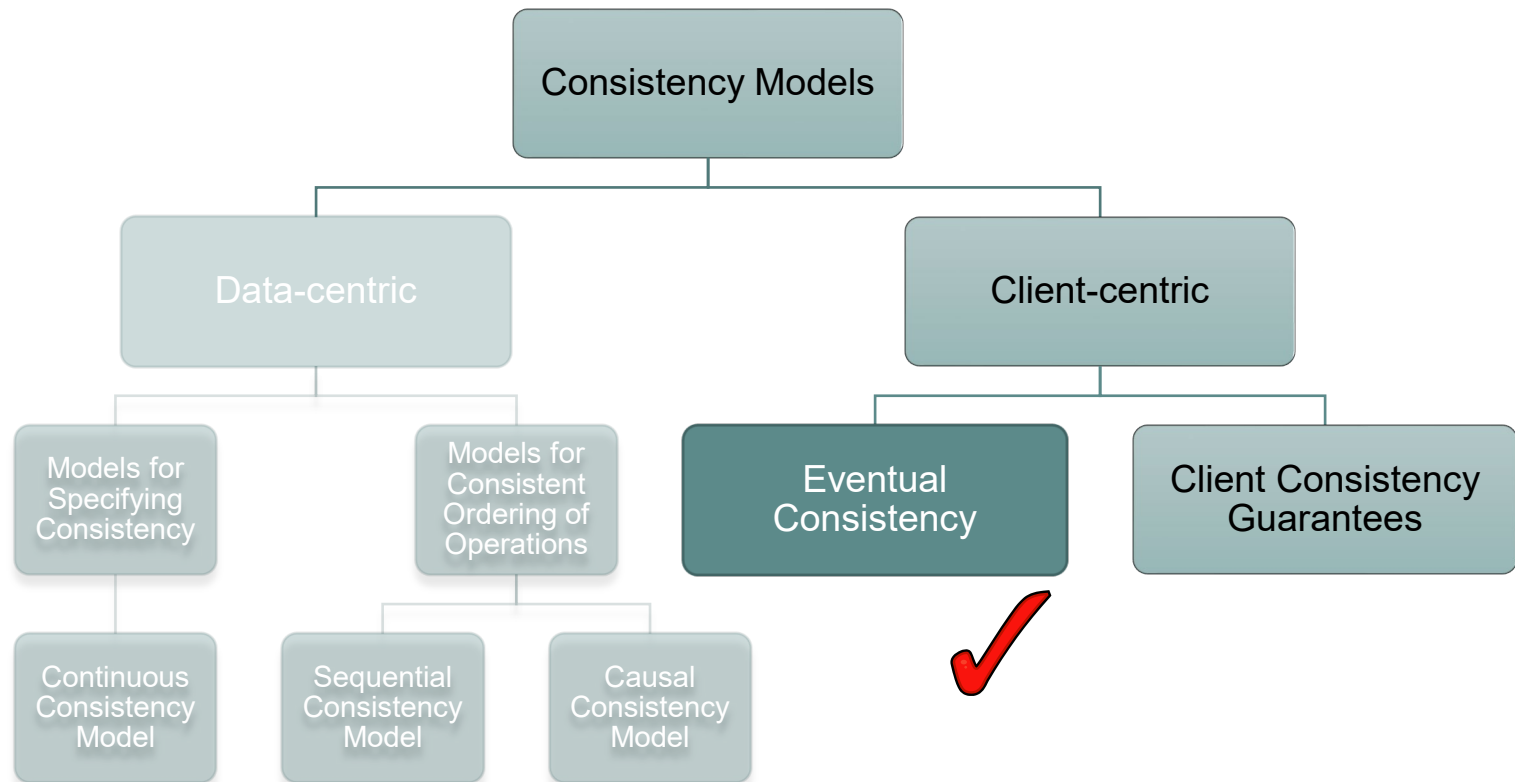
Overview



Client-Centric Consistency Models

- Data-centric models lead to excessive overheads in applications where:
 - ◆ a majority operations are reads, and
 - ◆ updates occur frequently, and are often from one client process
- For such applications, a weaker form of consistency called *Client-centric Consistency* is employed for improving efficiency
- Client-centric consistency models specify two requirements:
 1. Client Consistency Guarantees
 - A client events should be guaranteed some level of consistency while accessing the data value at different replicas
 2. Eventual Consistency
 - All the replicas should *eventually* converge on a final value

Overview



Eventual Consistency

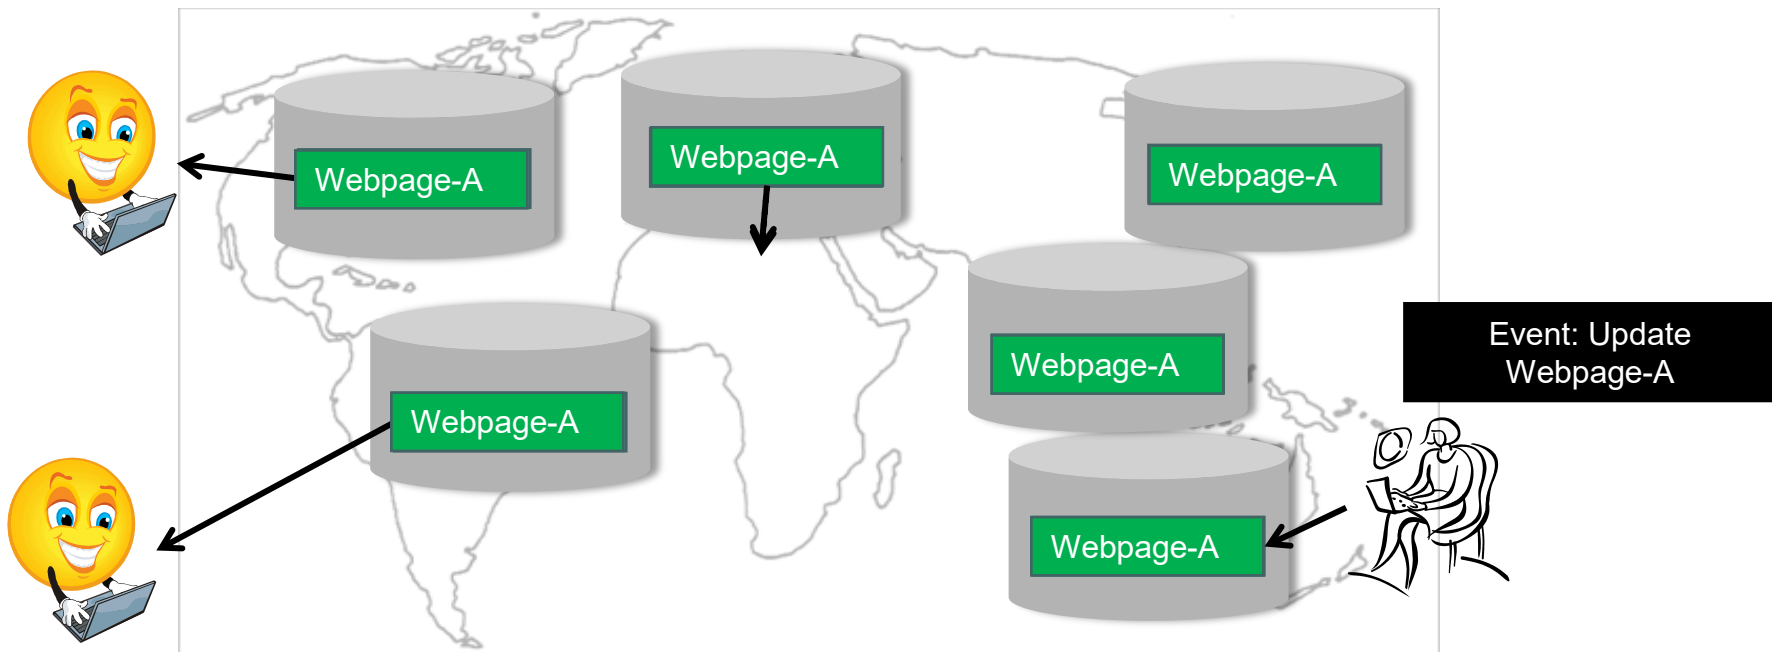
- Many applications can tolerate inconsistency for a long time
 - ◆ Webpage updates, Web Search – Crawling, indexing and ranking, Updates to DNS Server
- In such applications, it is acceptable and efficient if replicas in the data-store rarely exchange updates
- A data-store is termed as *Eventually Consistent* if:
 - ◆ All replicas will gradually become consistent in the absence of updates
- Typically, updates are propagated infrequently in eventually consistent data-stores

Designing Eventual Consistency

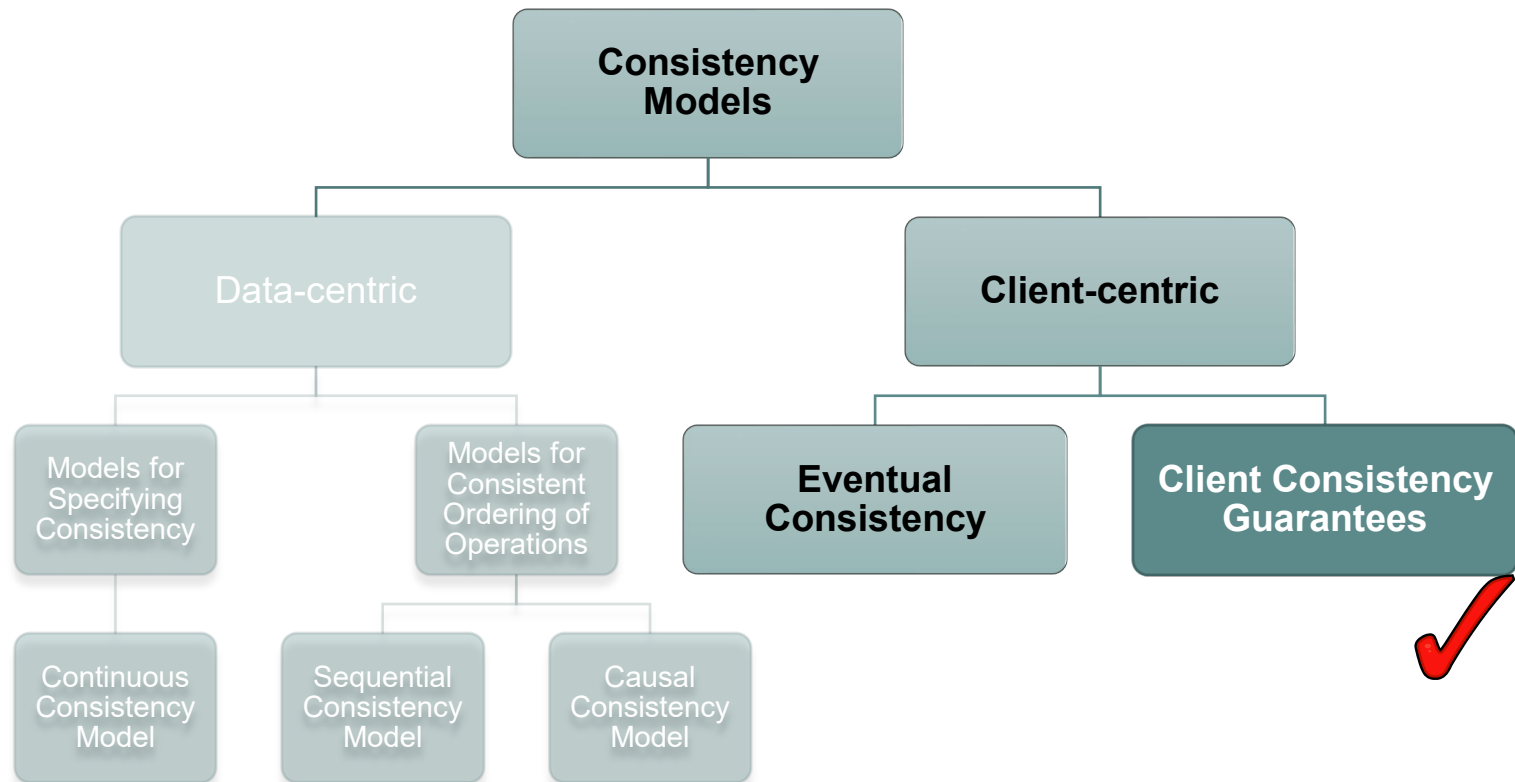
- In eventually consistent data-stores,
 - ◆ *Write-write conflicts* are rare
 - Two processes that write the same value are rare
 - Generally, one client updates the data value
 - e.g., One DNS server updates the name to IP mapping
 - Such rare conflicts can be handled through simple mechanisms, such as mutual exclusion
 - ◆ *Read-write conflicts* are more frequent
 - Conflicts where one process is reading a value, while another process is writing a value to the same variable
 - Eventual Consistency Design has to focus on efficiently resolving such conflicts

Challenges in Eventual Consistency

- Eventual Consistency is not good-enough when the client process accesses data from different replicas
 - ◆ We need consistency guarantees for a single client while accessing the data-store

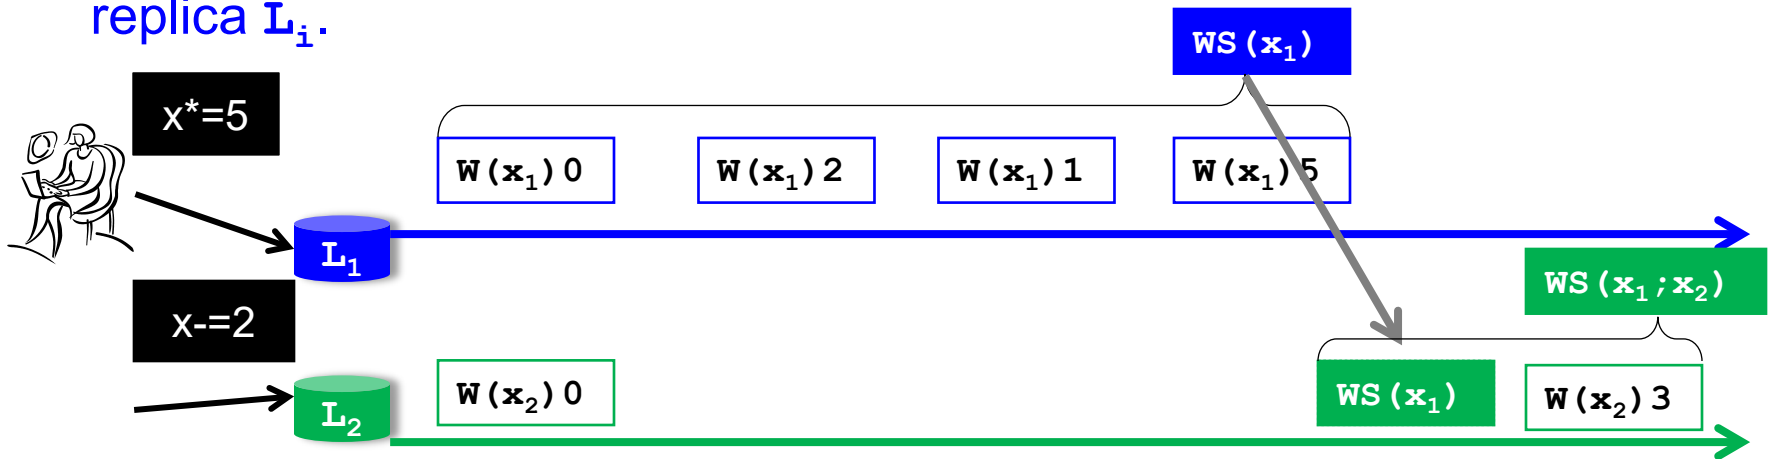


Overview



Client Consistency Guarantees

- Client-centric consistency provides guarantees for a single client for its accesses to a data-store
- Example: Providing consistency guarantee to a client process for data x replicated on two replicas. Let x_i be the local copy of a data x at replica L_i .



$WS(x_1)$ = Write Set for x_1 = Series of ops being done at some replica that reflects how L_1 updated x_1 till this time

$WS(x_1; x_2)$ = Write Set for x_1 and x_2 = Series of ops being done at some replica that reflects how L_1 updated x_1 and, later on, how x_2 is updated on L_2

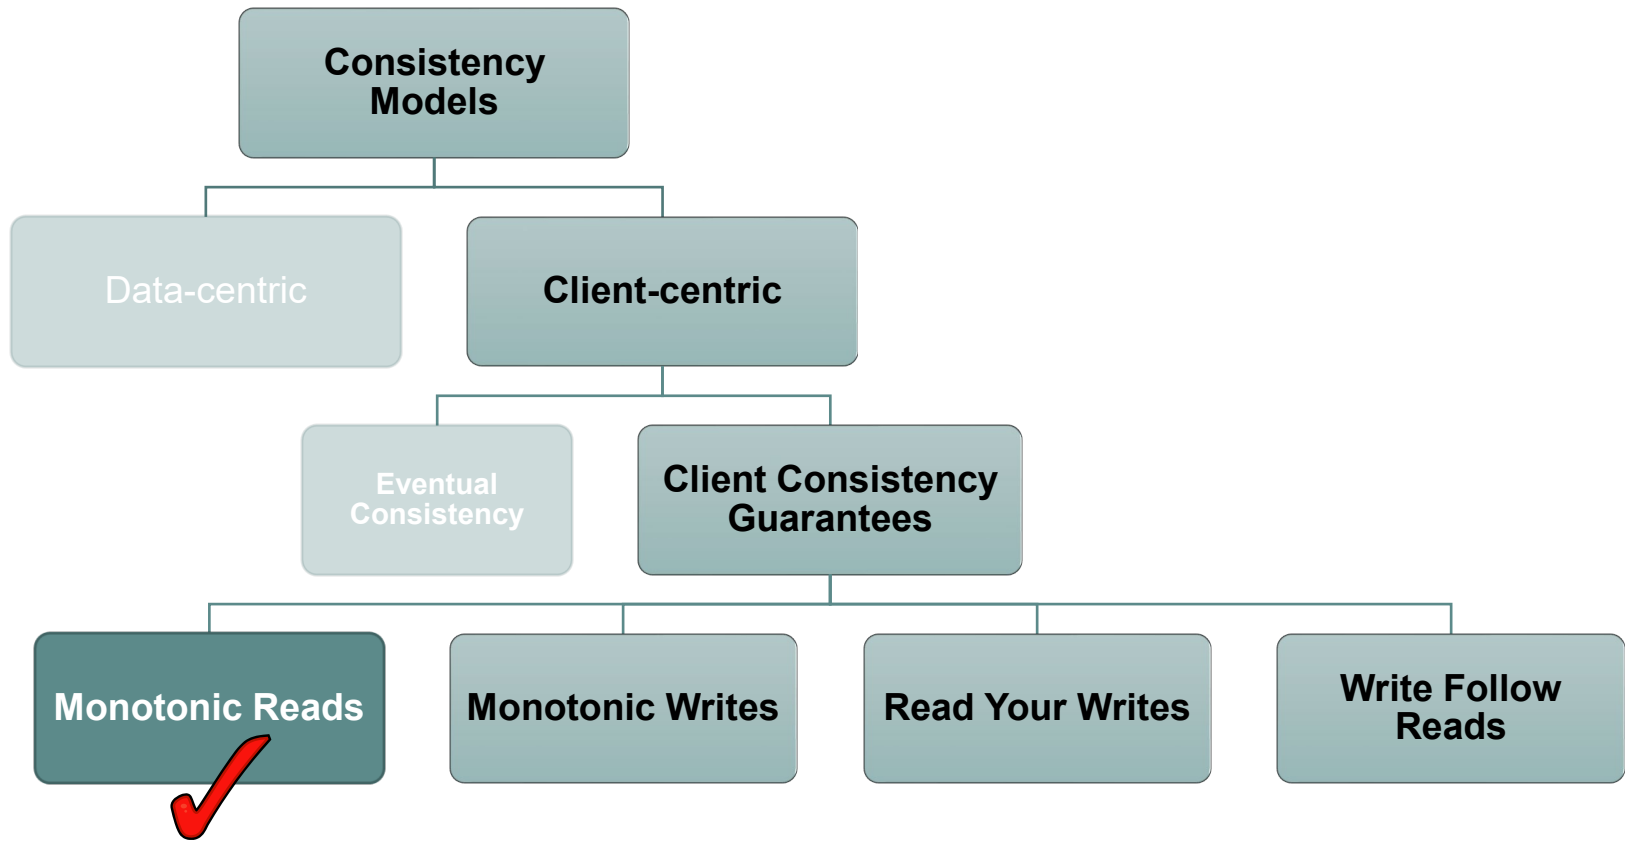
L_i = Replica i
 $R(x_i) b$ = Read variable x at replica i ; Result is b
 $W(x) b$ = Write variable x at replica i ; Result is b
 $WS(x_i)$ = Write Set

Client Consistency Guarantees

- We will study four types of client-centric consistency models¹
 1. Monotonic Reads
 2. Monotonic Writes
 3. Read Your Writes
 4. Write Follow Reads

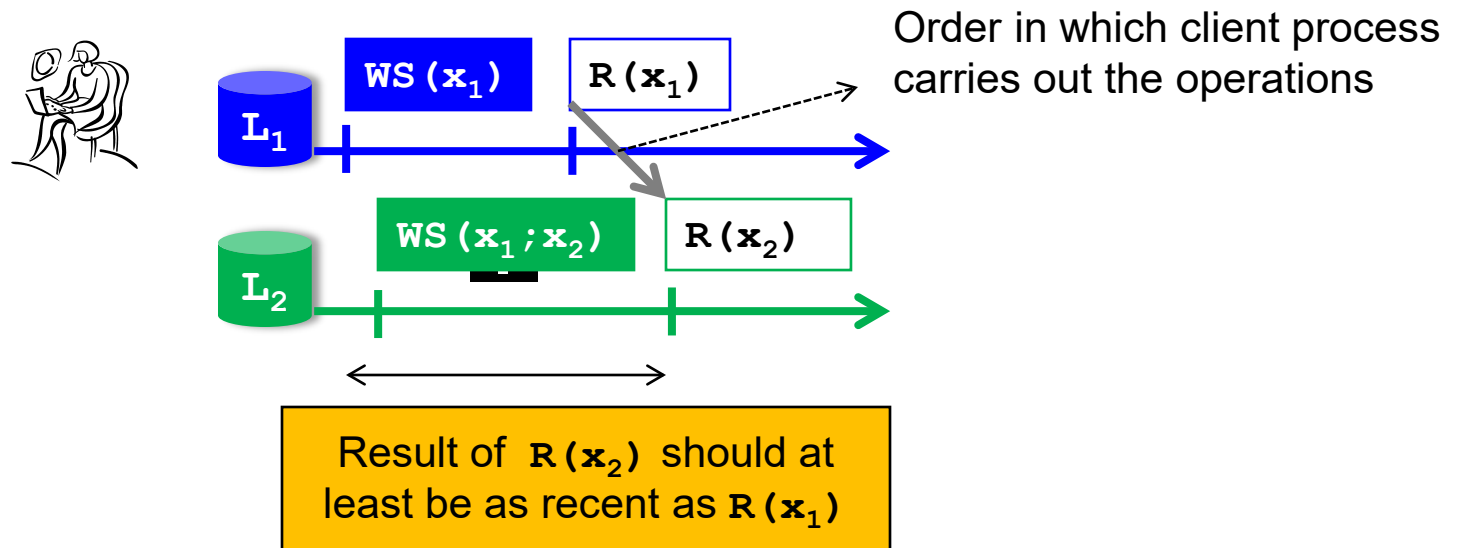
1. The work is based on the distributed database system built by Terry et al. [1]

Overview



Monotonic Reads

- The model provides guarantees on successive reads
- If a client process reads the value of data item x , then any successive read operation by that process should return the same or a more recent value for x



Monotonic Reads - Puzzle

Recognize data-stores that provide monotonic read guarantees

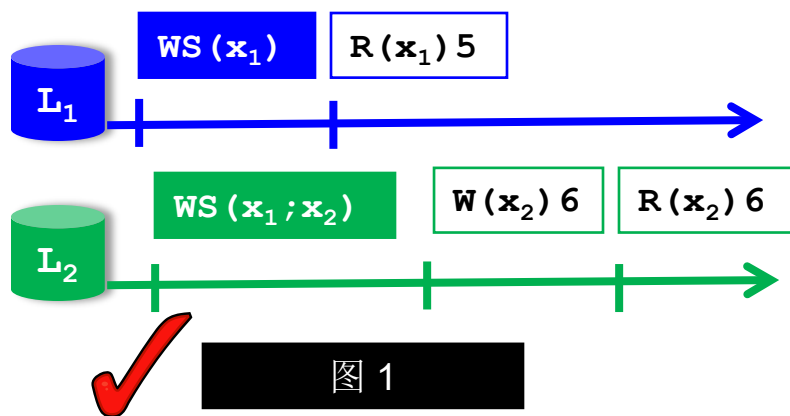


图 1

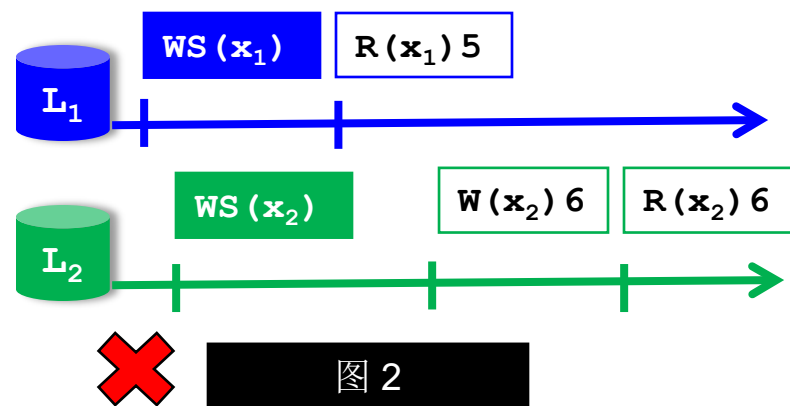


图 2

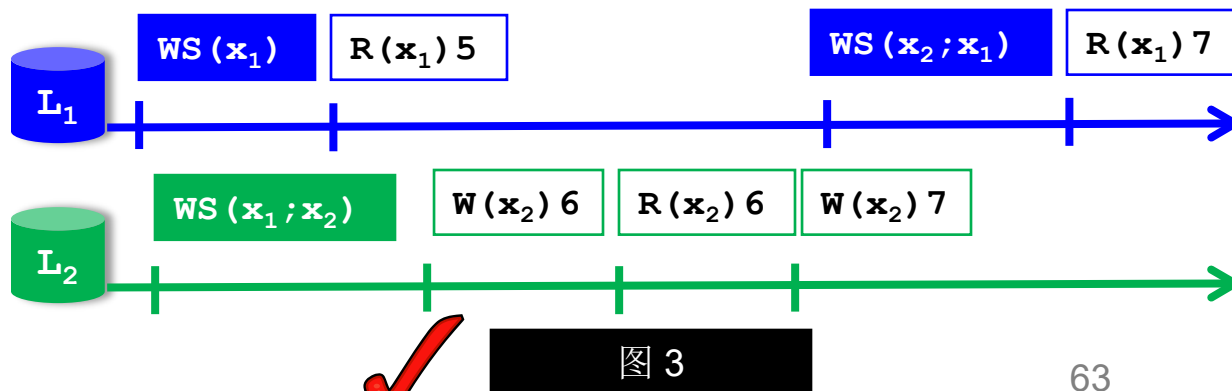


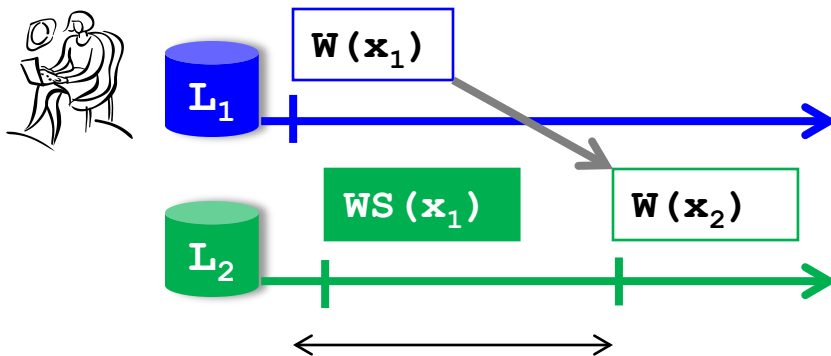
图 3

Overview

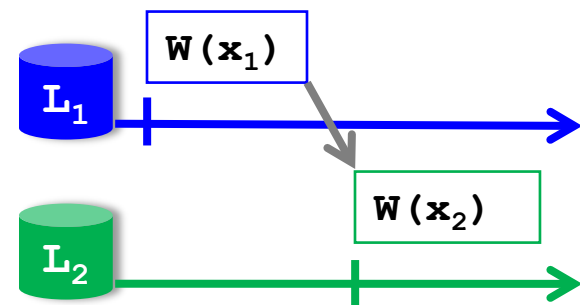


Monotonic Writes

- This consistency model assures that writes are monotonic
- A write operation by a client process on a data item x is completed before any successive write operation on x by the same process
 - ◆ A new write on a replica should wait for all old writes on any replica



$W(x_2)$ operation should be performed only after the result of $W(x_1)$ has been updated at L_2



The data-store does not provide monotonic write consistency

Monotonic Writes - An Example

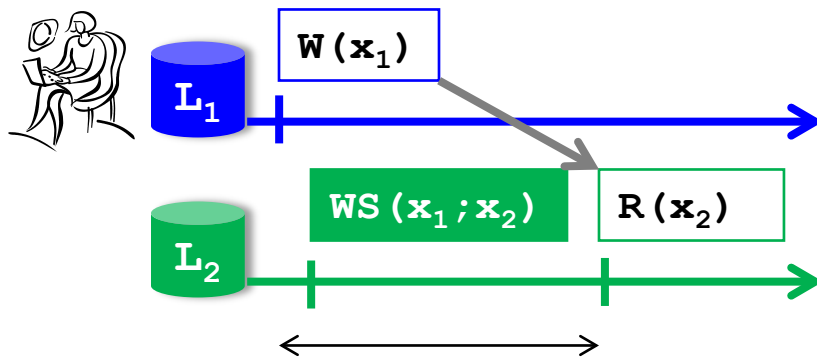
- Example: Updating individual libraries in a large software source code which is replicated
 - ◆ Updates can be propagated in a lazy fashion
 - ◆ Updates are performed on a part of the data item
 - Some functions in an individual library is often modified and updated
 - ◆ Monotonic writes: If an update is performed on a library, then all preceding updates on the same library are first updated
- Question: If the update overwrites the complete software source code, is it necessary to update all the previous updates?

Overview

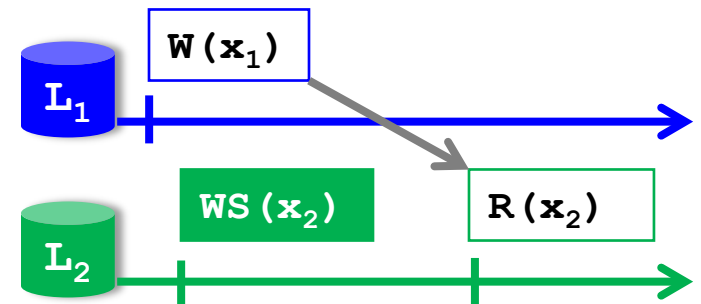


Read Your Writes

- The effect of a write operation on a data item x by a process will always be seen by a successive read operation on x by the same process
- Example scenario:
 - ◆ In systems where password is stored in a replicated data-base, the password change should be seen immediately

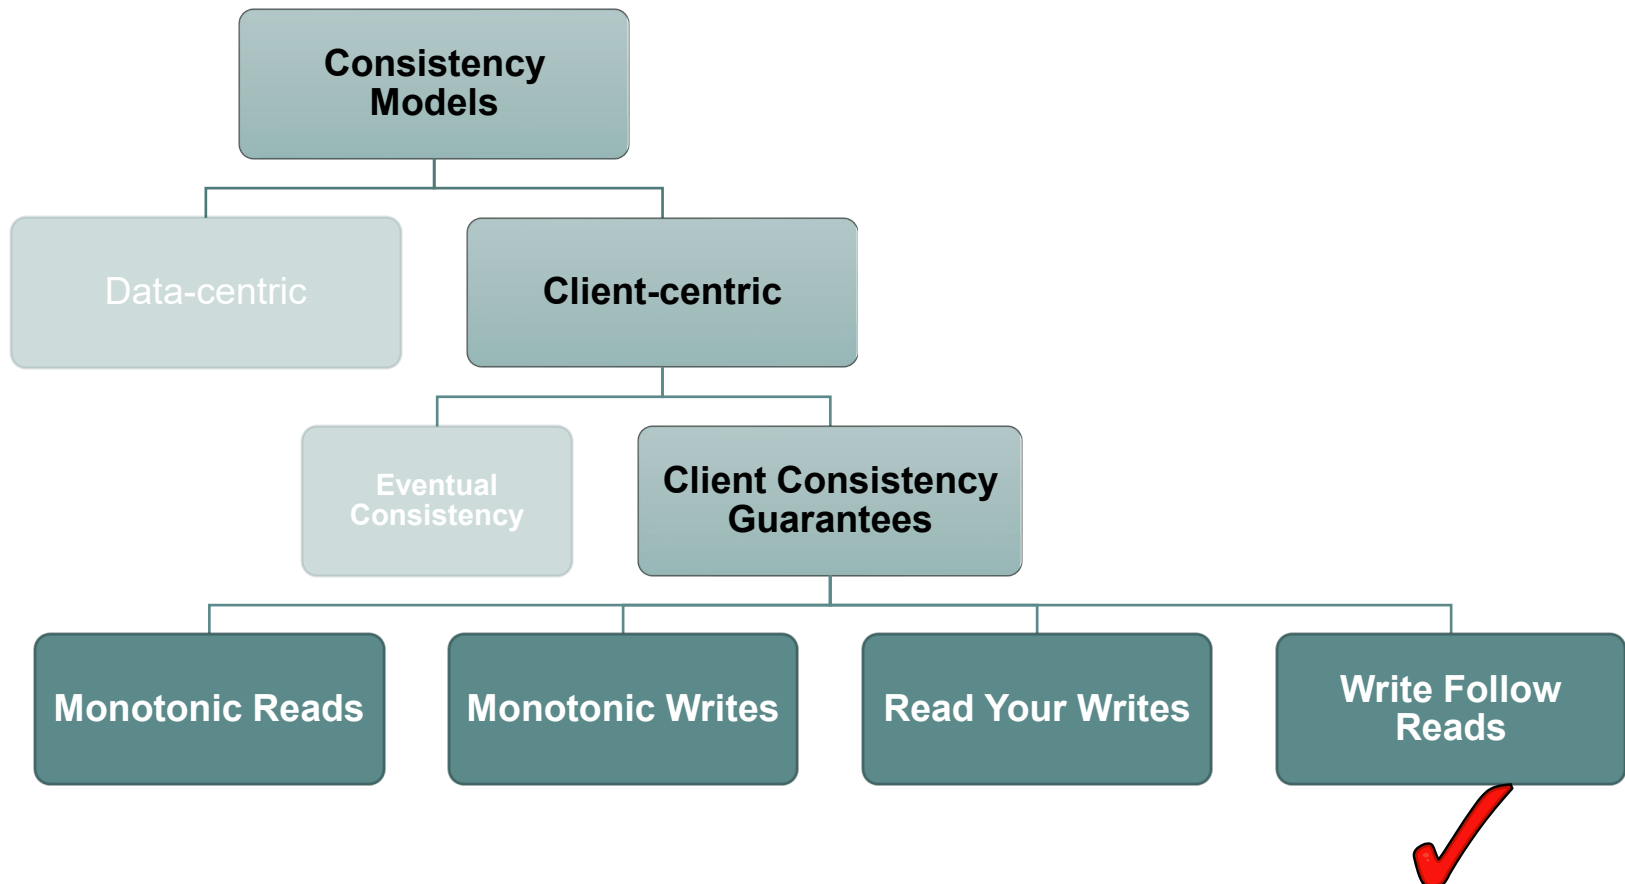


$R(x_2)$ operation should be performed only after the updating of the Write Set $WS(x_1)$ at L_2



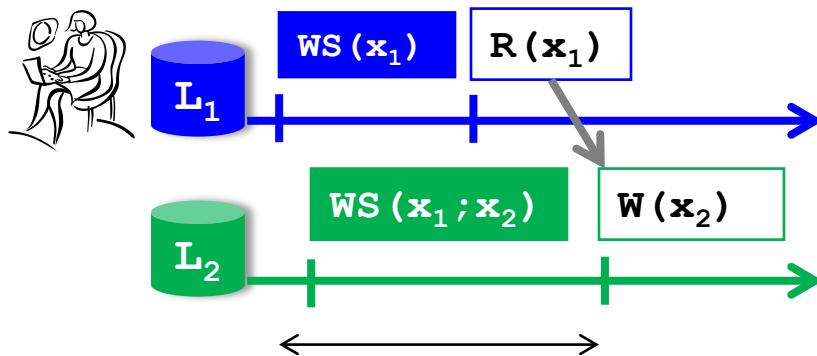
A data-store that does not provide *Read Your Write* consistency

Overview

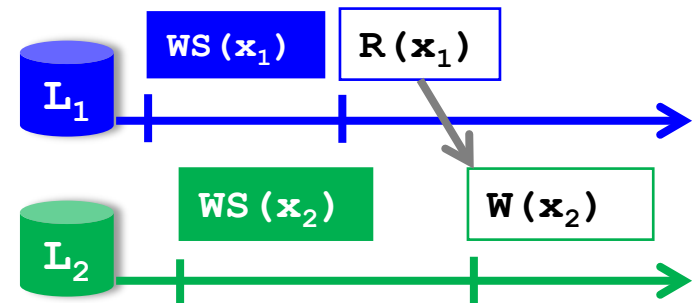


Write Follow Reads

- A write operation by a process on a data item x following a previous read operation on x by the same process is guaranteed to take place on the same or a more recent value of x that was read
- Example scenario:
 - ◆ Users of a newsgroup should post their comments only after they have read all previous comments



$W(x_2)$ operation should be performed only after all previous writes have been seen



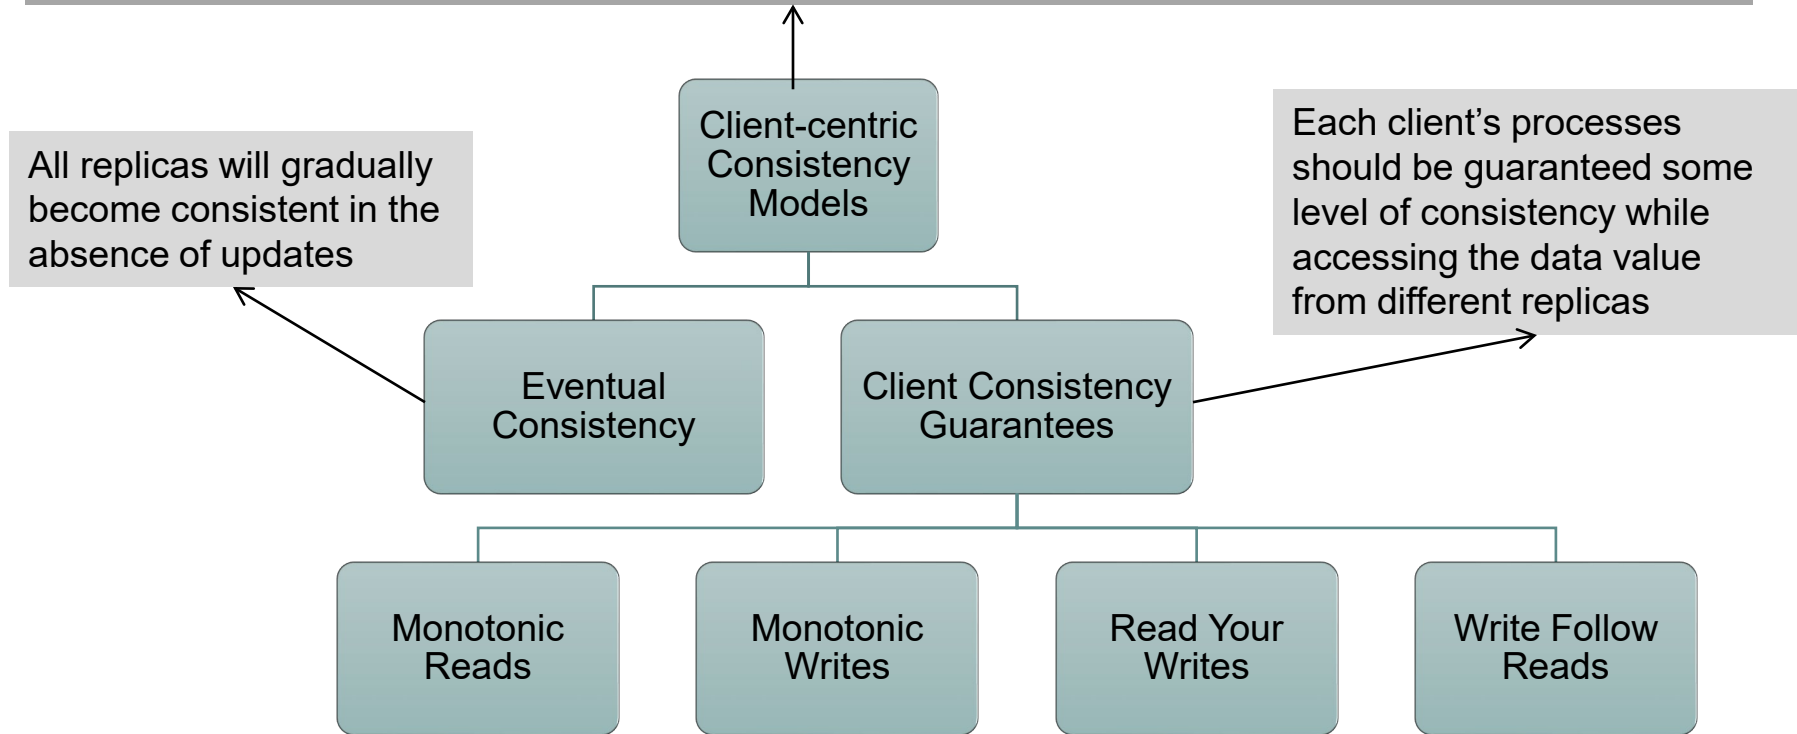
A data-store that does not guarantee Write Follow Read Consistency Model

Summary of Client-centric Consistency Models

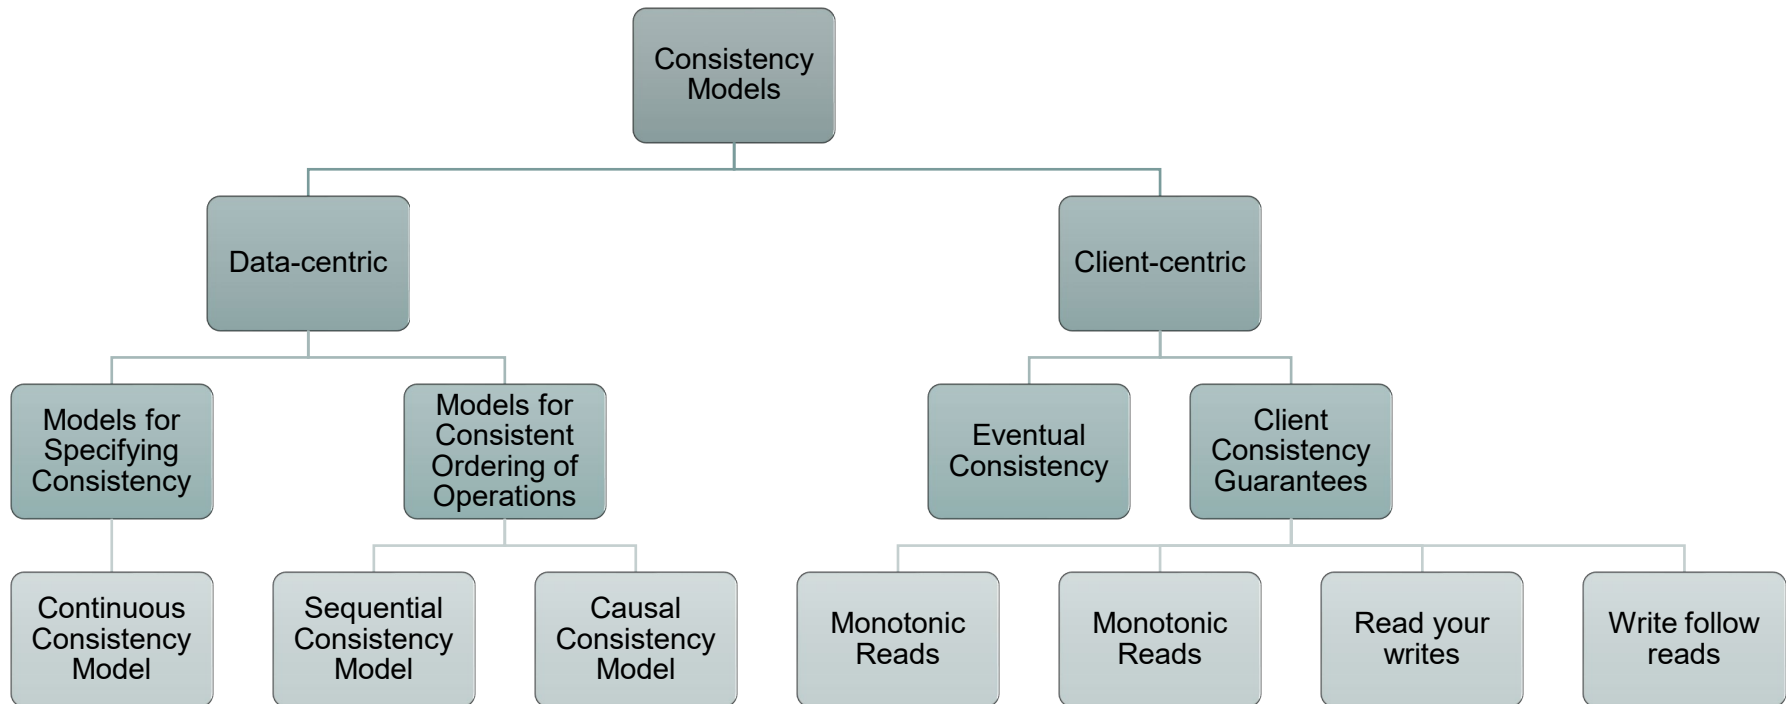
Client-centric Consistency Model defines how a data-store presents the data value to an individual client when the client process accesses the data value across different replicas.

It is generally useful in applications where:

- one client always updates the data-store.
- read-to-write ratio is high.



Topics Covered in Consistency Models



Summary of Consistency Models

- Different applications require different levels of consistency
 - ◆ Data-centric consistency models
 - Define how replicas in a data-store maintain consistency
 - ◆ Client-centric consistency models
 - Provide an efficient, but *weaker* form of consistency
 - One client process updates the data item, and many processes read the replica

Next Class

■ Replica Management

- ◆ Describes where, when and by whom replicas should be placed

■ Consistency Protocols

- ◆ We study “how” consistency is ensured in distributed systems

References

- [1] Terry, D.B., Demers, A.J., Petersen, K., Spreitzer, M.J., Theimer, M.M., Welch, B.B., "Session guarantees for weakly consistent replicated data", Proceedings of the Third International Conference on Parallel and Distributed Information Systems, 1994
- [2] Lili Qiu, Padmanabhan, V.N., Voelker, G.M., "On the placement of Web server replicas", Proceedings of IEEE INFOCOM 2001.
- [3] Rabinovich, M., Rabinovich, I., Rajaraman, R., Aggarwal, A., "A dynamic object replication and migration protocol for an Internet hosting service", Proceedings of IEEE International Conference on Distributed Computing Systems (ICDCS), 1999
- [4] <http://www.cdk5.net>