

分布式系统

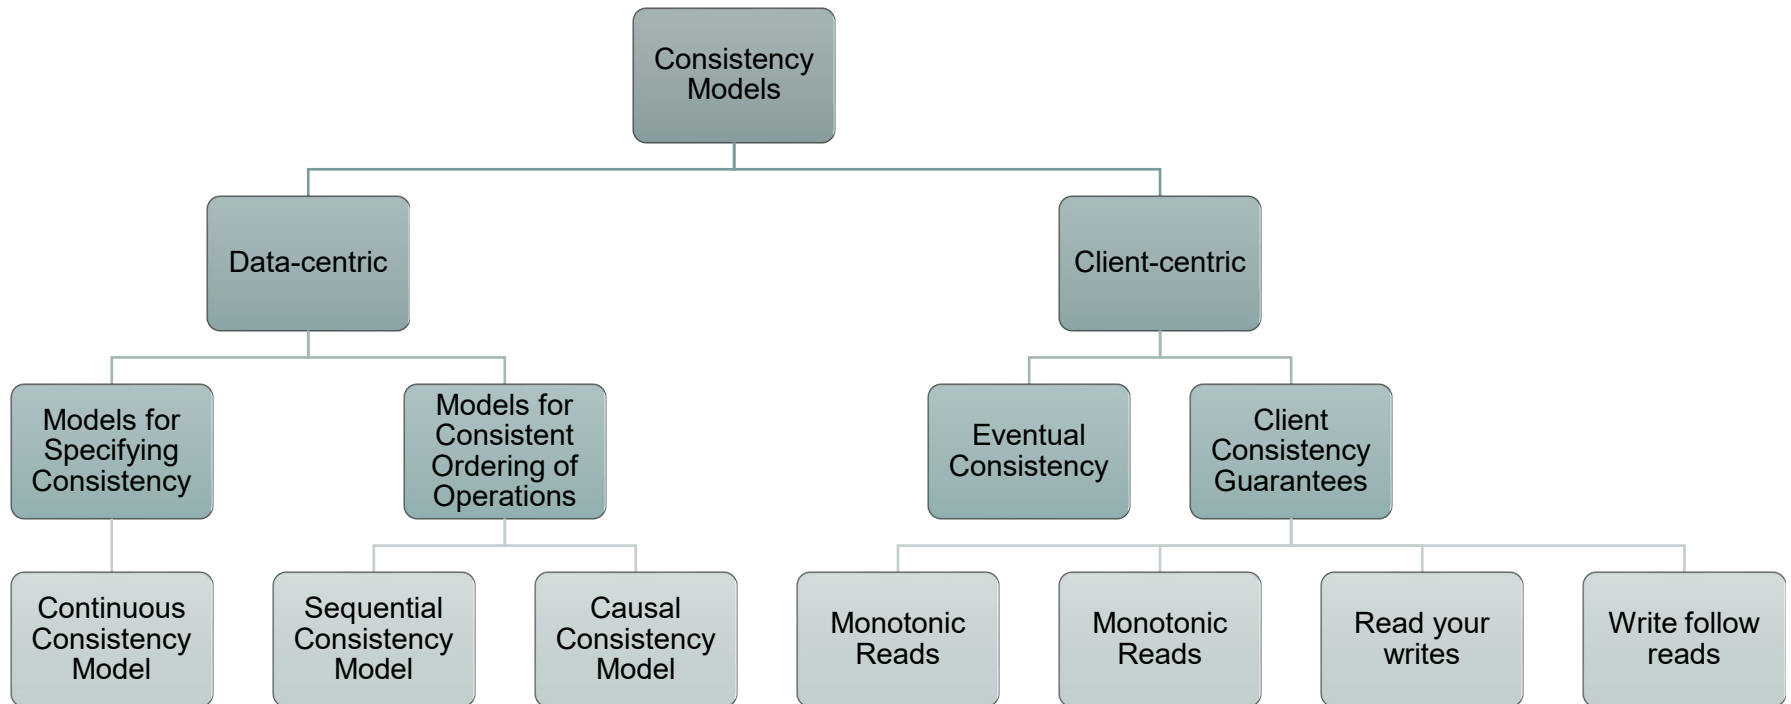
08-分布式系统的一致性与复制 Consistency and Replication in DS

Weixiong Rao 饶卫雄
Tongji University 同济大学软件学院
2023 秋季
wxrao@tongji.edu.cn

Today...

- Last Session
 - Consistency and Replication
 - Consistency Models: Data-centric and Client-centric
- **Today's session**
 - Consistency and Replication – Part III
 - Replica Management
 - Consistency Protocols

Recap: Topics covered in Consistency Models



Overview

- Consistency Models
- **Replica Management**
- Consistency Protocols

Replica Management

- Replica management describes where, when and by whom replicas should be placed
- We will study two problems under replica management
 1. Replica-Server Placement
 - ▣ Decides the best locations to place the replica server that can host data-stores
 2. Content Replication and Placement
 - ▣ Finds the best server for placing the contents

Overview

- Consistency Models
- Replica Management
 - ◆ Replica Server Placement
 - ◆ Content Replication and Placement
- Consistency Protocols

Replica Server Placement

- Factors that affect placement of replica servers:
 - ◆ What are the possible locations where servers can be placed?
 - Should we place replica servers close-by or distribute it uniformly?
 - ◆ How many replica servers can be placed?
 - What are the trade-offs between placing many replica servers vs. few?
 - ◆ How many clients are accessing the data from a location?
 - More replicas at locations where most clients access improves performance and fault-tolerance
- If K replicas have to be placed out of N possible locations, find the *best K out of N* locations ($K < N$)

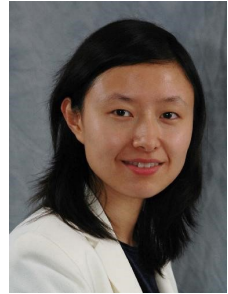
Replica Server Placement – An Example Approach

- Problem: K replica servers should be placed on some of the N possible replica sites such that
 - ◆ Clients have low-latency/high-bandwidth connections

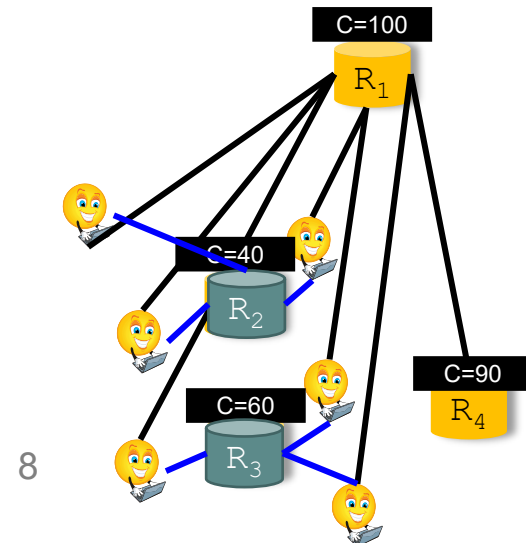
- Qiu *et al.* [2] suggested a Greedy Approach

1. Evaluate the cost of placing a replica on each of the N potential sites
 - + Examining the cost of C clients connecting to the replica
 - + Cost of a link can be $1/\text{bandwidth}$ or latency
2. Choose the lowest-cost site
3. In the second iteration, search for a second replica site which, in conjunction with the already selected site, yields the lowest cost
4. Iterate steps 2,3 and 4 until K replicas are chosen

邱锂力博士



- 微软亚洲研究院副院长，负责微软亚洲研究院（上海）的研究工作，以及与产学研各界的合作
- 国际计算机学会无线及移动系统专委 ACM SIGMOBILE的主席

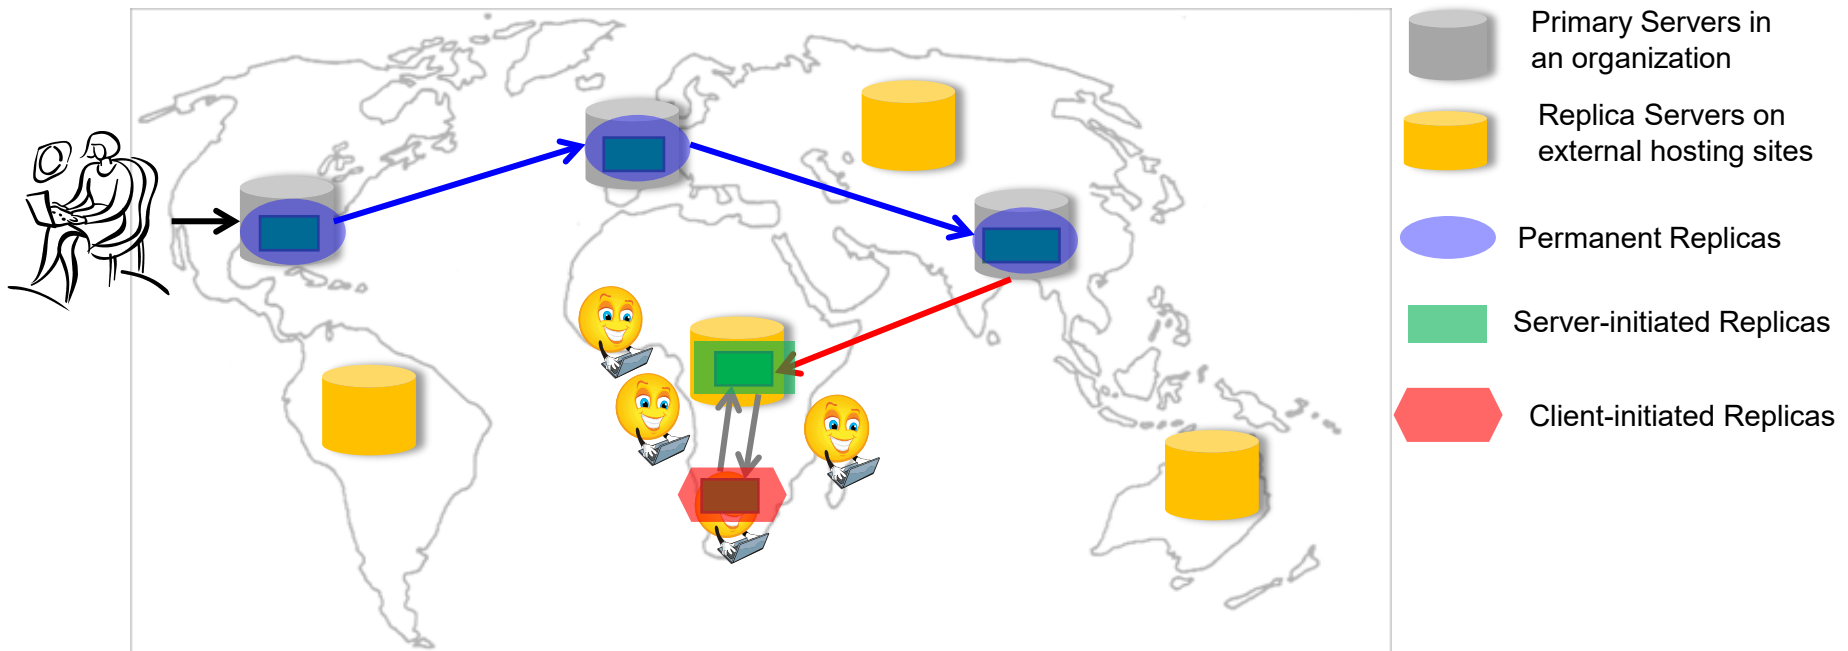


Overview

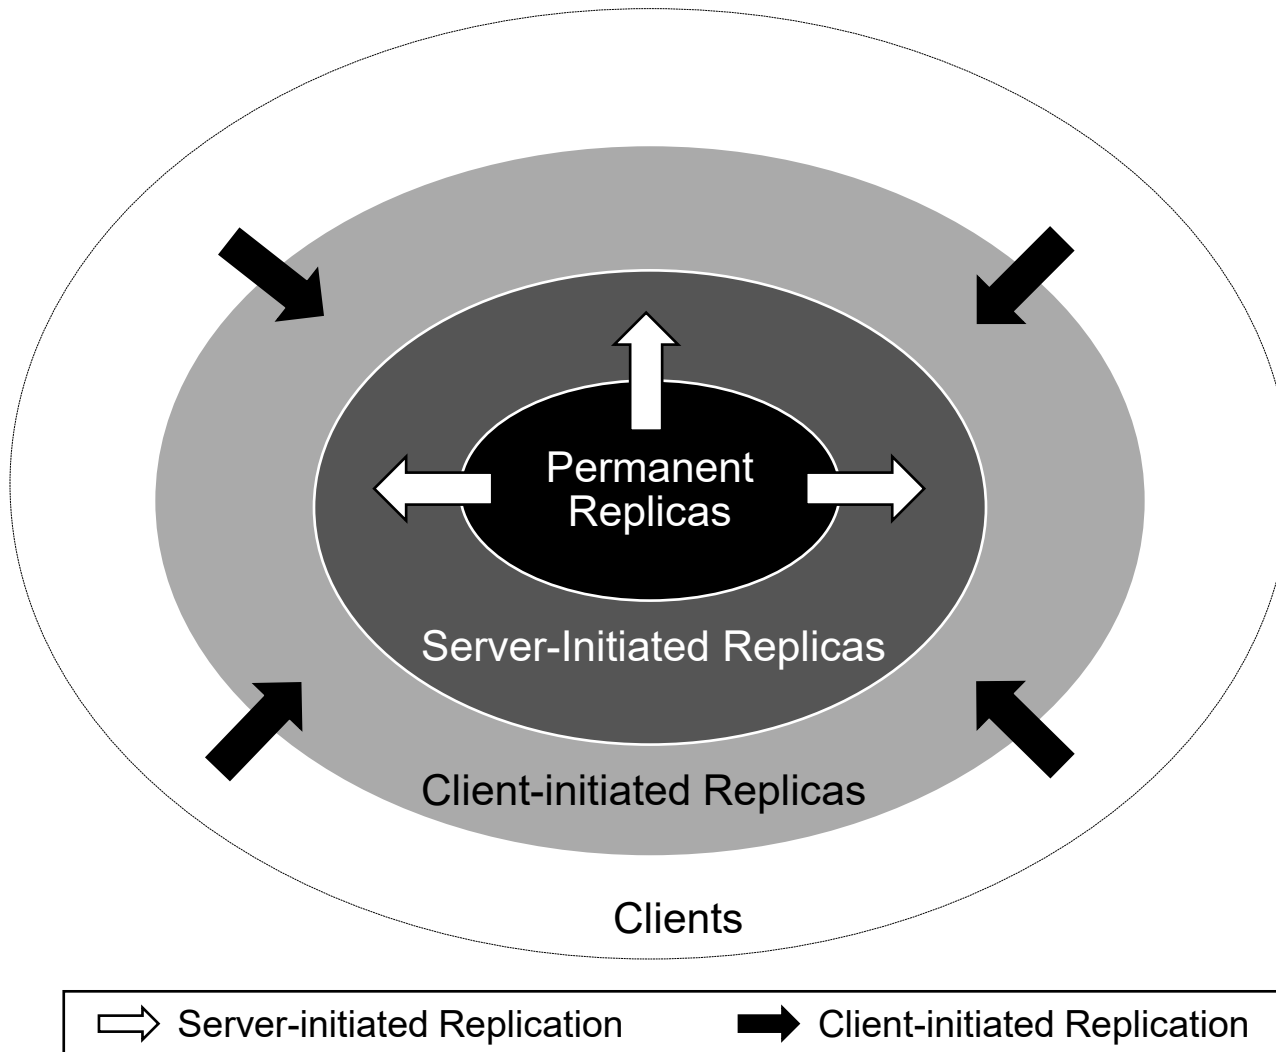
- Consistency Models
- Replica Management
 - ◆ Replica Server Placement
 - ◆ Content Replication and Placement
- Consistency Protocols

Content Replication and Placement

- In addition to the server placement, it is important:
 - ◆ how, when and by whom different data items (contents) are placed on possible replica servers
- Identify how webpage replicas are replicated:



Logical Organization of Replicas



1. Permanent Replicas

- Permanent replicas are the initial set of replicas that constitute a distributed data-store
- Typically, small in number
- There can be two types of permanent replicas:
 - ◆ Primary servers
 - One or more servers in an organization
 - Whenever a request arrives, it is forwarded into one of the primary servers
 - ◆ Mirror sites
 - Geographically spread, and replicas are generally statically configured
 - Clients pick one of the mirror sites to download the data

2. Server-initiated Replicas

- A third party (*provider*) owns the *secondary replica servers*, and they provide *hosting service*
 - ◆ The provider has a collection of servers across the Internet
 - ◆ The hosting service dynamically replicates files on different servers
 - Based on the popularity of the file in a region
- The permanent server chooses to host the data item on different *secondary replica servers*
- The scheme is efficient when updates are rare
- Examples of Server-initiated Replicas
 - ◆ Replicas in Content Delivery Networks (CDNs)

Dynamic Replication in Server-initiated Replicas

- **Dynamic replication at secondary servers:**
 - ◆ Helps to reduce the server load and improve client performance
 - ◆ But, replicas have to dynamically push the updates to other replicas
- + Rabinovich et al. [3] proposed a distributed scheme for replication:
 - + Each server keeps track of:
 - i. which is the closest server to the requesting client
 - ii. number of requests per file per closest server
 - + For example, each server Q keeps track of $\text{cnt}_Q(P, F)$ which denotes how many requests arrived at Q which are closer to server P (for a file F)
 - + If $\text{cnt}_Q(P, F) > 0.5 * \text{cnt}_Q(Q, F)$
 - + Request P to replicate a copy of file F
 - + If $\text{cnt}_P(P, F) < \text{LOWER_BOUND}$
 - + Delete the file at replica Q

If some other replica is nearer to the clients, request replication over that server

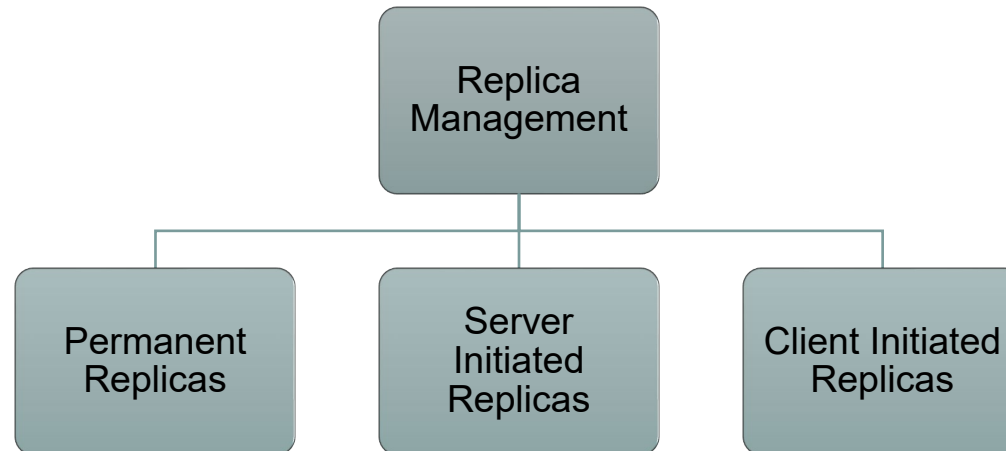
If the replication is not popular, delete the replica

3. Client-initiated Replicas

- Client-initiated replicas are known as *client caches*
- Client caches are used only to reduce the access latency of data
 - ◆ e.g., Browser caching a web-page locally
- Typically, managing a cache is entirely the responsibility of a client
 - ◆ Occasionally, data-store may inform client when the replica has become stale

Summary of Replica Management

- Replica management deals with placement of servers and content for improving performance and fault-tolerance



Till now, we know:

- how to place replica servers and content
- the required consistency models for applications

What else do we need to provide consistency in a distributed system?

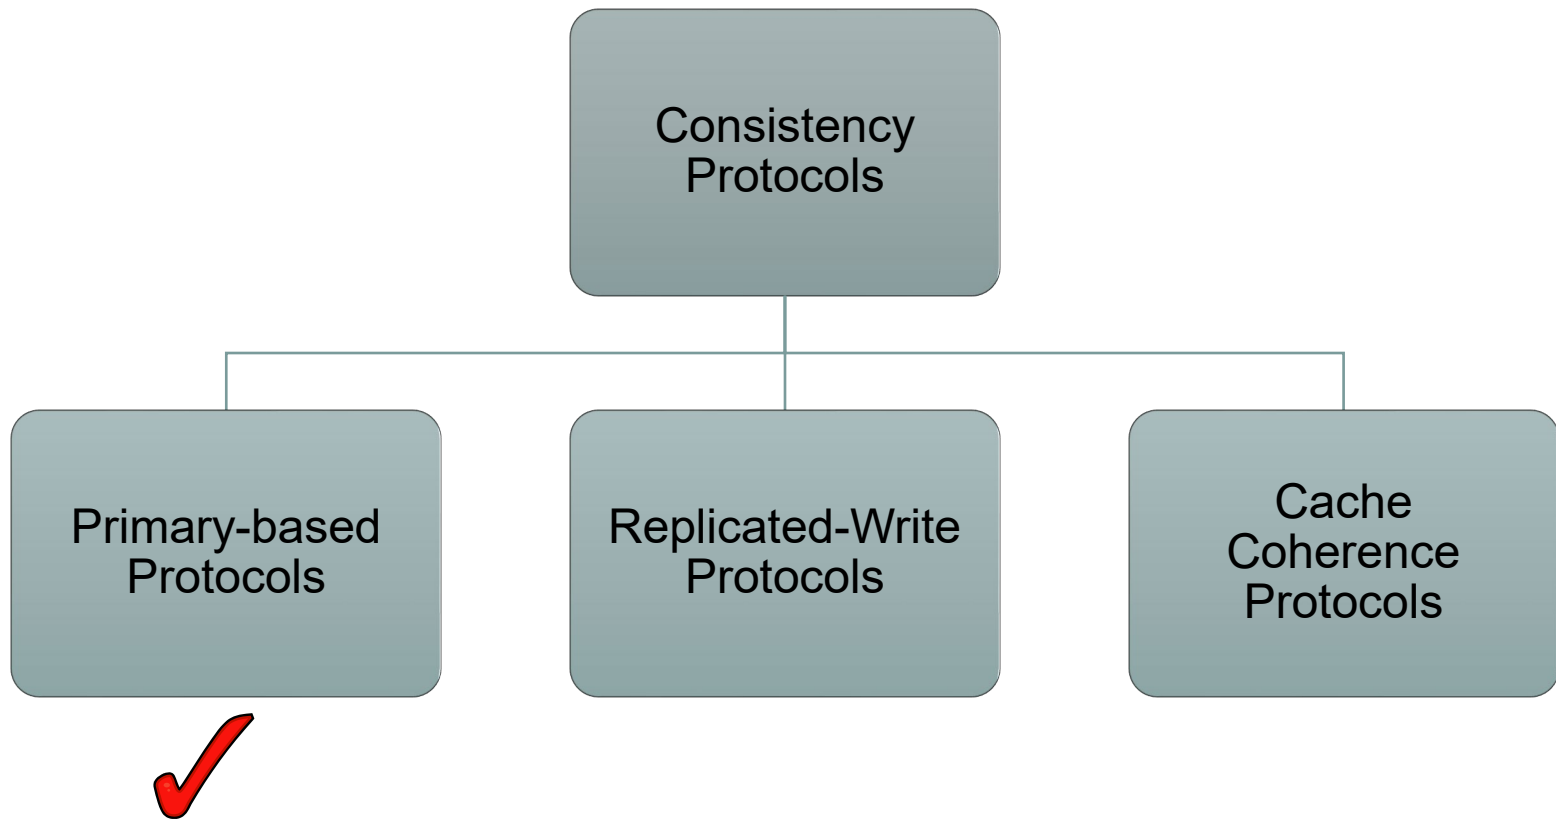
Overview

- Consistency Models
- Replica Management
- Consistency Protocols

Consistency Protocols

- A consistency protocol describes the implementation of a specific consistency model
- We are going to study three consistency protocols:
 - ◆ Primary-based protocols
 - One primary coordinator is elected to control replication across multiple replicas
 - ◆ Replicated-write protocols
 - Multiple replicas coordinate to provide consistency guarantees
 - ◆ Cache-coherence protocols
 - A special case of client-controlled replication

Overview of Consistency Protocols



Primary-based protocols

- In Primary-based protocols, a simple centralized design is used to implement consistency models
 - ◆ Each data-item **x** has an associated “*Primary Replica*”
 - ◆ Primary replica is responsible for coordinating write operations
- We will study one example of Primary-based protocols that implement Sequential Consistency Model
 - ◆ Remote-Write Protocol

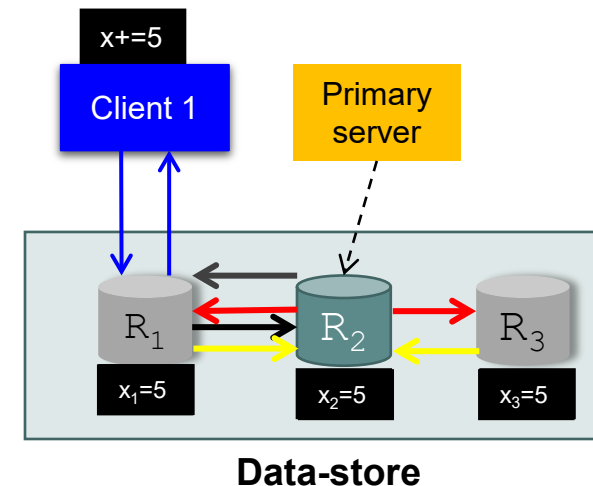
Remote-Write Protocol

■ Rules:

- ◆ All write operations are forwarded to the **primary replica**
- ◆ Read operations are carried out locally at each replica

■ Approach for write ops: (Budhiraja *et al.* [4])

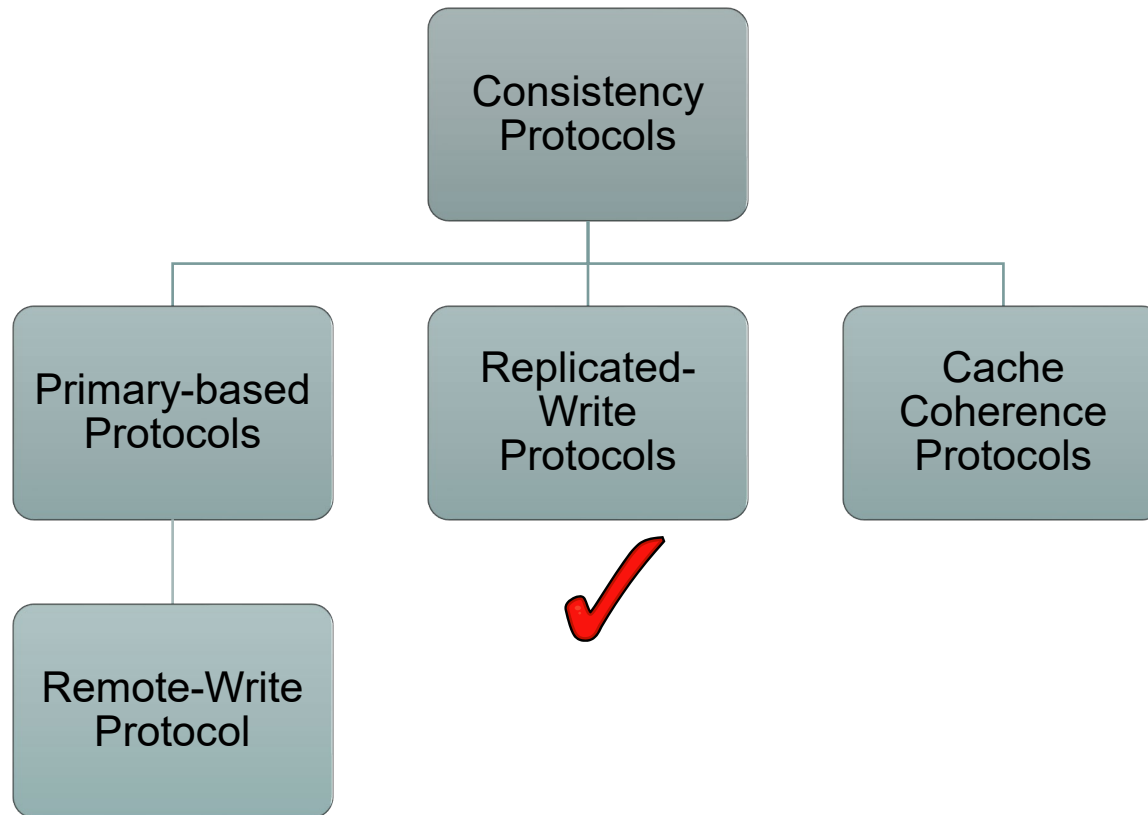
- + Client connects to some replica R_C
- + If the client issues write operation to R_C :
 - + R_C forwards the request to the primary replica R_P
 - + R_P updates its local value
 - + R_P forwards the update to other replicas R_i
 - + Other replicas R_i update, and send an ACK back to R_P
- + After R_P receives all ACKs, it informs the R_C that write operation is successful
- + R_C acknowledges to the client that write operation was successful



Remote-Write Protocol - Discussion

- Remote-Write provides
 - ◆ A simple way to implement sequential consistency
 - ◆ Guarantees that client see the most recent write operations
- However, latency is high in Remote-Write Protocols
 - ◆ Client blocks until all the replicas are updated
 - ◆ In what scenarios would you use remote-write protocols?
- Remote-Write Protocols are applied to distributed databases and file systems that require fault-tolerance
 - ◆ Replicas are placed on the same LAN to reduce latency

Overview of Consistency Protocols

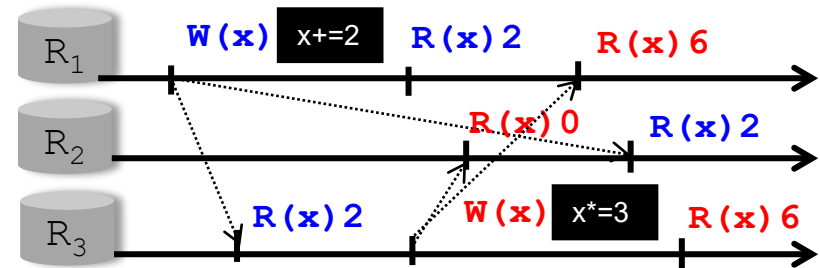
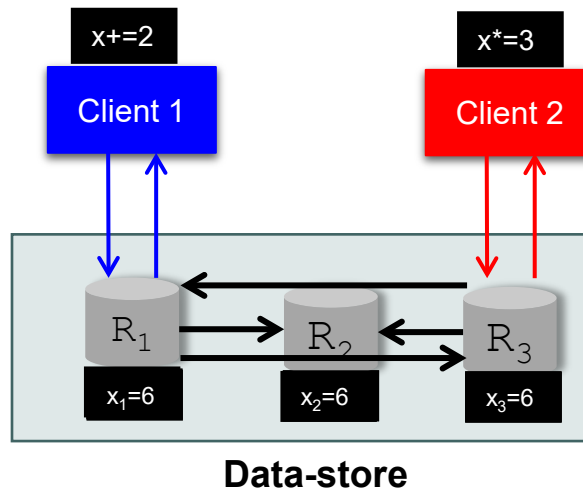


Replicated-Write Protocol

- In a replicated-write protocol, updates can be carried out at multiple replicas
- We will study one example replicated-write protocol called Active Replication Protocol
 - ◆ Here, clients write at any replica
 - ◆ The replica will propagate updates to other replicas

Active Replication Protocol

- When a client writes at a replica, the replica will send the write operation updates to all other replicas
- Challenges with Active Replication
 - ◆ Ordering of operations cannot be guaranteed across the replicas



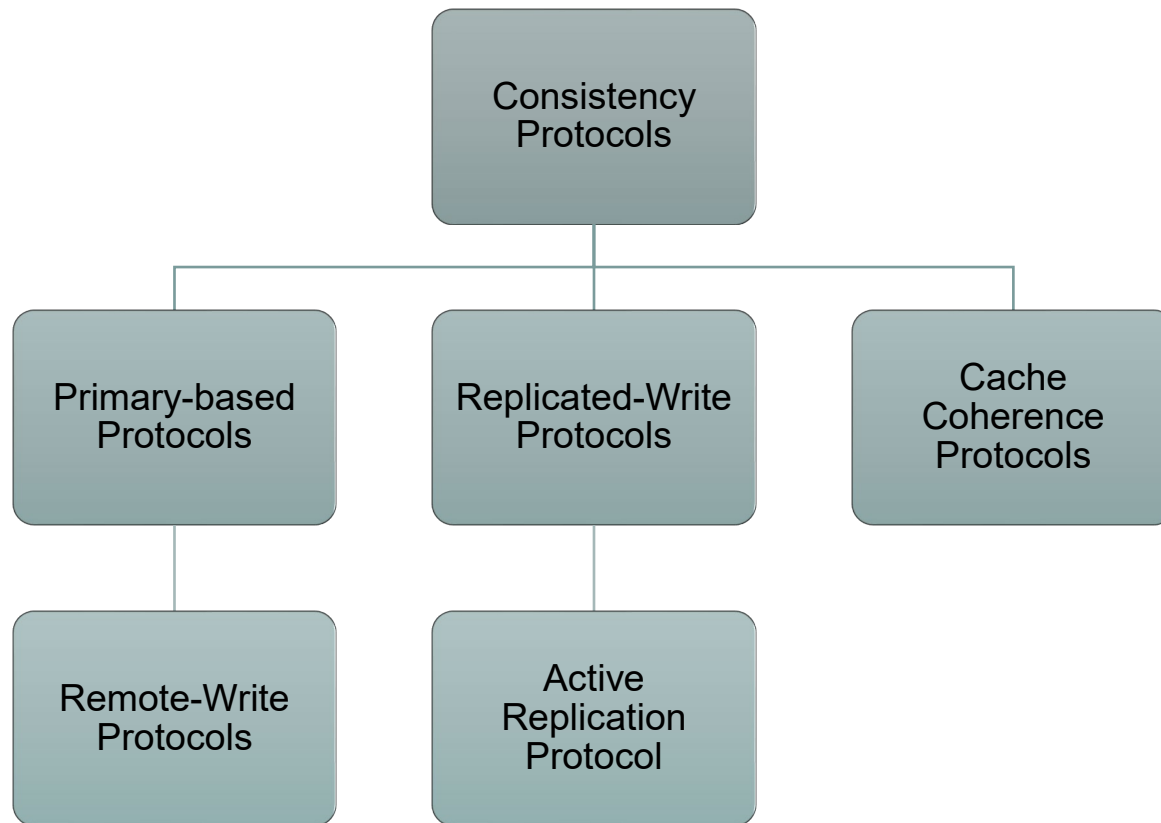
Centralized Active Replication Protocol

■ Approach

- There is a centralized coordinator called sequencer (**Seq**)
- When a client connects to a replica R_c and issues a write operation
 - ◆ R_c forwards the update to the **Seq**
 - ◆ **Seq** assigns a sequence number to the update operation
 - ◆ R_c propagates the sequence number and the operation to other replicas
- Operations are carried out at all the replicas in the order of the sequence number



Overview of Consistency Protocols



Cache Coherence Protocols

- Caches are special types of replicas
 - ◆ Typically, caches are client-controlled replicas
- Cache coherence refers to the consistency of data stored in caches
- How are the cache coherence protocols in shared-memory multiprocessor (SMP) systems different from those in Distributed Systems?
 - ◆ Coherence protocols in SMP assume cache states can be broadcasted efficiently
 - ◆ In DS, this is not possible because caches may reside on different machines

Cache Coherence Protocols (cont'd)

- Cache Coherence protocols determine how caches are kept consistent
- Caches may become inconsistent when data item is modified:
 1. at the server replicas, or
 2. at the cache

1. When Data is Modified at the Server

■ Two approaches for enforcing coherence:

1. Server-initiated invalidation

- Here, server sends all caches an invalidation message when data item is modified

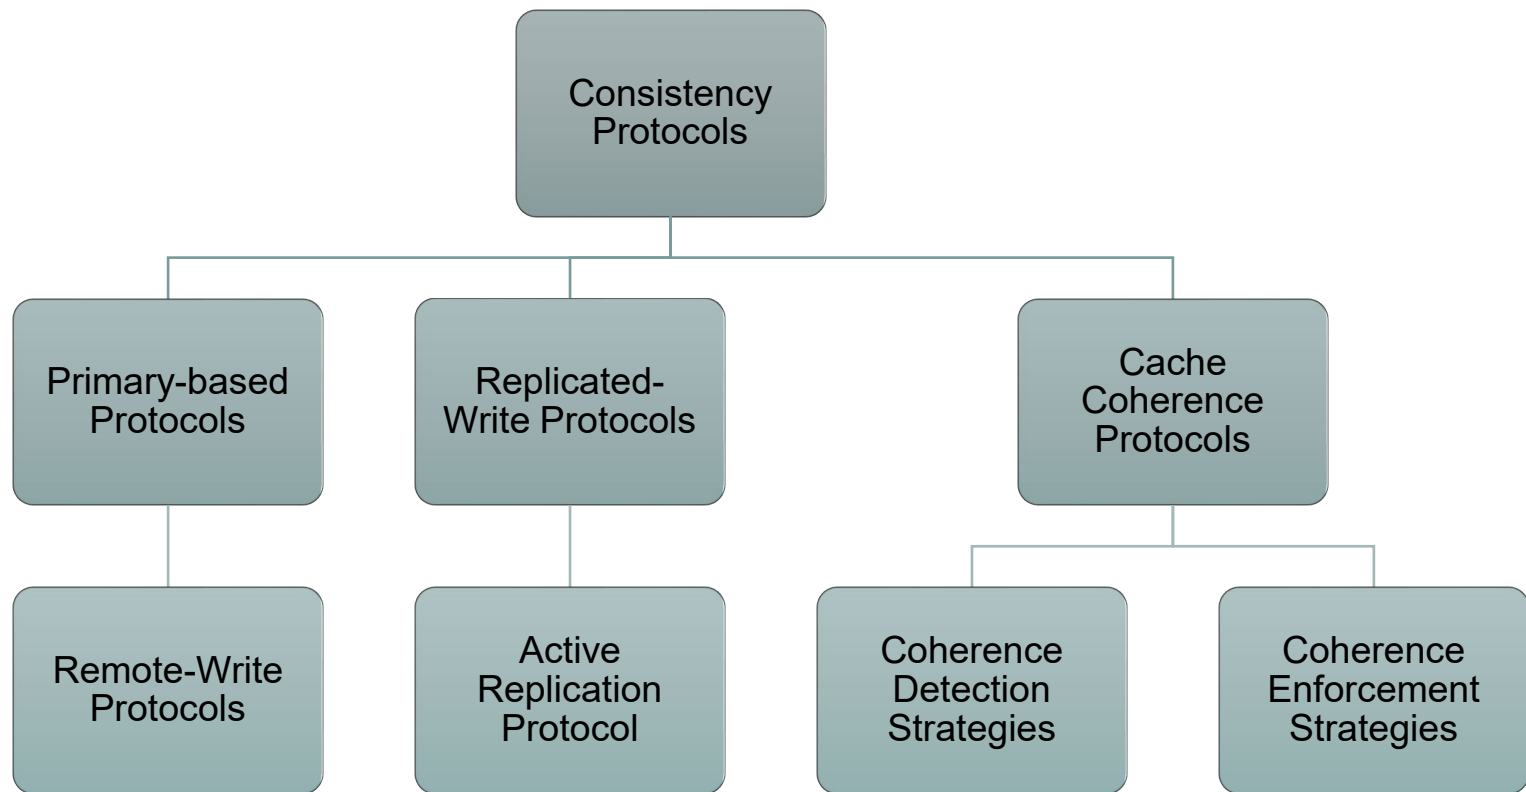
2. Server updates the cache

- Server will propagate the update to the cache

2. When Data is Modified at the Cache

- The enforcement protocol may use one of three techniques:
 - i. Read-only cache
 - The cache does not modify the data in the cache
 - The update is propagated to the server replica
 - ii. Write-through cache
 - Directly modify the cache, and forward the update to the server
 - iii. Write-back cache
 - The client allows multiple writes to take place at the cache
 - The client batches a set of writes, and will send the batched write updates to the server replica

Summary of Consistency Protocols



Consistency and Replication – Brief Summary

- Replication improves performance and fault-tolerance
- However, replicas have to be kept reasonably consistent

Consistency Models

- A contract between the data-store and process
- Types: Data-centric and Client-centric

Replication Management

- Describes where, when and by whom replicas should be placed
- Types: Replica Server Placement, Content Replication and Placement

Consistency Protocols

- Implement Consistency Models
- Types: Primary-based, Replicated-Write, Cache Coherence

Next Classes

- Fault-tolerance

- ◆ How to detect and deal with failures in Distributed Systems?

References

- [1] Terry, D.B., Demers, A.J., Petersen, K., Spreitzer, M.J., Theimer, M.M., Welch, B.B., "Session guarantees for weakly consistent replicated data", Proceedings of the Third International Conference on Parallel and Distributed Information Systems, 1994
- [2] Lili Qiu, Padmanabhan, V.N., Voelker, G.M., "On the placement of Web server replicas", Proceedings of IEEE INFOCOM 2001.
- [3] Rabinovich, M., Rabinovich, I., Rajaraman, R., Aggarwal, A., "A dynamic object replication and migration protocol for an Internet hosting service", Proceedings of IEEE International Conference on Distributed Computing Systems (ICDCS), 1999
- [4] Navin Budhiraja, Keith Marzullo. Fred B. Schneider. Sam Toueg, "The primary-backup approach", Distributed systems (2nd Ed.), ACM Press/Addison-Wesley Publishing Co., 1993
- [5] <http://www.cdk5.net>
- [6] http://en.wikipedia.org/wiki/Cache_coherence

A General Background

- Basic Concepts
- Failure Models
- Failure Masking by Redundancy

Today...

- Last Sessions
 - Consistency and Replication
 - Consistency Models: Data-centric and Client-centric
 - Replica Management
 - Consistency Protocols
- Today's session
 - Fault Tolerance
 - General background
 - Process resilience and failure detection

A General Background

- Basic Concepts
- Failure Models
- Failure Masking by Redundancy

Dependability 可信任性

Basics: A component provides services to clients. To provide services, the component may require the services from other components \Rightarrow a component may depend on some other component.

Specifically: A component C depends on C^* if the correctness of C 's behavior depends on the correctness of C^* 's behavior.

Some properties of dependability:

- **Availability**: Readiness for usage
- **Reliability**: Continuity of service delivery
- **Safety**: Very low probability of catastrophes
- **Maintainability**: How easy can a failed system be repaired

Note: For distributed systems, components can be either processes or channels

Terminology

Failure: When a component is not living up to its specifications, a failure occurs

失效

Error: That part of a component's state that can lead to a failure

错误

Fault: The cause of an error

故障

Fault prevention: prevent the occurrence of a fault

Fault tolerance: build a component in such a way that it can meet its specifications in the presence of faults (i.e., mask the presence of faults)

Fault removal: reduce the presence, number, seriousness of faults

Fault forecasting: estimate the present number, future incidence, and the consequences of faults

Failures, Due to What?

失效

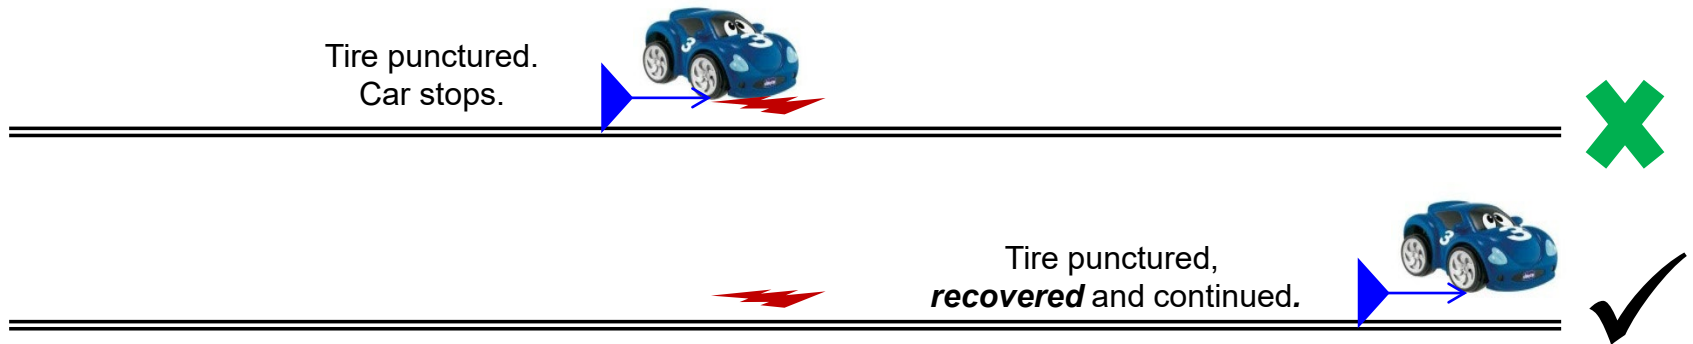
- Failures can happen due to a variety of reasons:
 - Hardware faults 失效
 - Software bugs
 - Operator errors 错误
 - Network errors/outages 断网
- A system is said to *fail* when it cannot meet its promises

Failures in Distributed Systems

- A characteristic feature of distributed systems that distinguishes them from single-machine systems is the notion of **partial failure**
- A partial failure may happen when a component in a distributed system fails
 - This failure may affect the proper operation of other components, while at the same time leaving yet other components unaffected

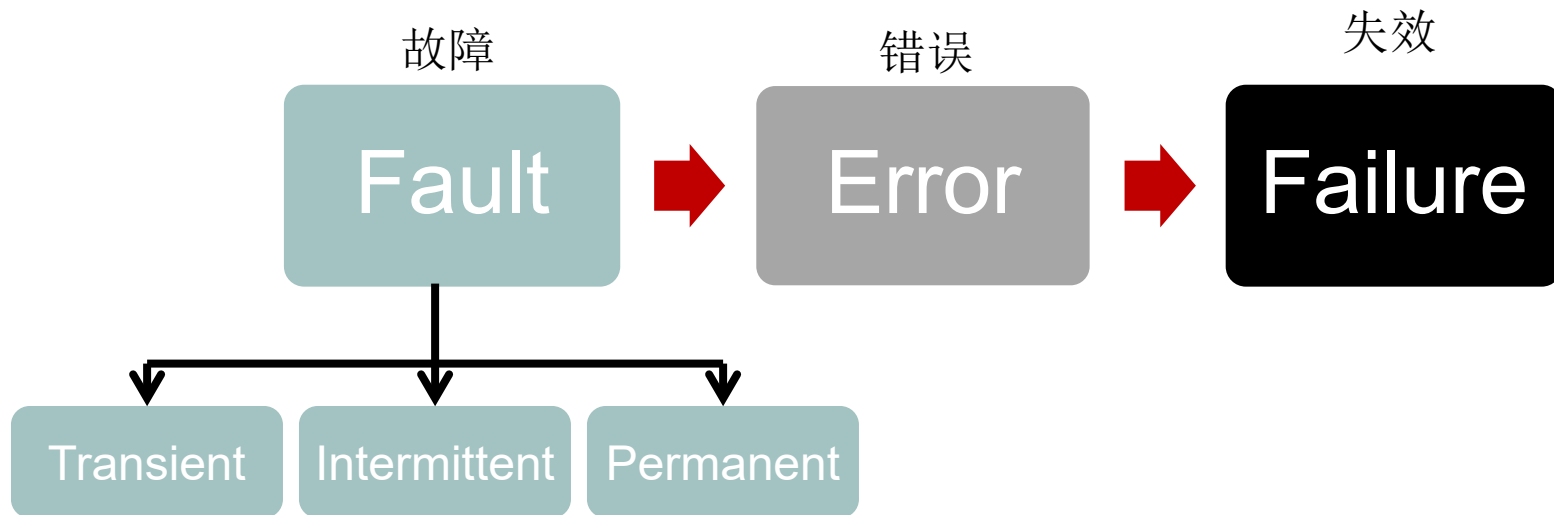
Goal and Fault-Tolerance

- An overall goal in distributed systems is to construct the system in such a way that it can automatically recover from partial failures*



- Fault-tolerance** is the property that enables a system to continue operating properly in the event of failures
- For example, TCP is designed to allow reliable two-way communication in a packet-switched network, even in the presence of communication links which are imperfect or overloaded

Faults, Errors and Failures



A system is said to be *fault tolerant* if it can provide its services even in the presence of *faults*

Fault Tolerance Requirements

- A robust fault tolerant system requires:
 1. No single point of failure
 2. Fault isolation/containment to the failing component
 3. Availability of reversion modes

Dependable Systems 可信系统

- Being fault tolerant is strongly related to what is called a *dependable system*



A General Background

- Basic Concepts
- **Failure Models**
- Failure Masking by Redundancy

Failure Models

失效
模式

Type of Failure	Description
<ul style="list-style-type: none">Crash Failure	<ul style="list-style-type: none">A server halts, but was working correctly until it stopped
<ul style="list-style-type: none">Omission Failure<ul style="list-style-type: none">Receive OmissionSend Omission	<ul style="list-style-type: none">A server fails to respond to incoming requests<ul style="list-style-type: none">A server fails to receive incoming messagesA server fails to send messages
<ul style="list-style-type: none">Timing Failure	<ul style="list-style-type: none">A server's response lies outside the specified time interval
<ul style="list-style-type: none">Response Failure<ul style="list-style-type: none">Value FailureState Transition Failure	<ul style="list-style-type: none">A server's response is incorrect<ul style="list-style-type: none">The value of the response is wrongThe server deviates from the correct flow of control
<ul style="list-style-type: none">Byzantine Failure	<ul style="list-style-type: none">A server may produce arbitrary responses at arbitrary times

A General Background

- Basic Concepts
- Failure Models
- Failure Masking by Redundancy

Faults Masking by Redundancy

- The key technique for masking faults is to use *redundancy*

Usually, extra bits are added to allow recovery from garbled bits

Information

Usually, extra processes are added to allow tolerating failed processes

Software

Redundancy

Hardware

Usually, extra equipment are added to allow tolerating failed hardware components

Time

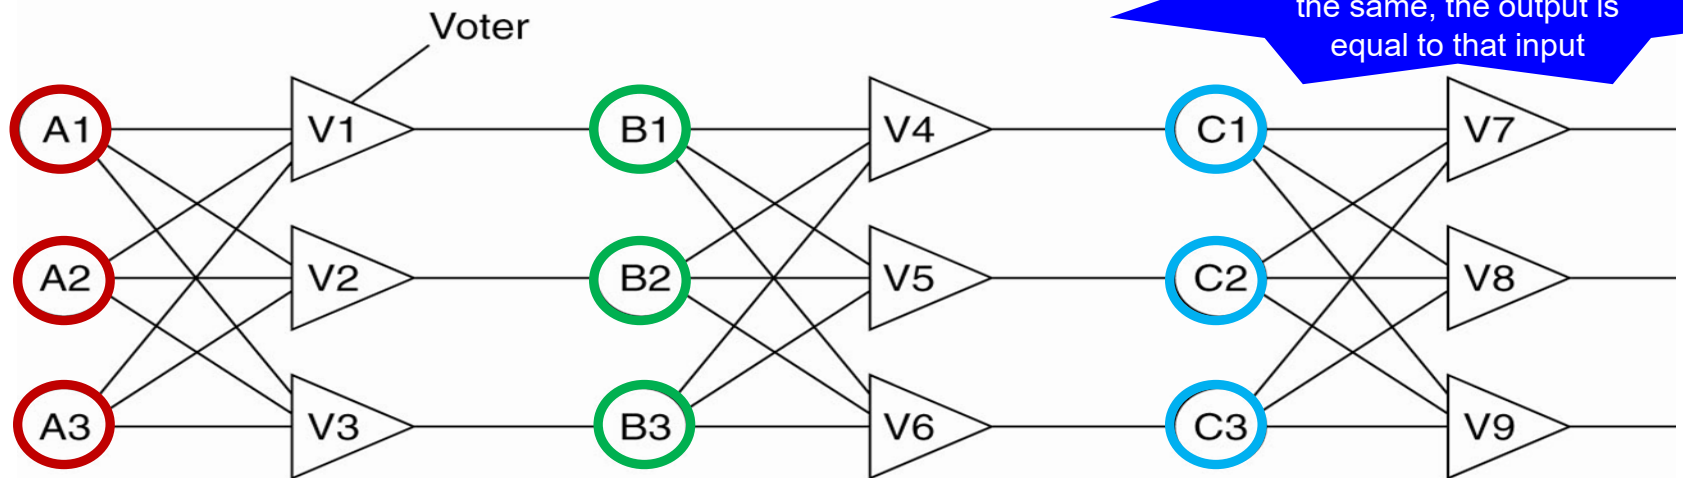
Usually, an action is performed, and then, if required, it is performed again

Triple Modular Redundancy



If one is faulty, the final result will be incorrect

A circuit with signals passing through devices A, B, and C, in sequence



If 2 or 3 of the inputs are the same, the output is equal to that input

Each device is replicated 3 times and after each stage is a triplicated voter

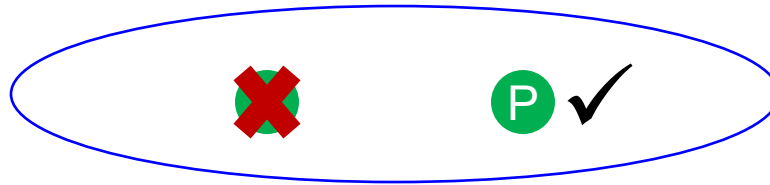
Process Resilience and Failure Detection

Process Resilience and Failure Detection

- Now that the basic issues of fault tolerance have been discussed, let us concentrate on how fault tolerance can actually be achieved in distributed systems
- The topics we will discuss:
 - How can we provide protection against process failures?
 - Process Groups
 - Reaching an agreement within a process group
 - How to detect failures?

Process Resilience

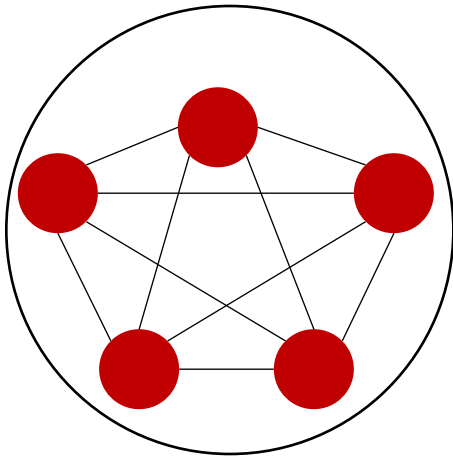
- The key approach to tolerating a faulty process is to organize several identical processes into a *group*



- If one process in a group fails, hopefully some other process can take over
- **Caveats:** 注意事项
 - A process can join a group or leave one during system operation
 - A process can be a member of several groups at the same time

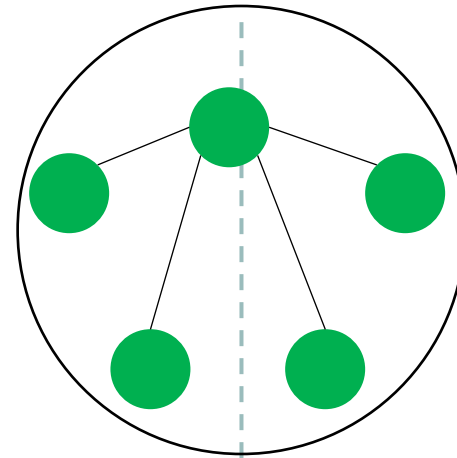
Flat Versus Hierarchical Groups

- An important distinction between different groups has to do with their internal structure



Flat Group:

- (+) Symmetrical
- (+) No single point of failure
- (-) Decision making is complicated



Hierarchical Group:

- (+) Decision making is simple
- (-) Asymmetrical
- (-) Single point of failure



K-Fault-Tolerant Systems

- A system is said to be *k-fault-tolerant* if it can survive faults in *k* components and still meet its specifications
- How can we achieve a *k-fault-tolerant* system?
 - This would require an agreement protocol applied to a process group

Agreement in Faulty Systems (1)

- A process group typically requires reaching an *agreement* in:
 - Electing a coordinator
 - Deciding whether or not to commit a transaction
 - Dividing tasks among workers
 - Synchronization
- When the communication and processes:
 - are perfect, reaching an agreement is often straightforward
 - are not perfect, there are problems in reaching an agreement

Agreement in Faulty Systems (2)

- **Goal:** have all non-faulty processes reach consensus on some issue, and establish that consensus within a finite number of steps
- Different assumptions about the underlying system require different solutions:
 - Synchronous versus asynchronous systems
 - Communication delay is bounded or not
 - Message delivery is ordered or not
 - Message transmission is done through unicasting or multicasting

Agreement in Faulty Systems (3)

- Reaching a distributed agreement is only possible in the following circumstances:

		Message Ordering					
		Unordered		Ordered			
Process Behavior	Synchronous	✓	✓	✓	✓	Bounded	Communication Delay
				✓	✓	Unbounded	
	Asynchronous				✓	Bounded	
					✓	Unbounded	
		Unicast	Multicast	Unicast	Multicast		
		Message Transmission					

Agreement in Faulty Systems (4)

- In practice most distributed systems assume that:
 - Processes behave asynchronously
 - Message transmission is unicast
 - Communication delays are unbounded
- Usage of ordered (reliable) message delivery is typically required
- The agreement problem has been originally studied by Lamport and referred to as the Byzantine Agreement Problem [Lamport *et al.*]

Byzantine Agreement Problem (1)

- Lamport assumes:
 - Processes are synchronous
 - Messages are unicast while preserving ordering
 - Communication delay is bounded
 - There are N processes, where each process i will provide a value v_i to the others
 - There are at most k faulty processes

Byzantine Agreement Problem (2)

□ Lamport's Assumptions:

- Processes are synchronous
- Messages are unicast while preserving ordering
- Communication delay is bounded
- There are N processes, where each process i will provide a value v_i to the others
- There are at most k faulty processes

		Message Ordering					
		Unordered		Ordered			
Process Behavior	Synchronous	✓	✓	✓	✓	Bounded	Communication Delay
				✓	✓	Unbounded	
	Asynchronous				✓	Bounded	
					✓	Unbounded	
		Unicast	Multicast	Unicast	Multicast		
		Message Transmission					

Lamport suggests that each process i constructs a vector V of length N , such that if process i is non-faulty, $V[i] = v_i$. Otherwise, $V[i]$ is undefined



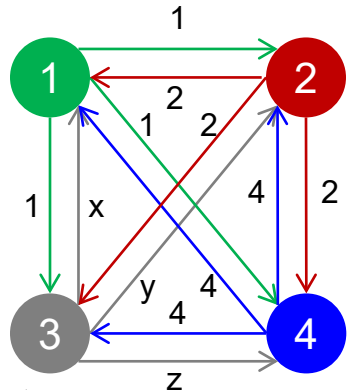
Byzantine Agreement Problem (3)

- Case I: $N = 4$ and $k = 1$

Step1: Each process sends its value to the others

Step2: Each process collects values received in a vector

Step3: Every process passes its vector to every other process



Faulty
process

1 Got(1, 2, x, 4)
2 Got(1, 2, y, 4)
3 Got(1, 2, 3, 4)
4 Got(1, 2, z, 4)

1 Got

(1, 2, y, 4)
(a, b, c, d)
(1, 2, z, 4)

2 Got

(1, 2, x, 4)
(e, f, g, h)
(1, 2, z, 4)

4 Got

(1, 2, x, 4)
(1, 2, y, 4)
(i, j, k, l)



Byzantine Agreement Problem (4)

Step 4:

- Each process examines the i th element of each of the newly received vectors
- If any value has a majority, that value is put into the result vector
- If no value has a majority, the corresponding element of the result vector is marked UNKNOWN

The algorithm reaches an agreement

Result Vector:

(1, 2, UNKNOWN, 4)

Result Vector:

(1, 2, UNKNOWN, 4)

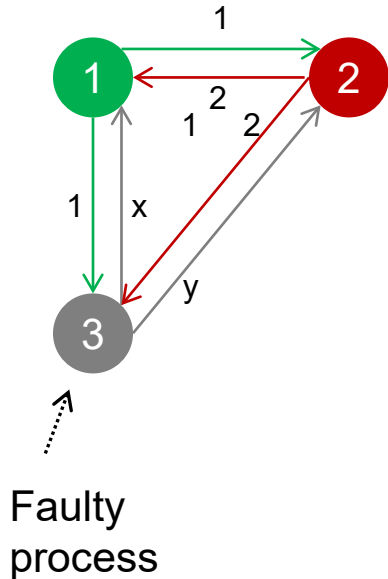
Result Vector:

(1, 2, UNKNOWN, 4)

Byzantine Agreement Problem (5)

- Case II: $N = 3$ and $k = 1$

Step1: Each process sends its value to the others



Step2: Each process collects values received in a vector

1 Got(1, 2, x)
2 Got(1, 2, y)
3 Got(1, 2, 3)

Step3: Every process passes its vector to every other process

1 Got
(1, 2, y)
(a, b, c)

2 Got
(1, 2, x)
(d, e, f)



Byzantine Agreement Problem (6)

Step 4:

- Each process examines the i th element of each of the newly received vectors
- If any value has a majority, that value is put into the result vector
- If no value has a majority, the corresponding element of the result vector is marked UNKNOWN

1 Got

(1, 2, y)
(a, b, c)

The algorithm has
failed to produce
an agreement

2 Got

(1, 2, x)
(d, e, f)

Result Vector:

(UNKNOWN, UNKNOWN, UNKNOWN)

Result Vector:

(UNKNOWN, UNKNOWN, UNKNOWN)

Concluding Remarks on the Byzantine Agreement Problem

- In their paper, *Lamport et al.* (1982) proved that in a system with k faulty processes, an agreement can be achieved only if $2k+1$ correctly functioning processes are present, for a total of $3k+1$.
 - i.e., An agreement is possible only if more than two-thirds of the processes are working properly.
- *Fisher et al.* (1985) proved that in a distributed system in which ordering of messages **cannot** be guaranteed to be delivered within a known, finite time, no agreement is possible even if only one process is faulty.

Process Failure Detection

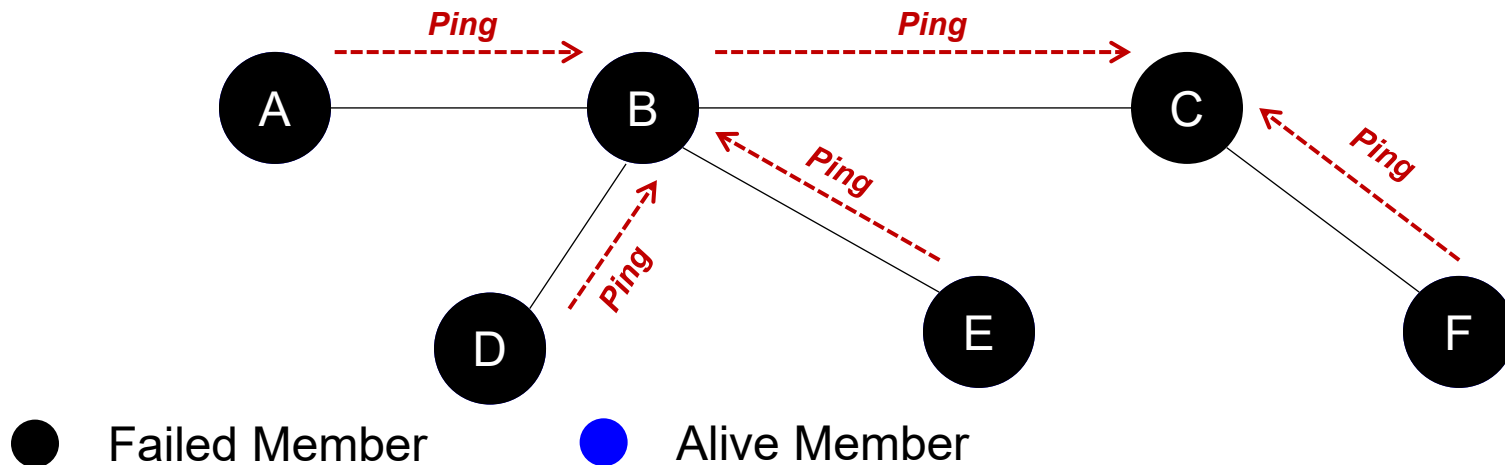
- Before we properly mask failures, we generally need to detect them
- For a group of processes, non-faulty members should be able to decide who is still a member and who is not
- Two policies:
 - Processes actively send “are you alive?” messages to each other (i.e., *pinging each other*)
 - Processes passively wait until messages come in from different processes

Timeout Mechanism

- In failure detection a [timeout mechanism](#) is usually involved
- Specify a **timer**, after a period of time, trigger a timeout
- However, due to unreliable networks, simply stating that a process has failed because it does not return an answer to a ping message may be wrong

Example: FUSE

- In FUSE, processes can be joined in a group that spans a WAN
- The group members create a spanning tree that is used for monitoring member failures
- An active (pinging) policy is used where a single node failure is rapidly promoted to a group failure notification



Failure Considerations

- There are various issues that need to be taken into account when designing a failure detection subsystem:
 1. Failure detection can be done as a side-effect of regularly exchanging information with neighbors (e.g., ***gossip-based information dissemination***)
 2. A failure detection subsystem should ideally be able to distinguish network failures from node failures
 3. When a member failure is detected, how should other non-faulty processes be informed

Recovery

- So far, we have mainly concentrated on algorithms that allow us to tolerate faults
- However, once a failure has occurred, it is essential that the process where the failure has happened can *recover* to a *correct state*
- In what follows we focus on:
 - What it actually means to recover to a correct state
 - When and how the state of a distributed system can be recorded and recovered, by means of *checkpointing* and *message logging*

Recovery

- Error Recovery
- Checkpointing
- Message Logging

Recovery

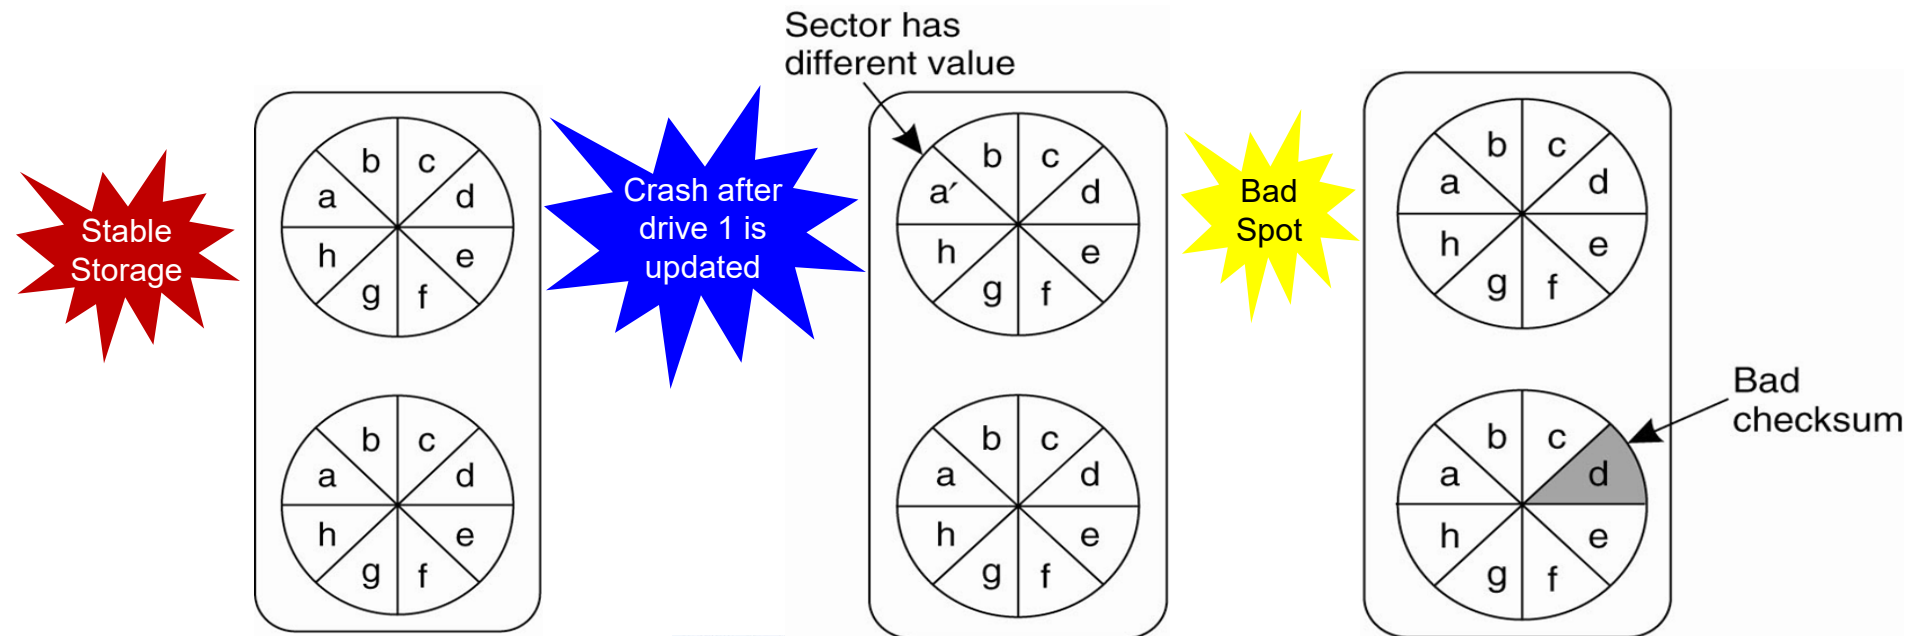
- Error Recovery
- Checkpointing
- Message Logging

Error Recovery

- Once a failure has occurred, it is essential that the process where the failure has happened can recover to a correct state
- Fundamental to fault tolerance is the recovery from an error
- The idea of *error recovery* is to replace an erroneous state with an error-free state
- There are essentially two forms of error recovery:
 1. *Backward recovery*
 2. *Forward recovery*

1. Backward Recovery (1)

- In backward recovery, the main issue is to bring the system from its present erroneous state back to a previously correct state
- It is necessary to record the system's state *from time to time* onto a stable storage, and to restore such a recorded state when things go wrong



1. Backward Recovery (2)

- Each time (part of) the system's present state is recorded, a checkpoint is said to be made
- Problems with backward recovery:
 - Restoring a system or a process to a previous state is generally expensive in terms of performance
 - Some states can never be rolled back (e.g., typing in UNIX `rm -fr *`)

2. Forward Recovery

- When the system detects that it has made an error, forward recovery reverts the system state to error time and corrects it, to be able to move forward
- Forward recovery is typically faster than backward recovery but requires that it has to be known in advance which errors may occur
- Some systems make use of both forward and backward recovery for different errors or different parts of one error

Recovery

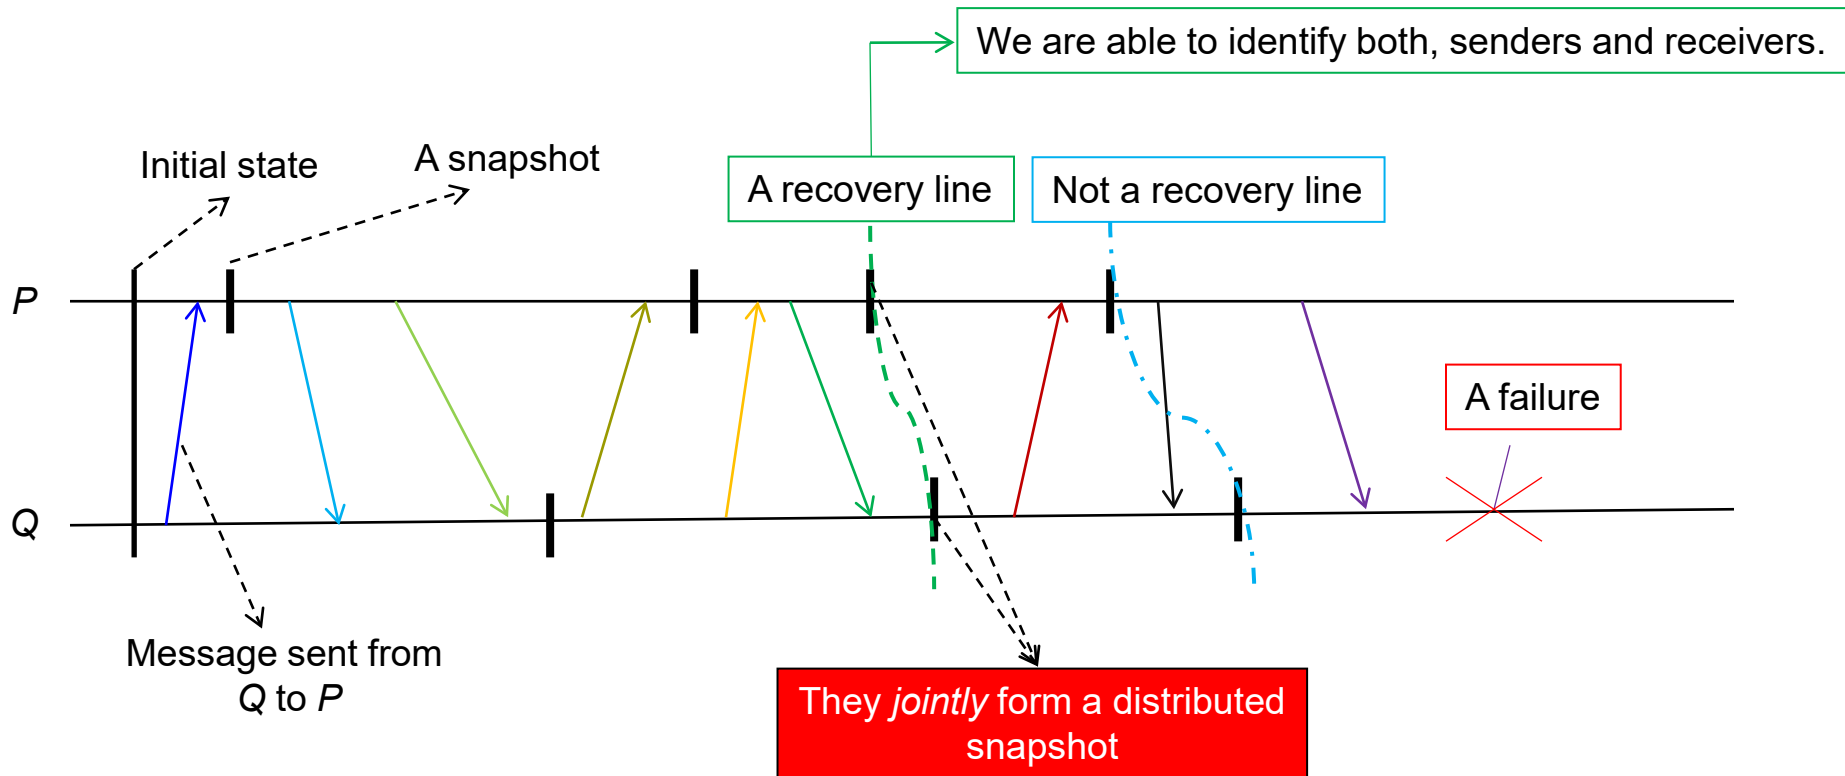
- Error Recovery
- Checkpointing
- Message Logging

Why Checkpointing?

- In a fault-tolerant distributed system, backward recovery requires that the system regularly saves its state onto a stable storage
- This process is referred to as **checkpointing**
- In particular, checkpointing consists of storing a distributed snapshot of the current application state (i.e., a consistent global state), and later on, use it for restarting the execution in case of a failure

Recovery Line

- In a distributed snapshot, if a process P has recorded the receipt of a message, then there should be also a process Q that has recorded the sending of that message

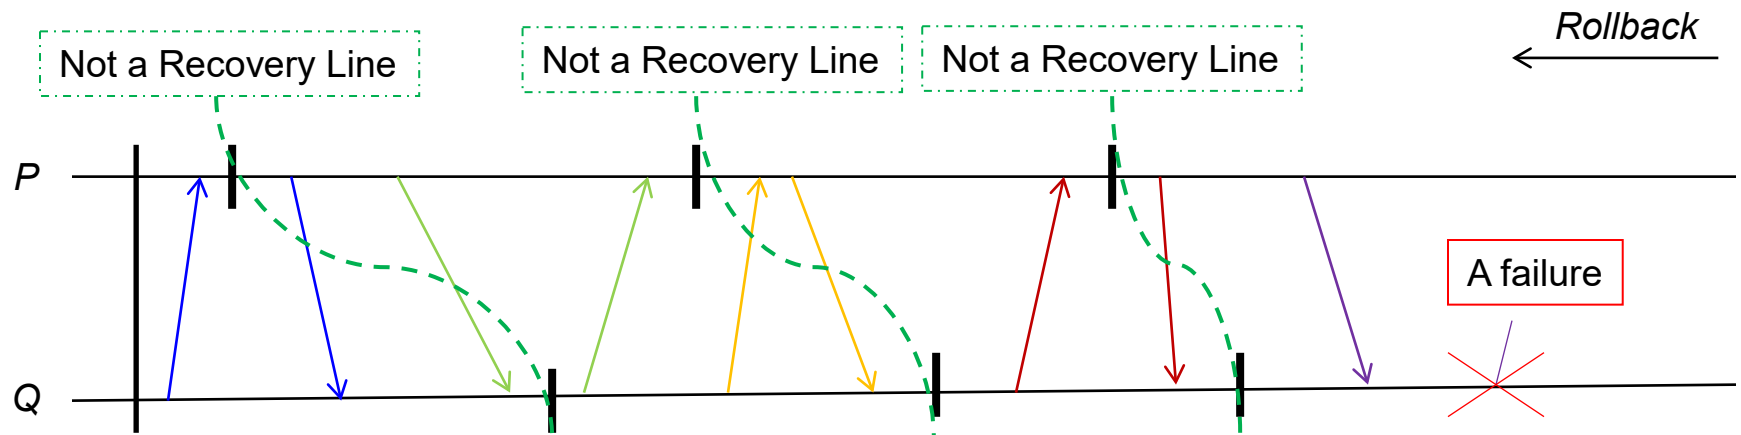


Checkpointing

- Checkpointing can be of two types:
 1. Independent Checkpointing: each process simply records its local state from time to time in an uncoordinated fashion
 2. Coordinated Checkpointing: all processes synchronize to jointly write their states to local stable storages

Domino Effect

- Independent checkpointing may make it difficult to find a recovery line, leading potentially to a domino effect resulting from *cascaded* rollbacks



- With coordinated checkpointing, the saved state is automatically globally consistent, hence, domino effect is inherently avoided

Recovery

- Error Recovery
- Checkpointing
- Message Logging

Why Message Logging?

- Considering that checkpointing is an expensive operation, techniques have been sought to reduce the number of checkpoints, but still enable recovery
- An important technique in distributed systems is [message logging](#)
- The basic idea is that if transmission of messages can be replayed, we can still reach a globally consistent state but without having to restore that state from stable storage
- In practice, the combination of having fewer checkpoints and message logging is more efficient than having to take many checkpoints

Message Logging

- Message logging can be of two types:
 1. Sender-based logging: A process can log its messages before sending them off
 2. Receiver-based logging: A receiving process can first log an incoming message before delivering it to the application
- When a sending or a receiving process crashes, it can restore the most recently checkpointed state, and from there on replay the logged messages (important for non-deterministic behaviors)

Replay of Messages and Orphan Processes

- Incorrect replay of messages after recovery can lead to *orphan* processes. This should be avoided

